



Final Exam Review

Introduction to Computer Graphics

Yu-Ting Wu

About Final Exam

- The final exam will be held physically on **Dec. 26th**, 2022
 - 09:10 – 12:00
 - Room 103
 - 35% of your final grade
- **Scope (theoretical part only)**
 - Scene representation
 - Camera
 - Lighting and shading
 - Textures
 - Transparency
 - GPU graphics pipeline
 - Deferred shading
 - Shadow map
 - Ambient occlusion
 - Ray tracing

About the 17th and 18th Weeks

- There is **NO** class in the 17th week (Jan. 2nd, 2023)
- There is a **physical class** (and recording) in the 18th week
 - Advanced shaders in OpenGL 3.0 and 4.0
 - A case study of the rendering system in Unity
 - You can take back your answering sheet for the final exam

About Homework

- The grade of your HW2 has been announced
 - If you have questions, feel free to mail me
- **The final deadline for all homework is Jan. 8th, 2023**
 - **You can resubmit previous HW if you think it is worth it (but let me know by mail)**

Computer Graphics Overview

- **Digitally synthesize and manipulate** a virtual world



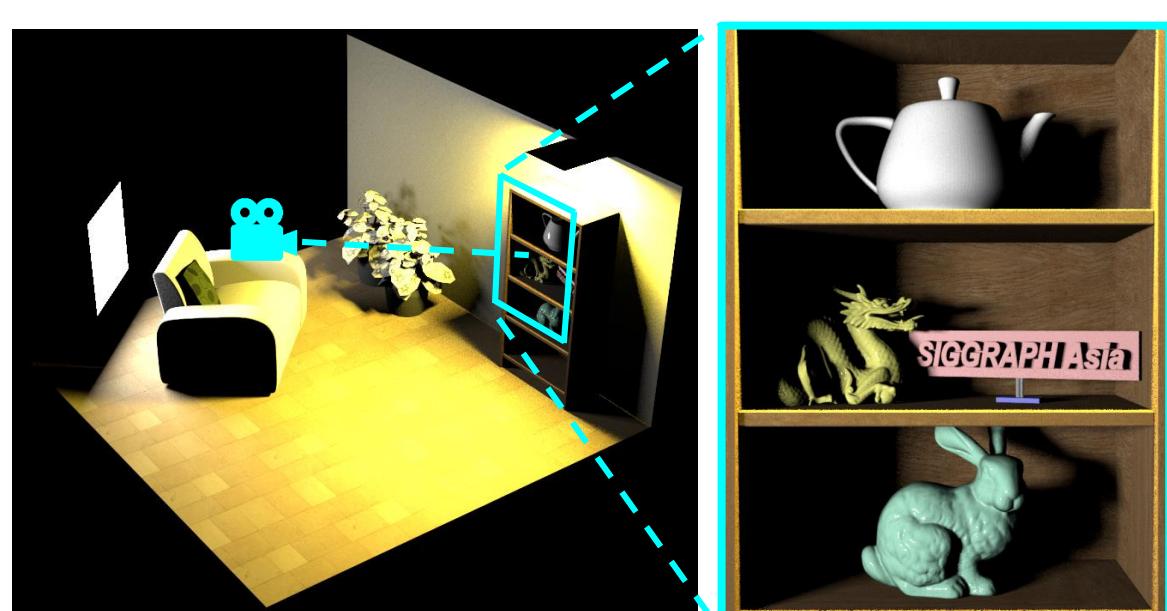
Computer Graphics Overview (cont.)

- **Digitally synthesize and manipulate** a virtual world
 - Model **geometry** of the 3D objects using **triangle meshes**
 - Model **materials** of the 3D objects (*Phong* lighting model) and simulate **lighting** (point/directional/spot lights)
 - Simulate more realistic materials (**microfacet models** and **transparency**) and lighting phenomena (shadow and **ambient occlusion**)
 - Simulate more complex light paths (**global illumination**)



Digital Image Synthesis

- Most displays are 2D, so we need to generate images from the 3D world
- Just like taking a picture with a camera in our daily lives
 - But with a **virtual camera** and a **virtual film**



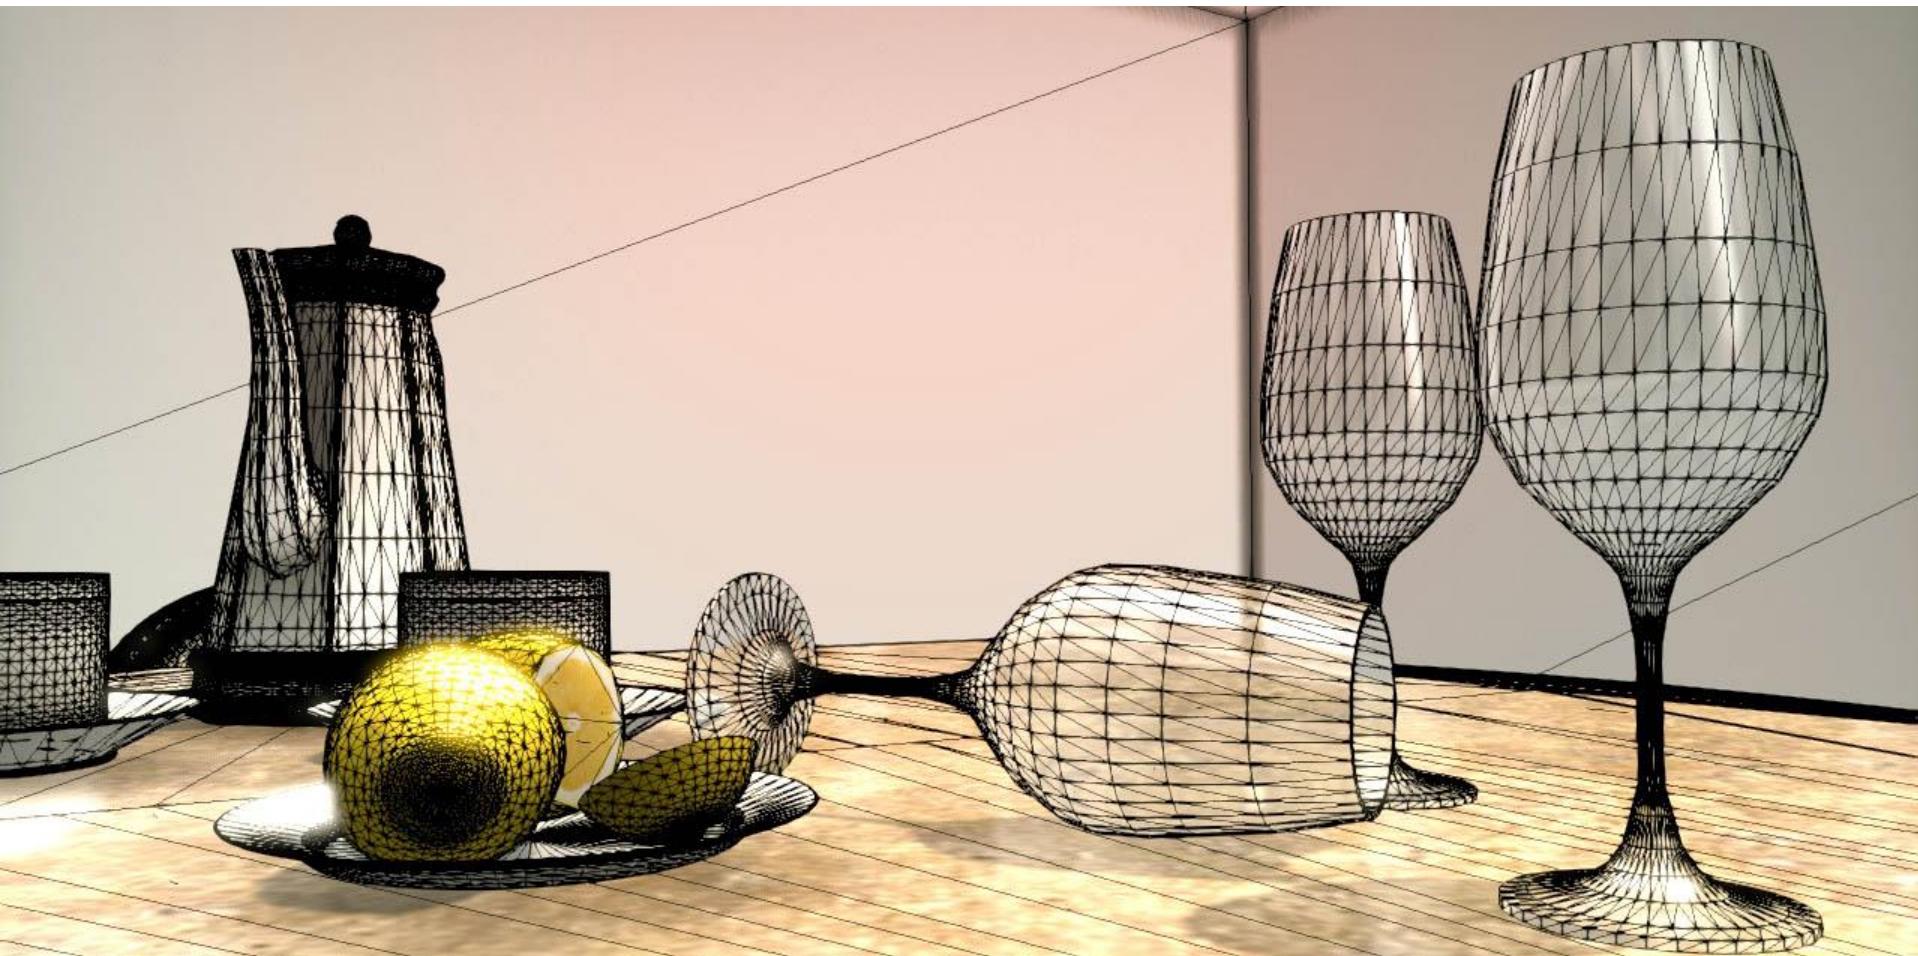
3D virtual world

rendered image

Scene Representation

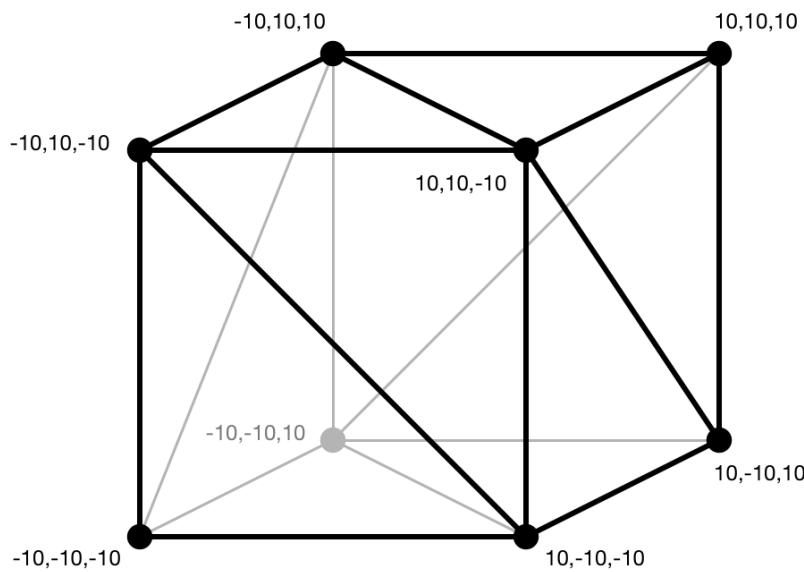
Triangle Mesh

- Use triangle meshes to model the geometry of 3D objects



Triangle Mesh (cont.)

- Vertex attributes: 3D position, normal, texture coordinate
- Vertex adjacencies: which vertices form a triangle



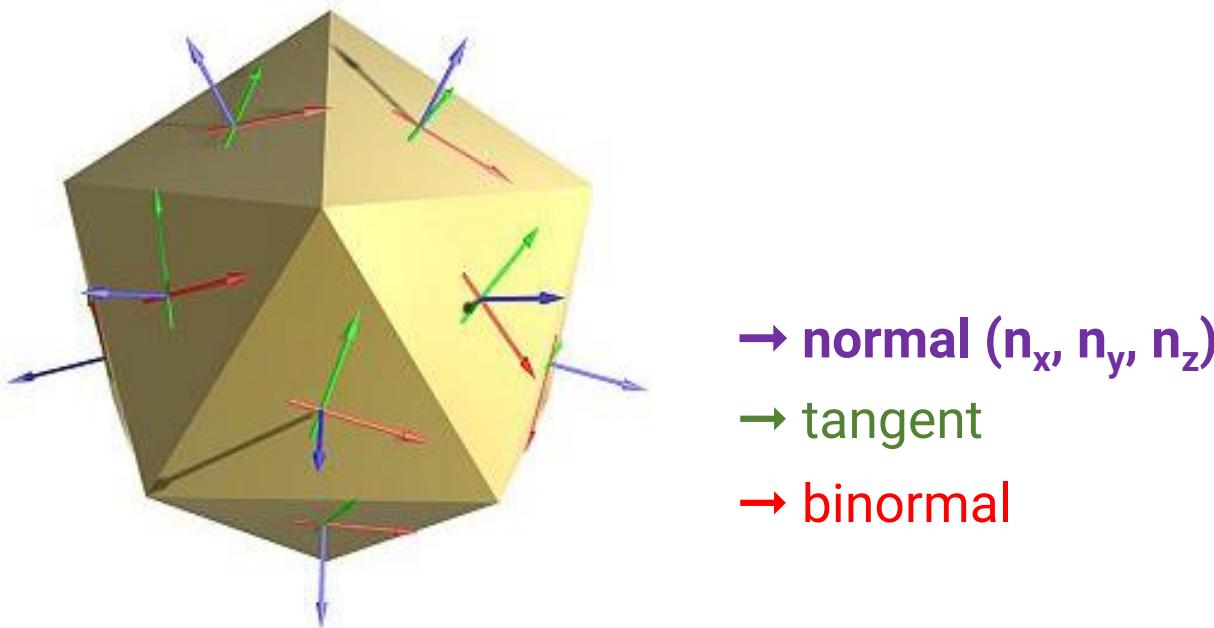
12 triangles



10K triangles

Triangle Mesh (cont.)

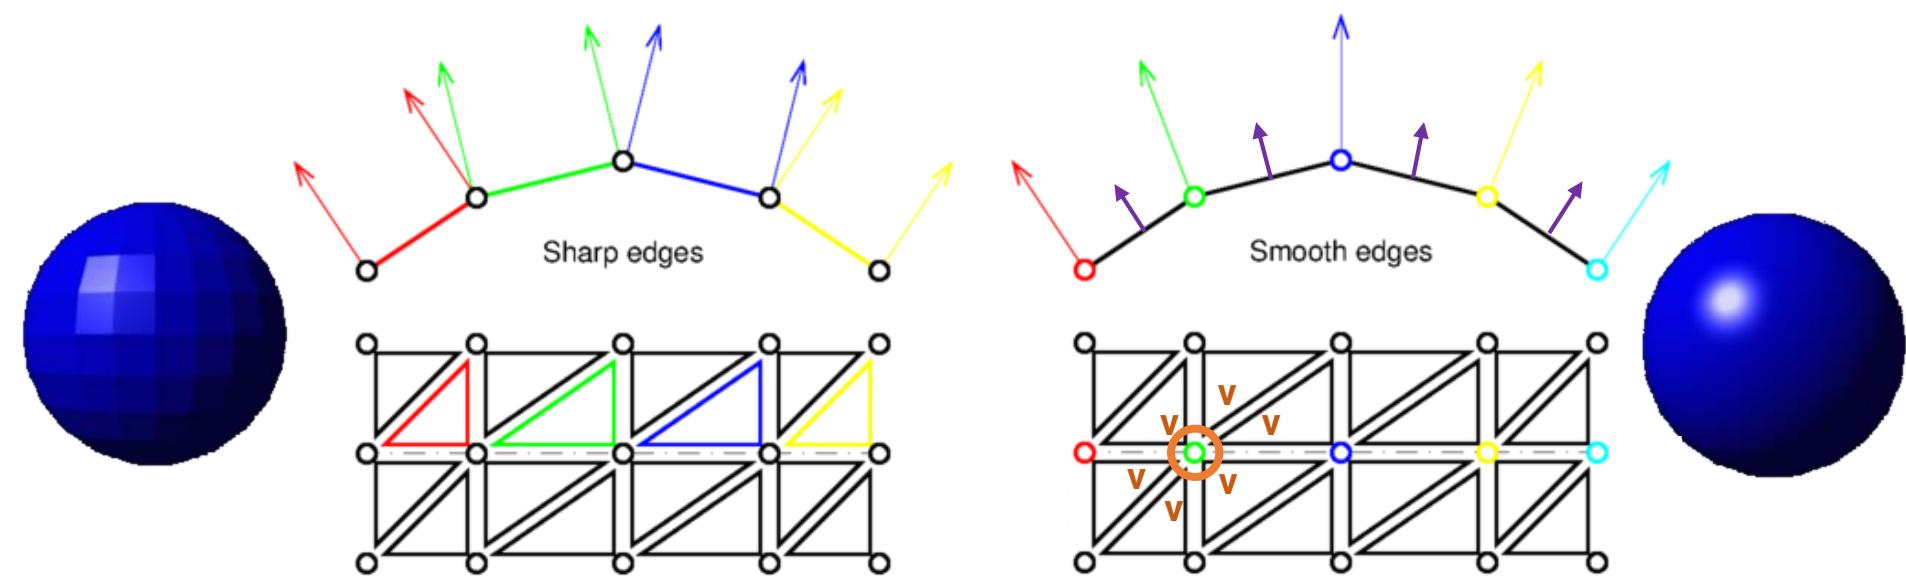
- A **surface normal** is a vector that is **perpendicular** to a surface at a particular position
- Represent the orientation of the face
- The length of a normal should be equal to **1**



Triangle Mesh (cont.)

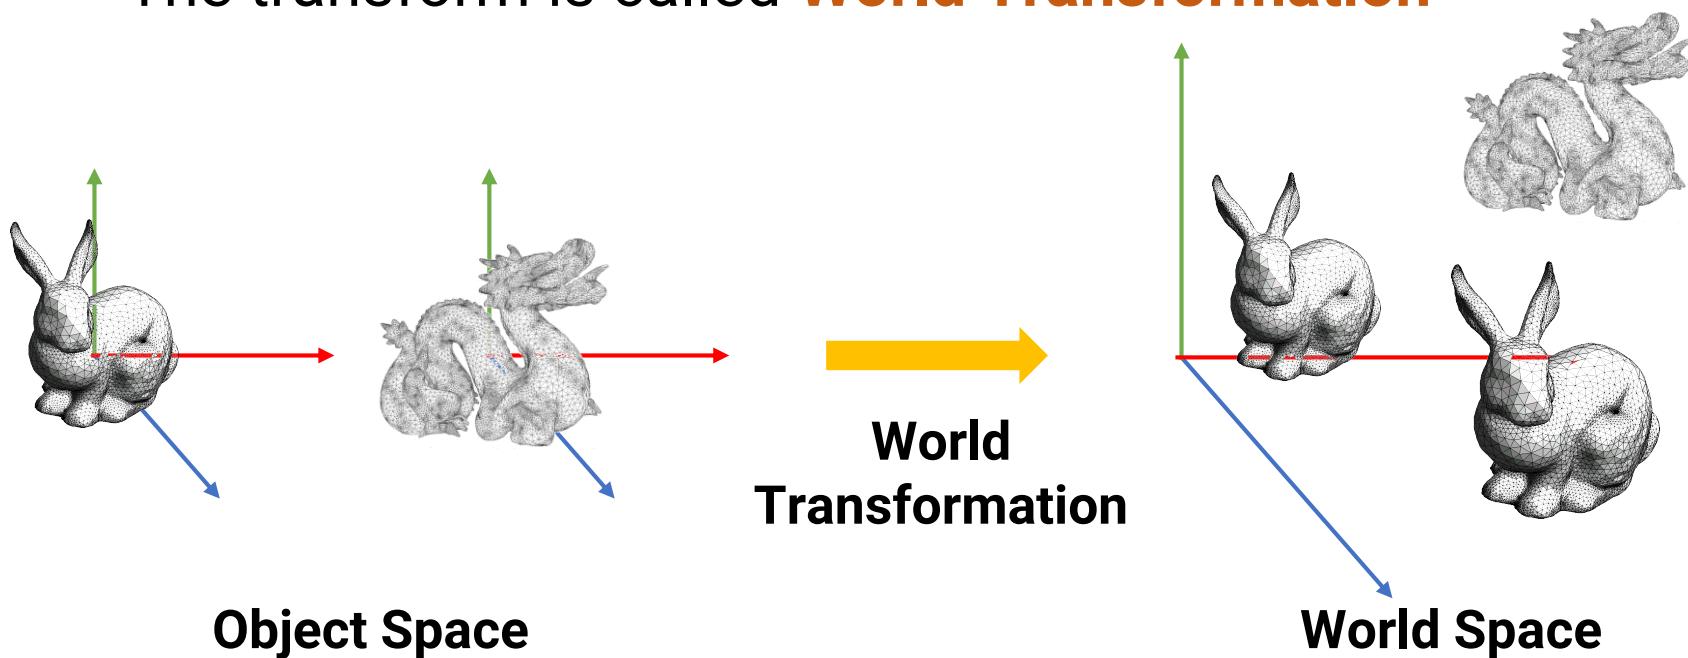
- **Vertex normal**

- Computed by **averaging** the surface normals of the faces that contain that vertex
- Can achieve more smooth shading



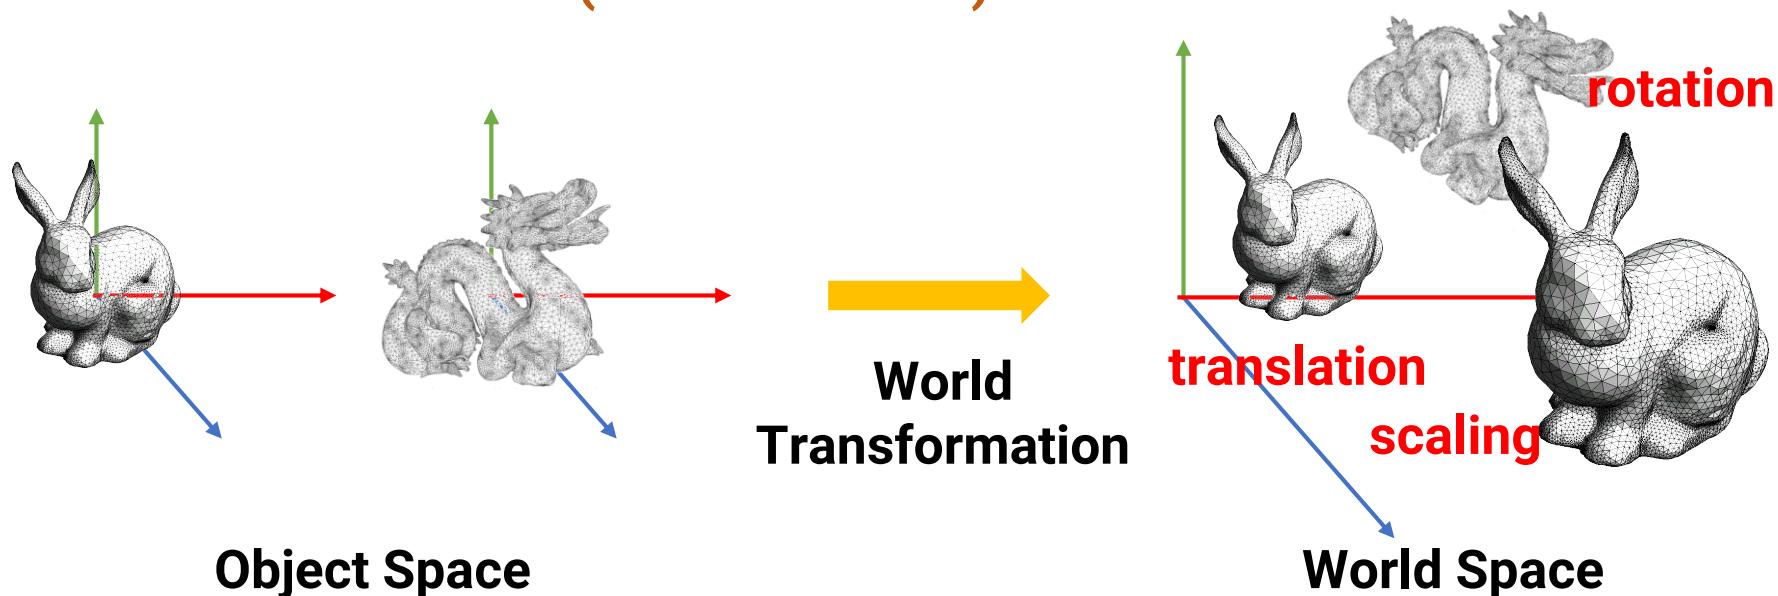
Object Space and World Space

- 3D models are usually defined in **Object Space**
- When building a scene, each object is transformed into a **global** and **unique** space called **World Space**
- The transform is called **World Transformation**



Object Space and World Space (cont.)

- Common 3D transforms for **World Transformation**
 - Translation
 - Scaling
 - Rotation
 - **Combination (concatenation) of transforms**



Common 3D Transforms

- To allow for concatenation, a transform is represented by a **4x4 matrix**

$$\begin{bmatrix} 1 & 0 & 0 & t_x \\ 0 & 1 & 0 & t_y \\ 0 & 0 & 1 & t_z \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

translation

$$\begin{bmatrix} s_x & 0 & 0 & 0 \\ 0 & s_y & 0 & 0 \\ 0 & 0 & s_z & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

scaling

$$\begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & \cos\theta & -\sin\theta & 0 \\ 0 & \sin\theta & \cos\theta & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

rotation w.r.t x-axis

$$\begin{bmatrix} \cos\theta & 0 & \sin\theta & 0 \\ 0 & 1 & 0 & 0 \\ -\sin\theta & 0 & \cos\theta & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

rotation w.r.t y-axis

$$\begin{bmatrix} \cos\theta & -\sin\theta & 0 & 0 \\ \sin\theta & \cos\theta & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

rotation w.r.t z-axis

Common 3D Transforms

- Concatenation of transforms

$$\begin{bmatrix} \cos\theta & 0 & \sin\theta & 0 \\ 0 & 1 & 0 & 0 \\ -\sin\theta & 0 & \cos\theta & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} s_x & 0 & 0 & 0 \\ 0 & s_y & 0 & 0 \\ 0 & 0 & s_z & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 & t_x \\ 0 & 1 & 0 & t_y \\ 0 & 0 & 1 & t_z \\ 0 & 0 & 0 & 1 \end{bmatrix}$$

rotation w.r.t y-axis

scaling

translation

Homogeneous Coordinate

- To apply a transform, we need to use **homogeneous coordinate**

- For a point

$$(x, y, z) \rightarrow (x, y, z, w = 1)$$

- For a vector (normal)

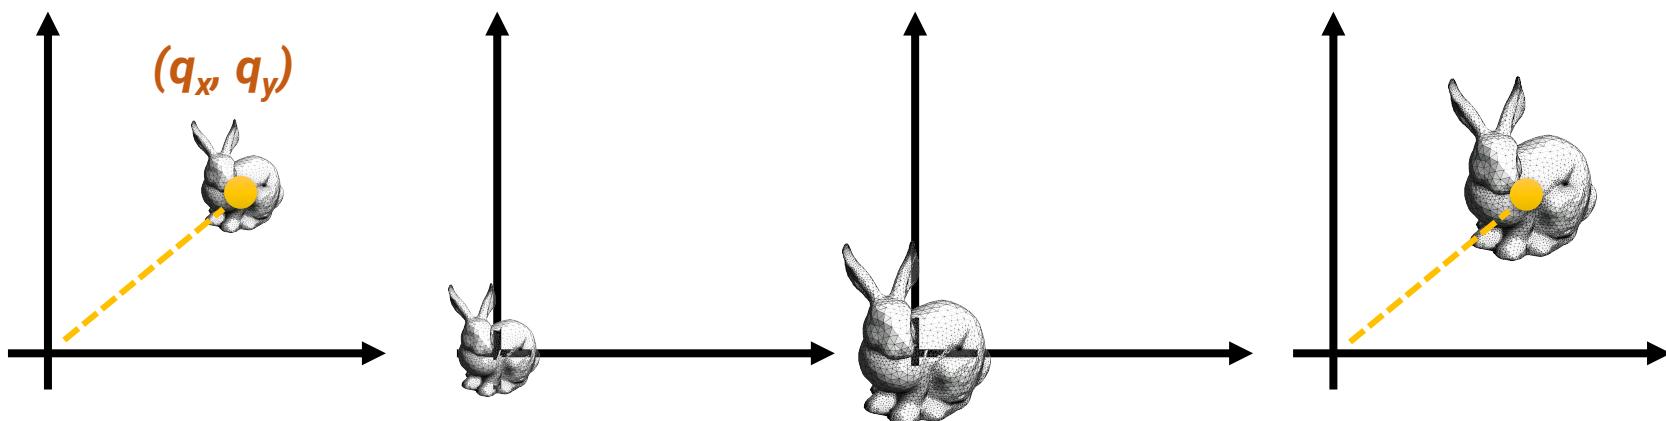
$$(x, y, z) \rightarrow (x, y, z, w = 0) \quad \text{to avoid translation}$$

- After transformation, the xyz components are **divided by the w component** for a transformed point, or **normalized** for a transformed vector (normal)

Issue of Scaling Transform

- If the object is not centered at the origin, it will get a shift
- To scale about an **arbitrary pivot point $Q(q_x, q_y, q_z)$**
 - Translate the objects so that Q will coincide with the origin: $T(-q_x, -q_y, -q_z)$
 - Scale the object: $S(s_x, s_y, s_z)$
 - Translate the object back: $T(q_x, q_y, q_z)$
 - The final scaling matrix can be written as $T(q)S(s)T(-q)$

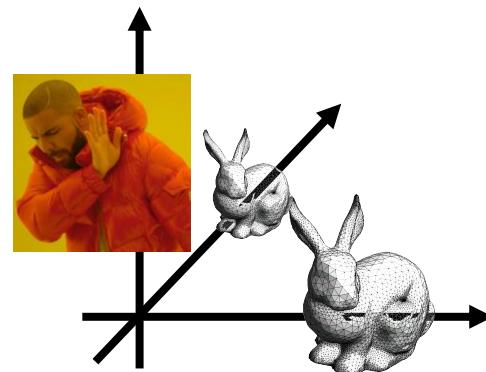
Concatenation
of matrices



Issues of Rotation Transform

- The standard rotation matrix is used to rotate w.r.t the x, y, and z axis
- **To rotate locally**
 - Translate the objects so that its center Q will coincide with the origin: $T(-q_x, -q_y, -q_z)$
 - Rotate the object: $R(\theta)$
 - Translate the object back: $T(q_x, q_y, q_z)$
 - The final rotation matrix can be written as $T(q)R(\theta)T(-q)$

Concatenation
of matrices



Why use World Transformation?

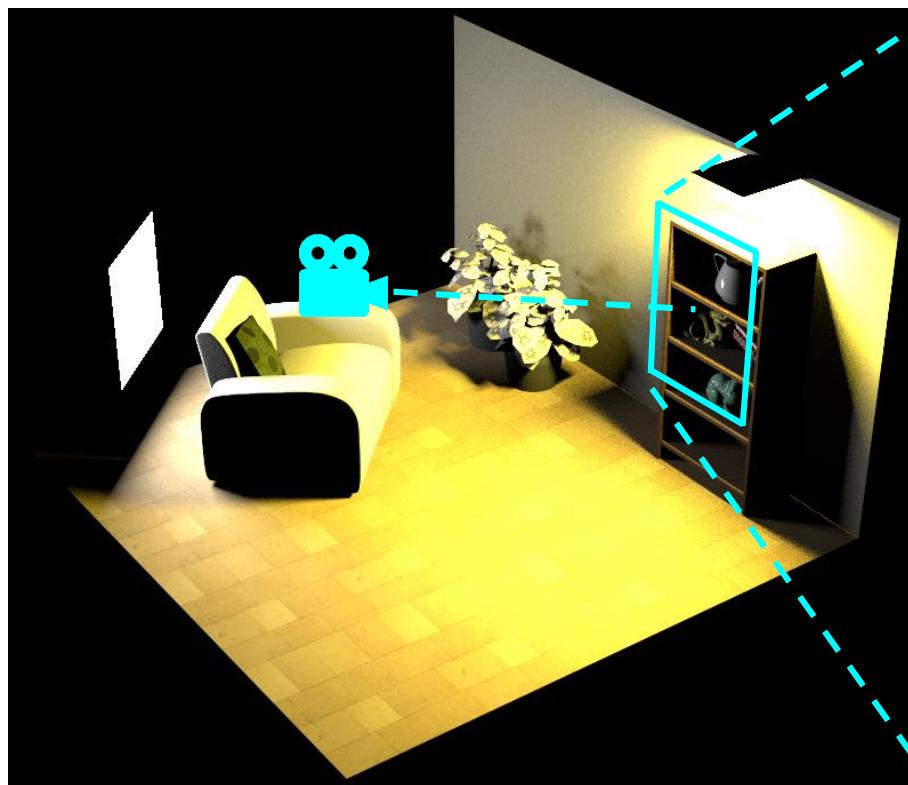
- **Reuse model:** design a model and use it in several scenes
- **Memory saving:** store a 4x4 matrix instead of duplication of the entire models



Rendering

Recap: Digital Image Synthesis

- Generate an image from the 3D world with a **virtual camera**



3D virtual world



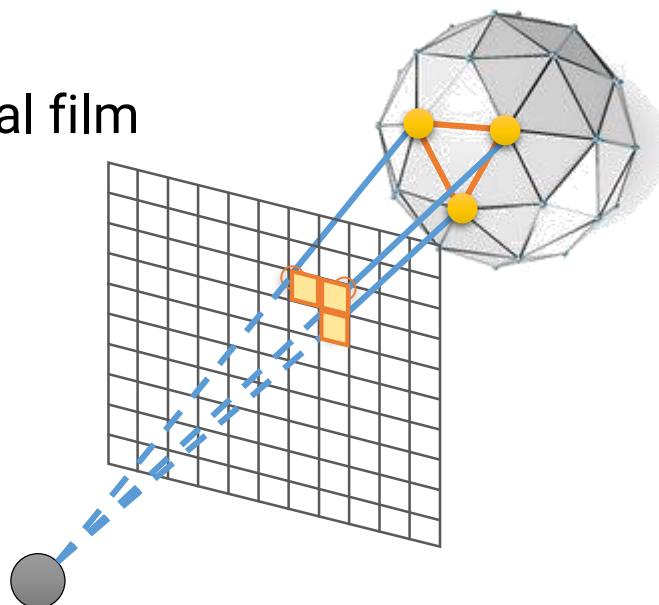
rendered image

Rasterization and Ray Tracing

- Two ways for generating synthetic images

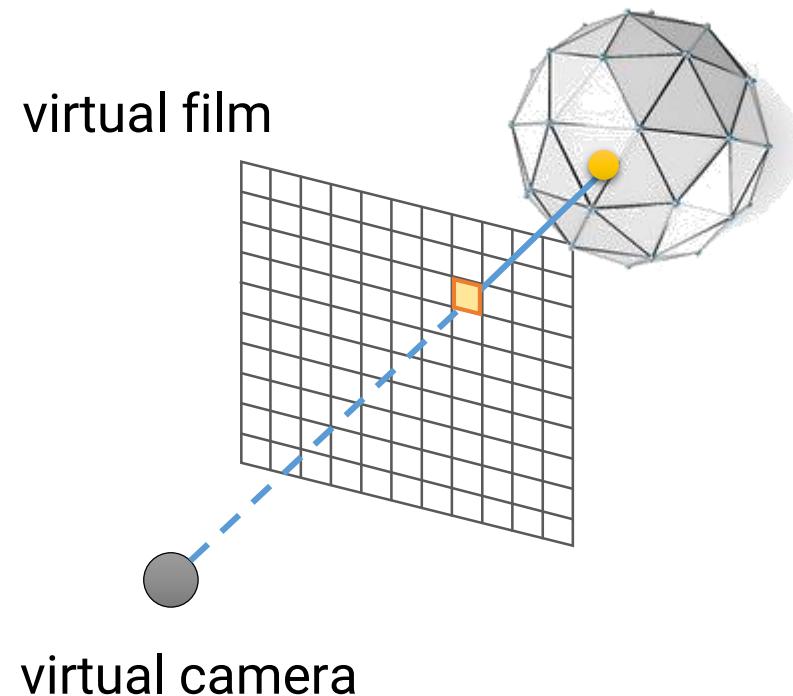
Rasterization

virtual film



Ray tracing

virtual film



virtual camera

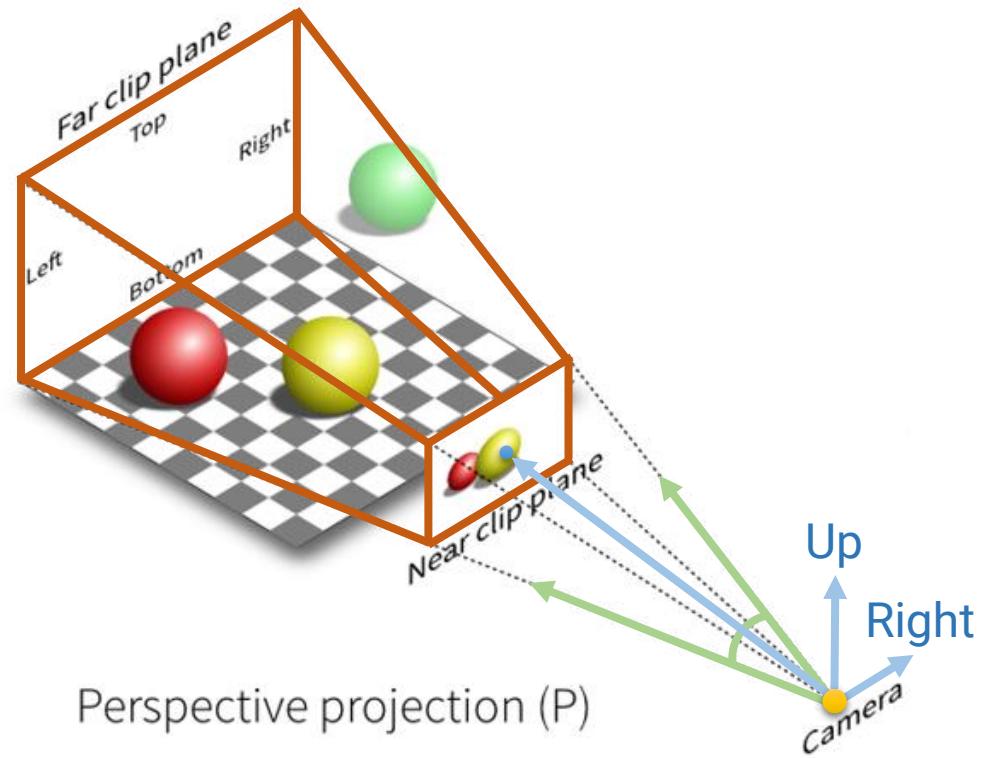
Camera Properties

- The film is **in front of** the camera (to avoid up-side-down)
- **Basic properties**

- Camera position
- Viewing direction
- Camera local frame
- Field of view
- Aspect ratio

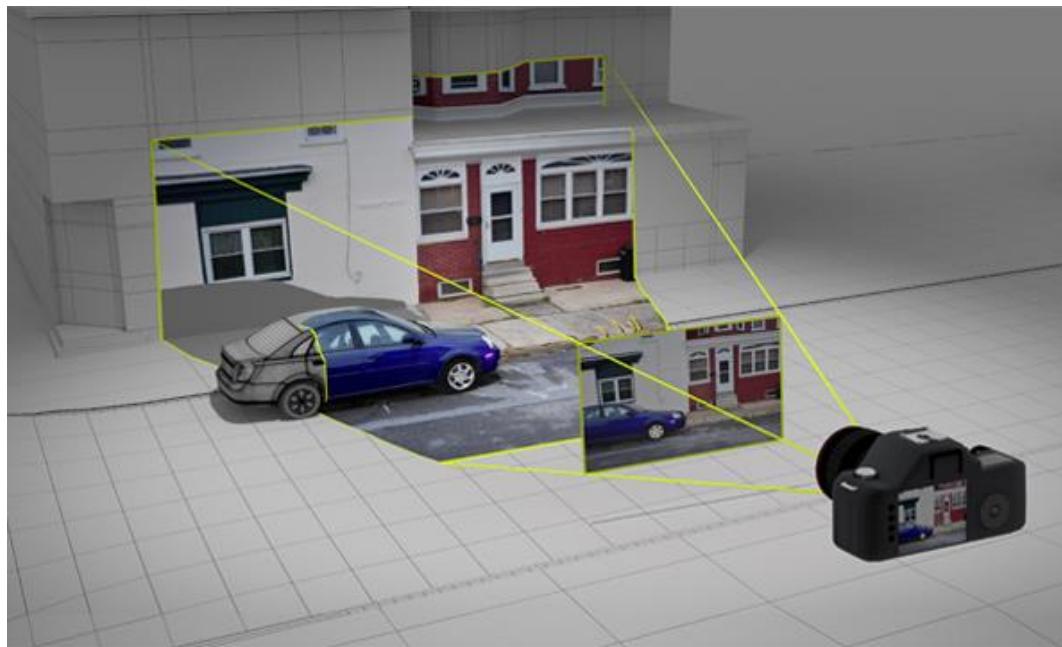
viewing volume (view frustum)

- **Advanced properties**
- Shutter speed
- Lens system



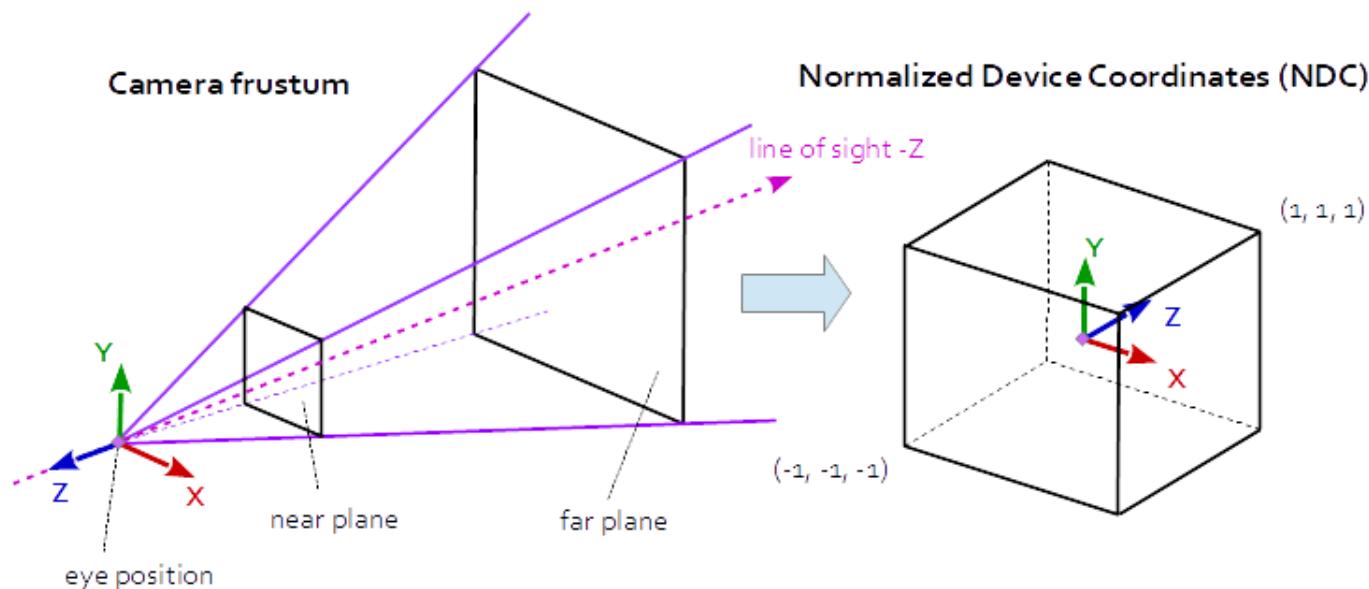
Camera (View) Space

- The camera can be at an arbitrary position and have an arbitrary viewing direction in **World Space**
- This makes the **projection** or **ray generation** difficult in terms of mathematics



Camera (View) Space (cont.)

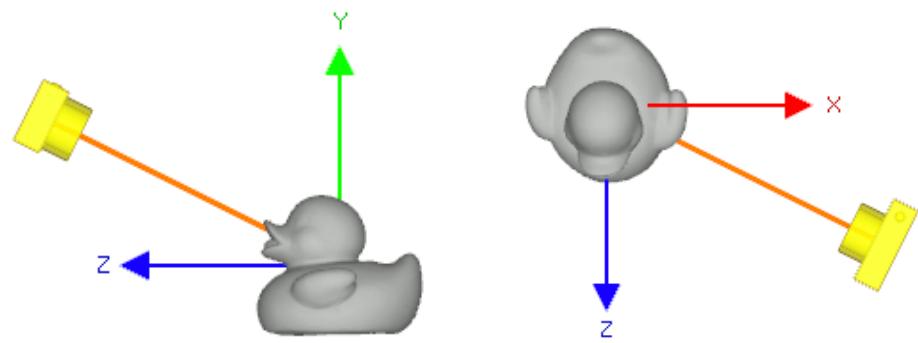
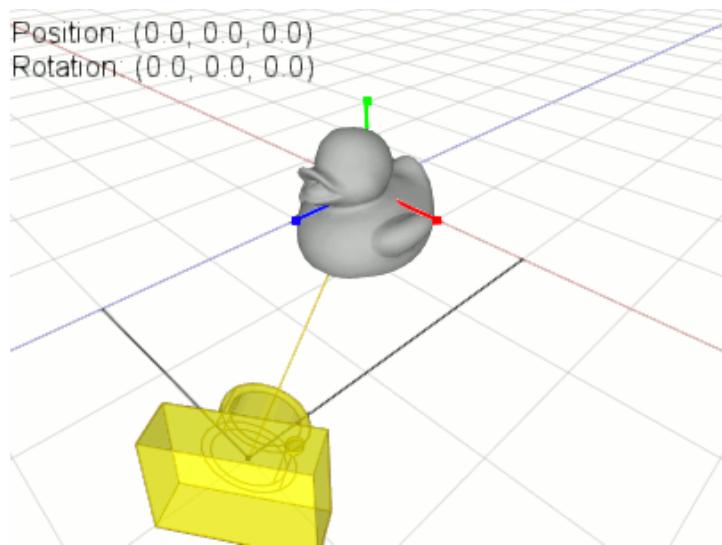
- To keep the math of projection or ray generation simpler, we additionally define a **camera (view, eye) space**
 - In the camera space, the camera is **at the origin $(0, 0, 0)$** and **looking at the negative Z-axis**



- Vertices are then projected **from Camera Space to NDC**

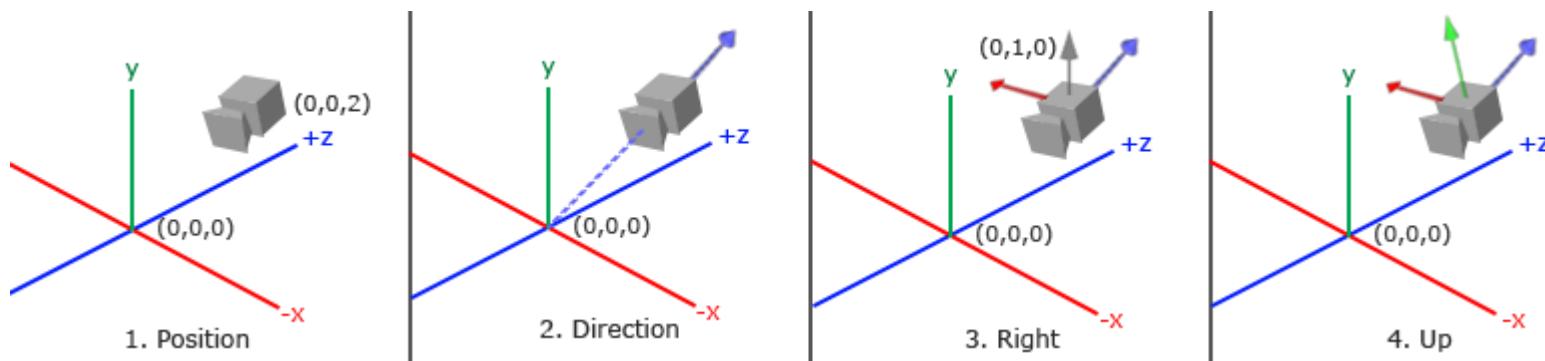
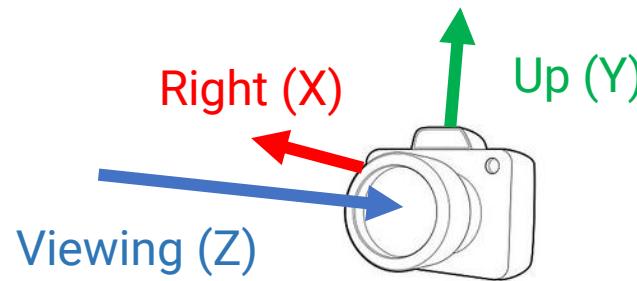
Camera (View) Transformation

- An object (with all its vertices) is transformed from **World Space** to **Camera Space** by applying **Camera Transformation (4x4 matrix)**
- The camera matrix is a combination of **translation** and **rotation** matrix



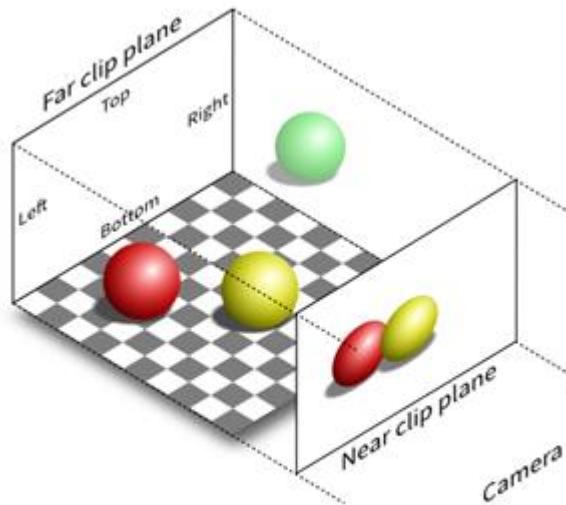
Camera (View) Transformation

- To do this, we need to define the camera's local frame
 - Viewing direction (Z)
 - Right vector (X)
 - Up vector (Y)
- Process
 - Move it with the inverse translation of the camera's position
 - Rotate the object to match the camera's local frame

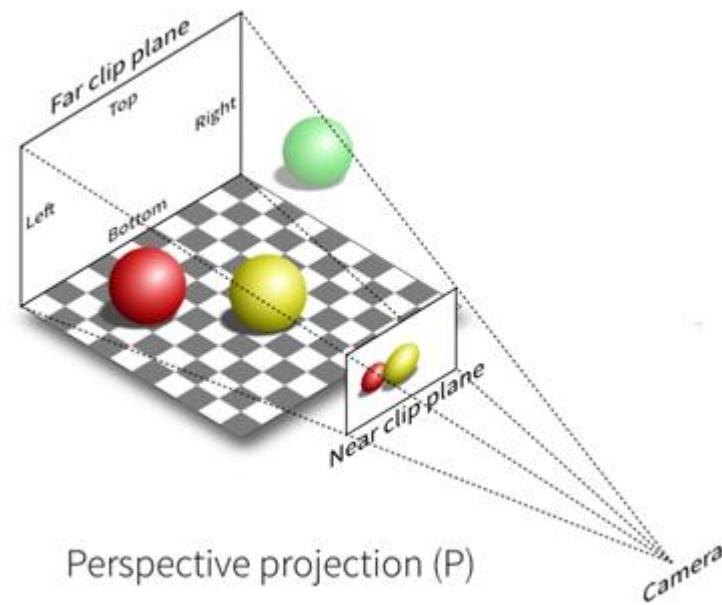


Camera Projection

- Two common projections in Rasterization
 - **Orthographic**: parallel projection with projectors perpendicular to the projection plane
 - **Perspective**: mimic the vision of human eyes; the objects that are farther away appear much smaller



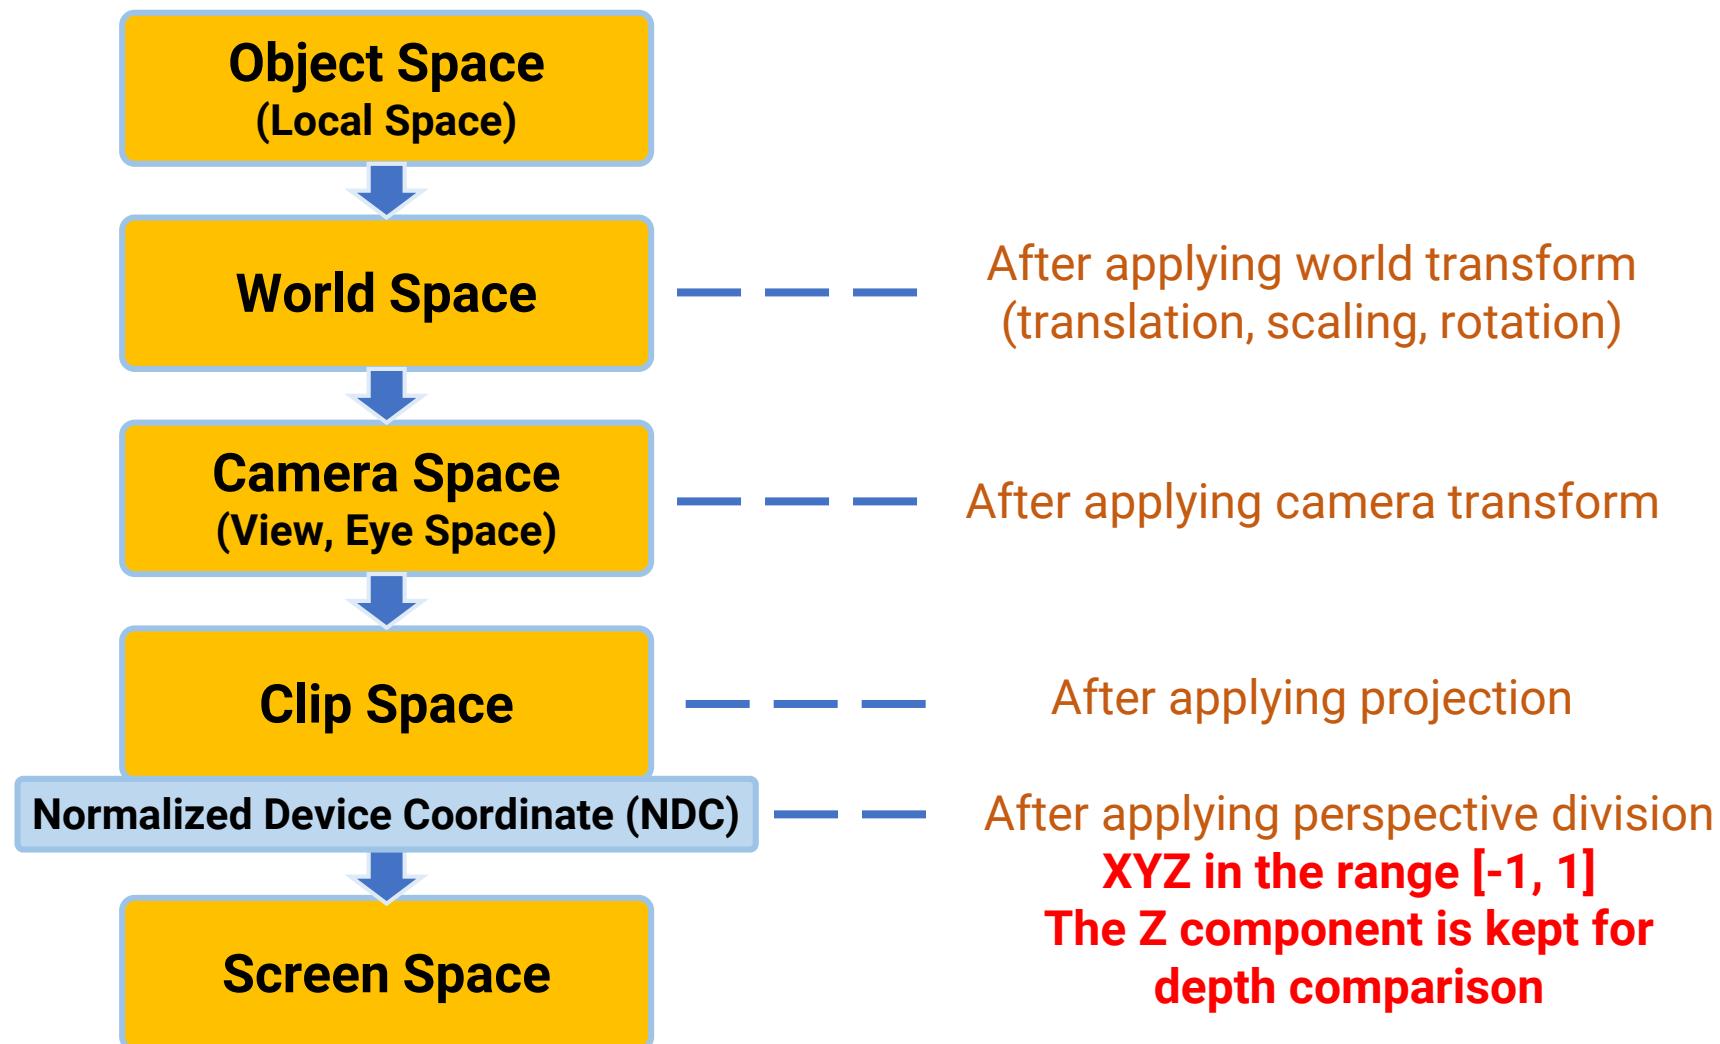
Orthographic projection (O)



Perspective projection (P)

GPU Rendering Pipeline

Full Vertex Transformation



Running on GPU! (Rasterization)

- In complex scenes there are massive vertices, each requires
 - Vertex transformation
 - Shading
- Both tasks are relatively simple
- Better tailored to graphics hardware (GPU)



GPU Rendering Pipeline

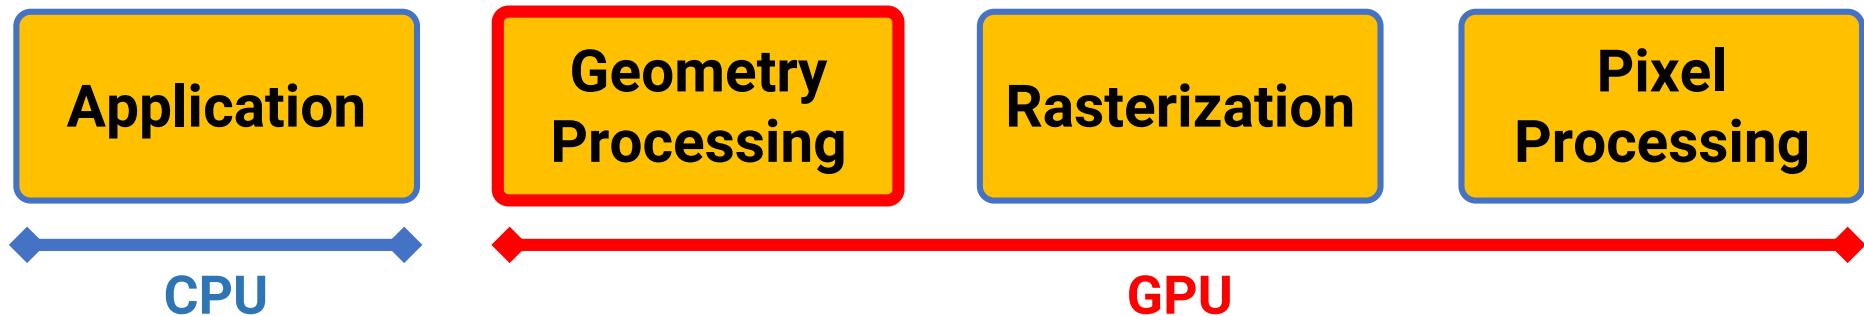
- Responsible for the fixed routines of bringing triangles to pixels
- Can be roughly categorized into 4 stages



- Physical simulation
- Animation
- Collision detection
- Global acceleration
- etc.

GPU Graphics Pipeline Overview (cont.)

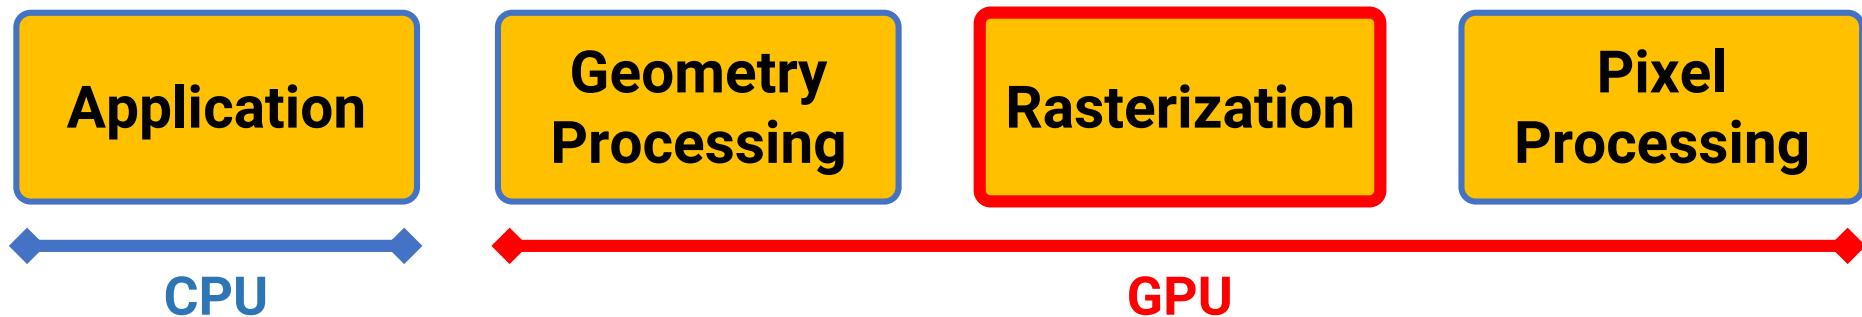
- Responsible for the fixed routines of bringing triangles to pixels
- Can be roughly categorized into 4 stages



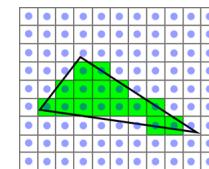
- Vertex transform and projection
- Vertex lighting and shading (rarely used now)
- Geometry assembly
- Clipping
- Culling

GPU Graphics Pipeline Overview (cont.)

- Responsible for the fixed routines of bringing triangles to pixels
- Can be roughly categorized into 4 stages



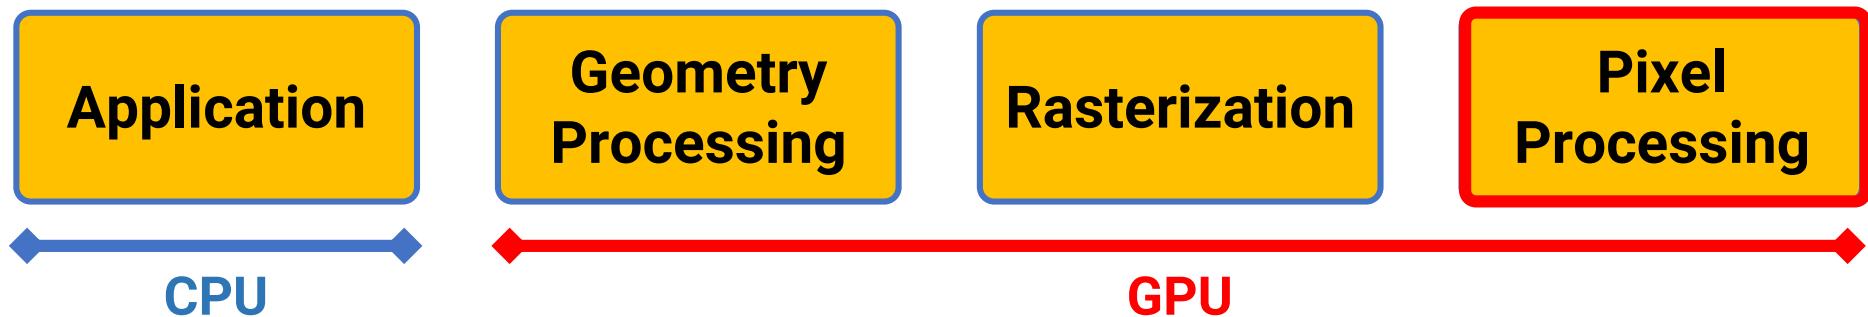
- Triangle setup
- Fragments (pixels) generation
 - Attribute interpolation



**from continuous
to discrete**

GPU Graphics Pipeline Overview (cont.)

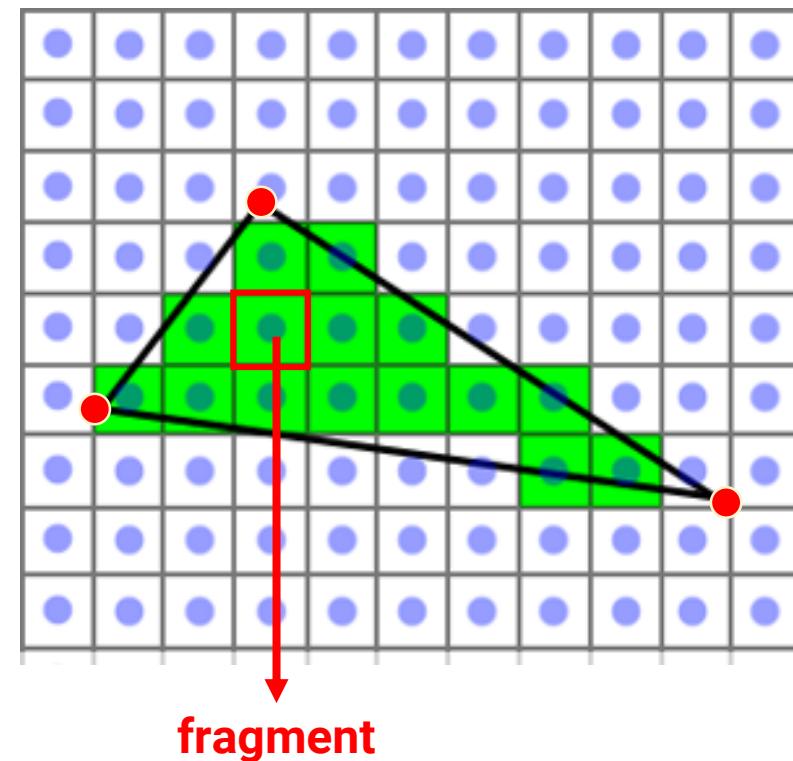
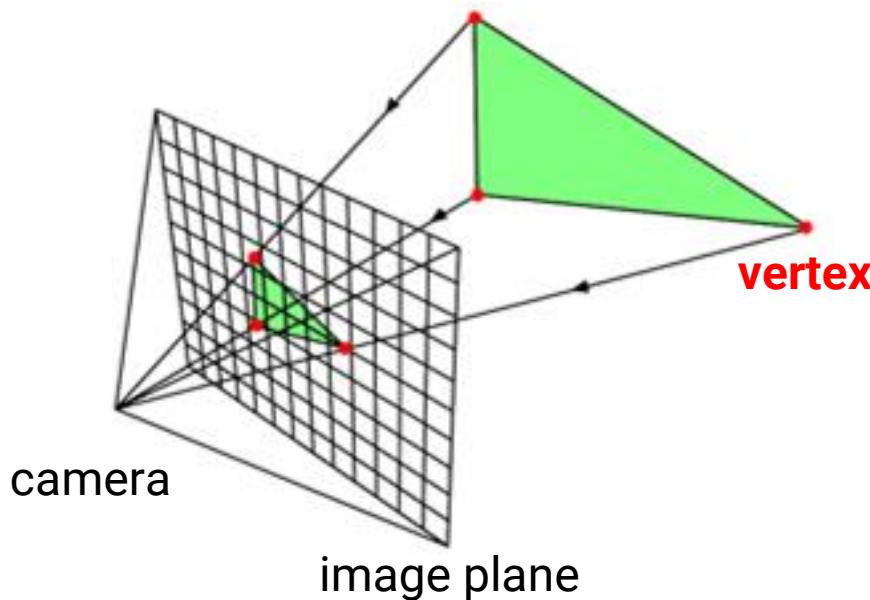
- Responsible for the fixed routines of bringing triangles to pixels
- Can be roughly categorized into 4 stages



- Pixel shading / Texturing
- Depth testing
- Alpha blending

Rasterization

- Convert **triangles (continuous)** into **fragments (discrete)**, which eventually become the individual **screen pixels**)
- Vertex attributes are **interpolated** across the face

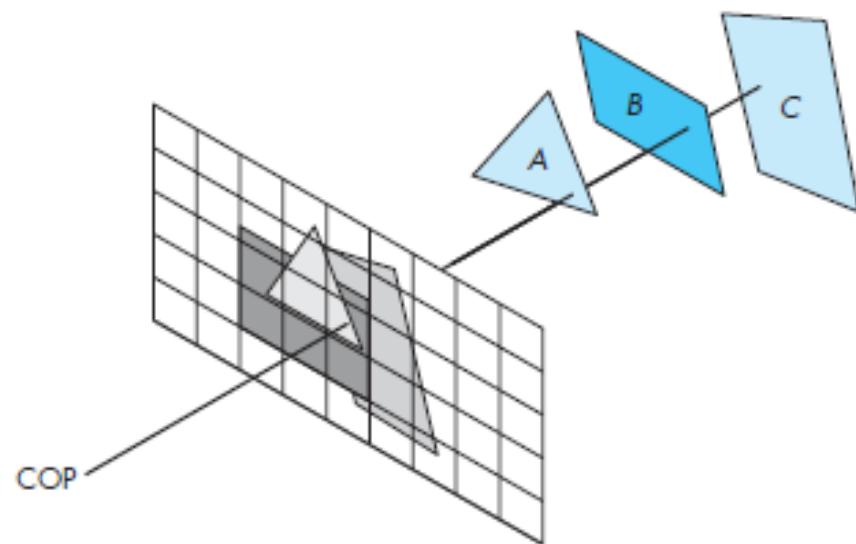


Rasterization (cont.)

- Convert **triangles (continuous)** into **fragments (discrete)**, which eventually become the individual **screen pixels**)
- Vertex attributes are **interpolated** across the face, including
 - (Lighting) color if per-vertex lighting is used
 - Texture coordinate
 - Position for per-fragment lighting
 - Normal (after OpenGL 2.0)
 - Anything you want to interpolate

Depth Comparison

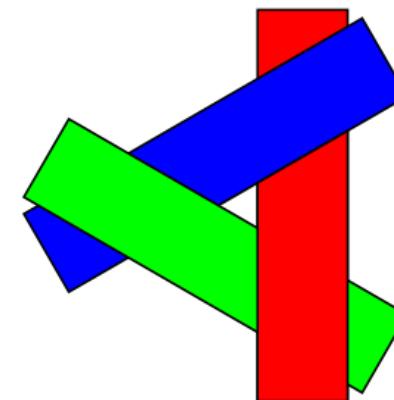
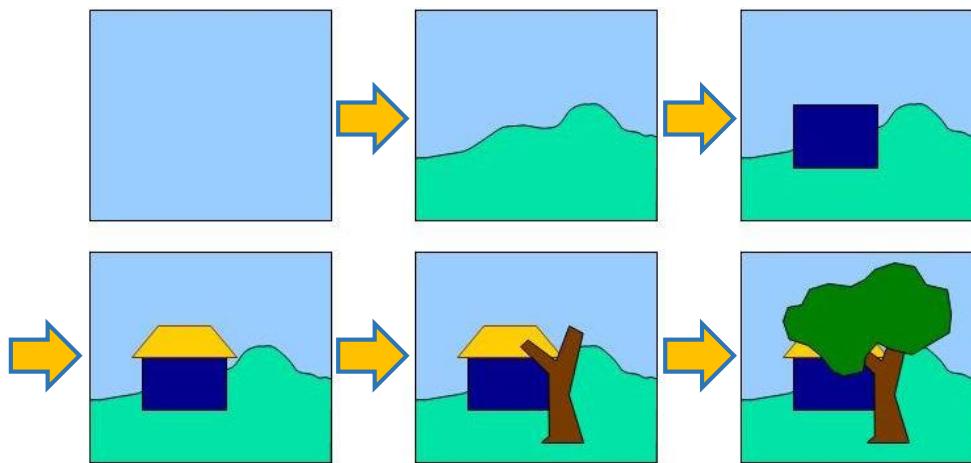
- Several surfaces can project to the same screen location
 - Only the **closest** surfaces from the camera can be seen at each pixel
- How to hide occluded surfaces?
 - Painter's algorithm
 - Z-buffer



Depth Comparison (cont.)

- **Painter's algorithm**

- Draw objects in an order based on their distances from the camera
- The farthest object first

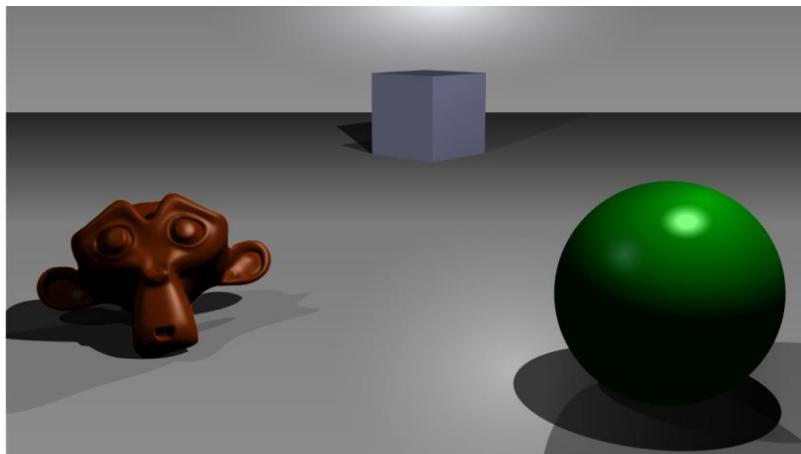


How about this?
(cyclical overlapping)

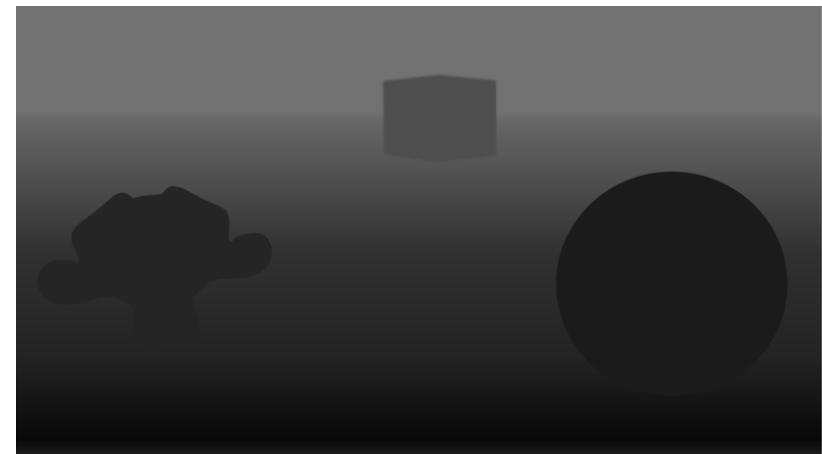
Depth Comparison (cont.)

- **Z-buffer (current solution)**

- An additional buffer used to maintain the z value of the closest surface to a pixel
- **Discard** fragments if they have larger depth values than the ones stored in their corresponding positions in the Z buffer



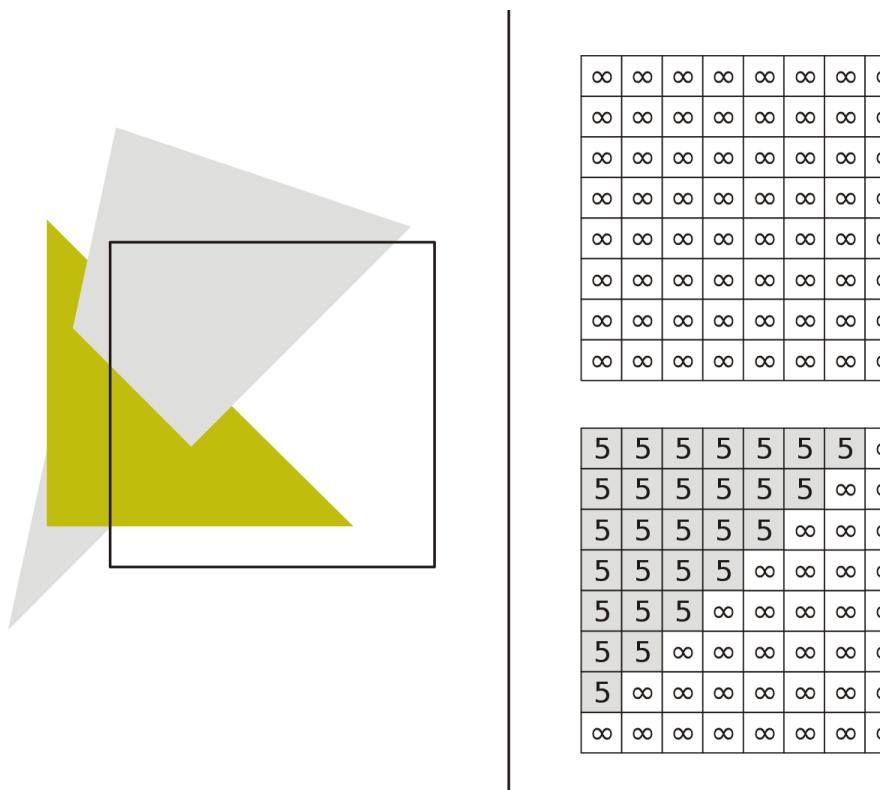
color frame buffer



Z (depth) buffer

Depth Comparison (cont.)

- Z-buffer (current solution)
 - An example of Z-buffer update

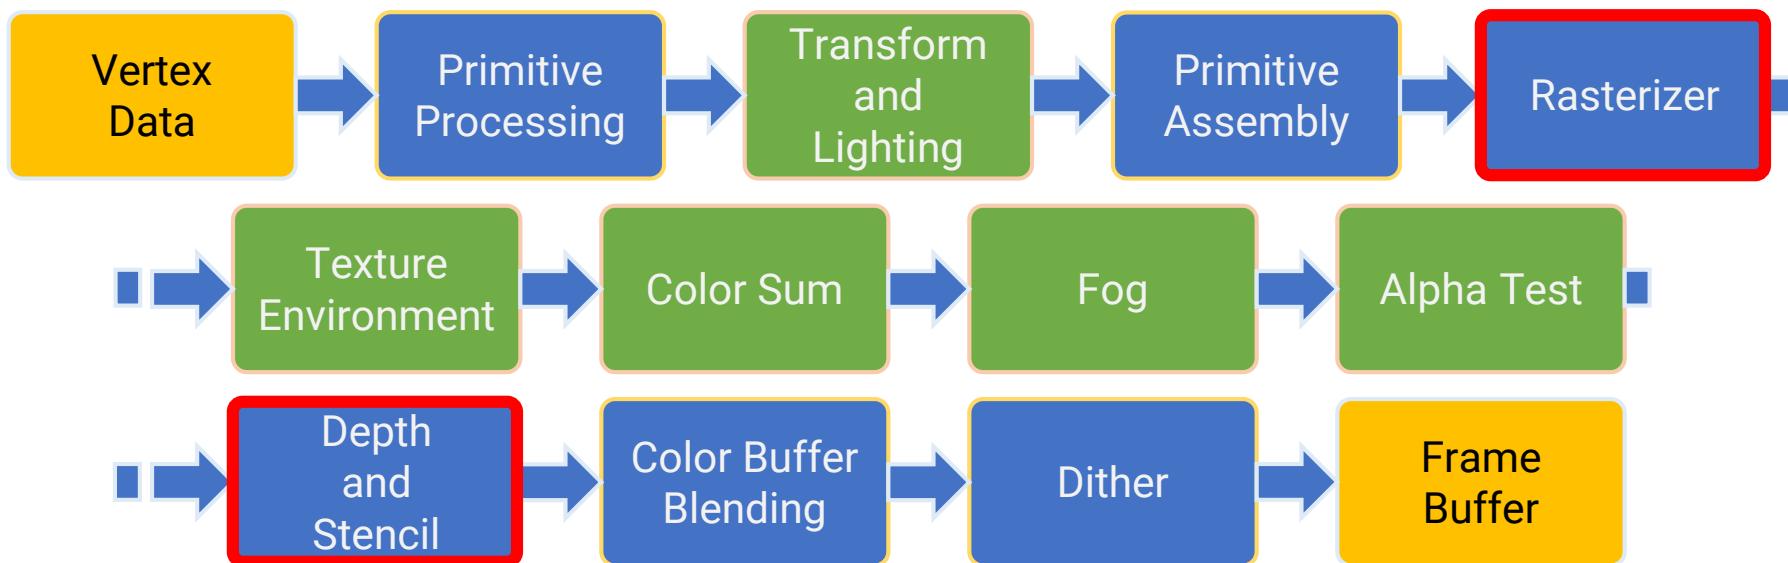


$$\begin{array}{c}
 \begin{array}{|c|c|c|c|c|c|c|} \hline
 5 & 5 & 5 & 5 & 5 & 5 & 5 \\ \hline
 5 & 5 & 5 & 5 & 5 & 5 & \\ \hline
 5 & 5 & 5 & 5 & 5 & & \\ \hline
 5 & 5 & 5 & 5 & & & \\ \hline
 5 & 5 & 5 & & & & \\ \hline
 5 & 5 & & & & & \\ \hline
 5 & & & & & & \\ \hline
\end{array} \\
+
\end{array}$$

$$+ \begin{array}{r} 7 \\ 6 \ 7 \\ 5 \ 6 \ 7 \\ 4 \ 5 \ 6 \ 7 \\ 3 \ 4 \ 5 \ 6 \ 7 \\ 2 \ 3 \ 4 \ 5 \ 6 \ 7 \end{array}$$

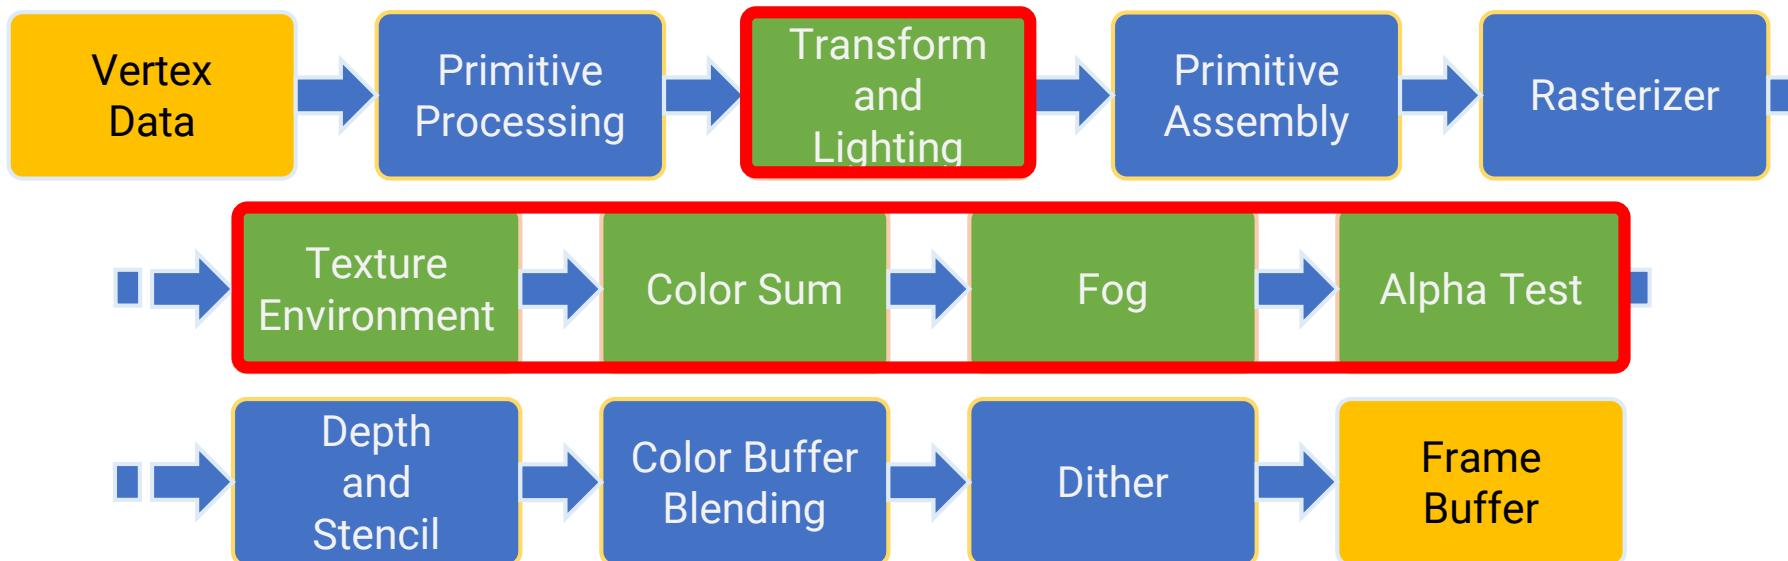
Depth Comparison (cont.)

- Used when OpenGL was first introduced
- All the functions performed by OpenGL are **fixed** and **could not** be modified except through the manipulation of the **rendering states**



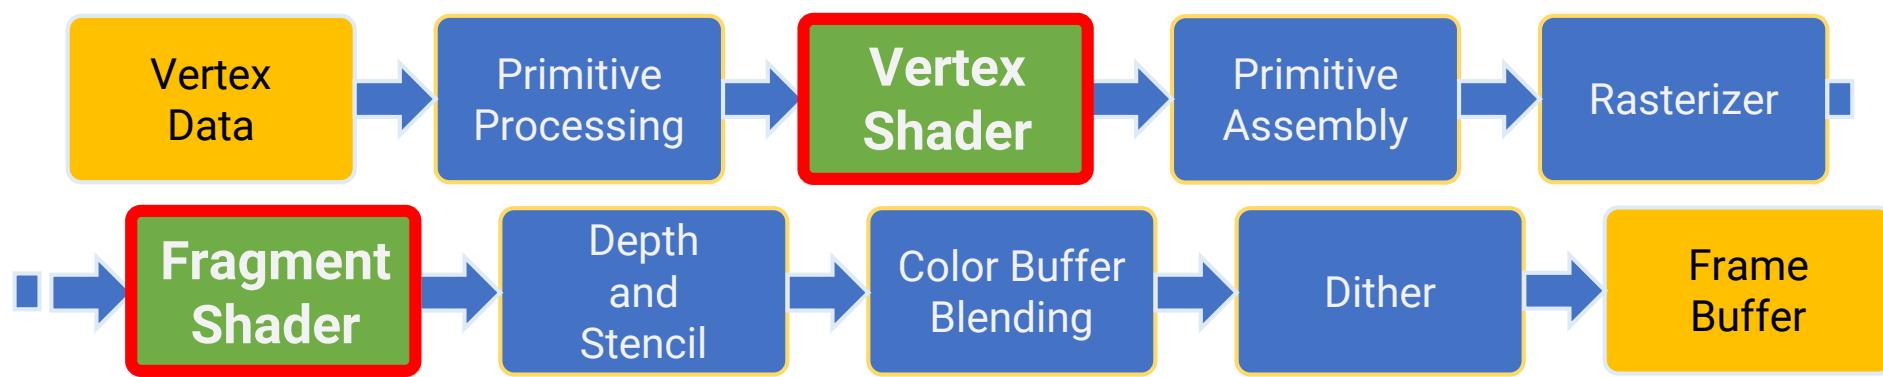
OpenGL (1.1) Fixed Function Pipeline

- All the functions performed by OpenGL are **fixed** and **could not** be modified except through the manipulation of the **rendering states**
- To provide flexibility, the stages shown in **green** have been replaced by **shaders**



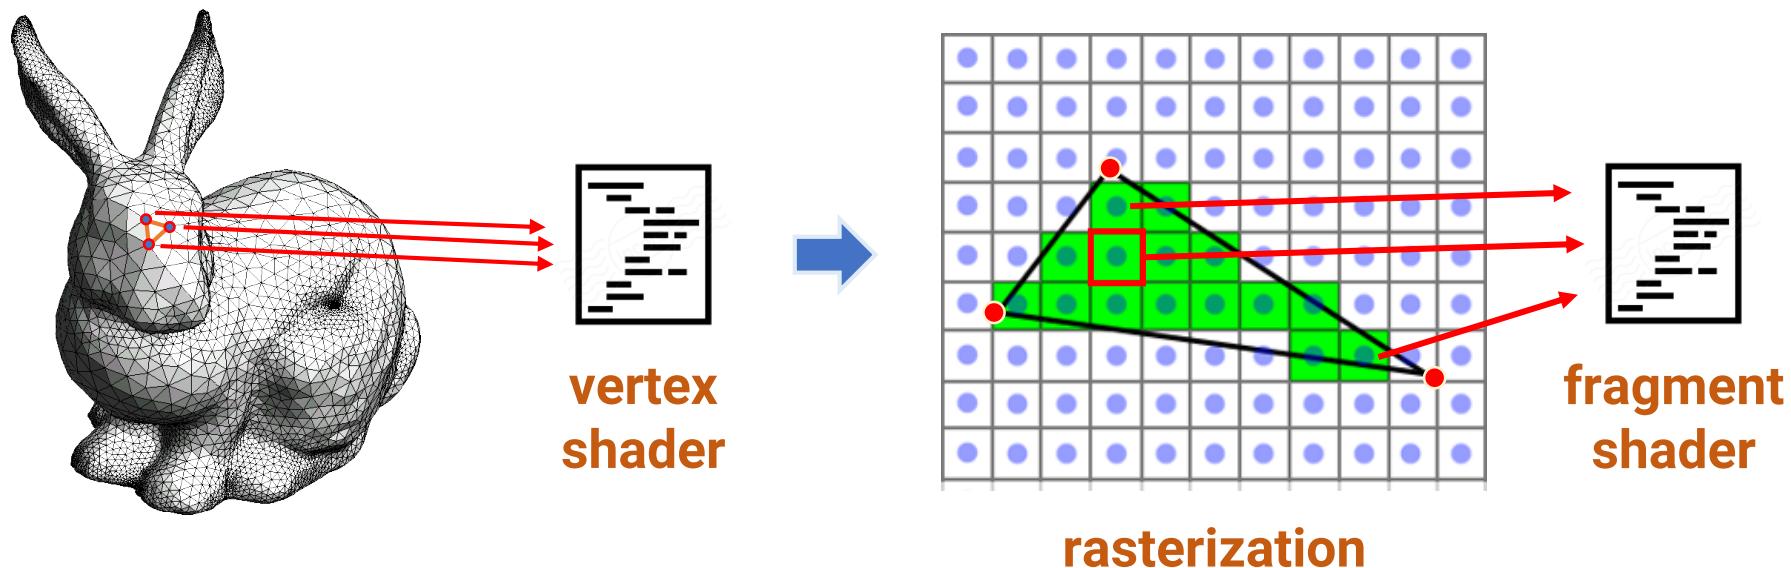
OpenGL (2.0) Graphics Pipeline

- Released in 2004
- Provide the ability to **programmatically** define the vertex transformation and lighting and the fragment operations (with small GPU programs called **shaders**)



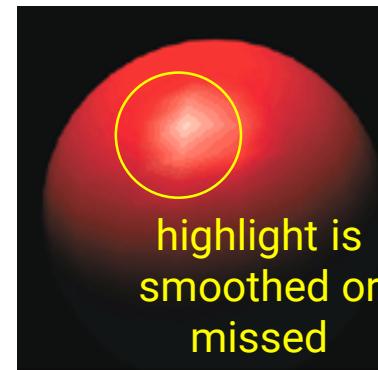
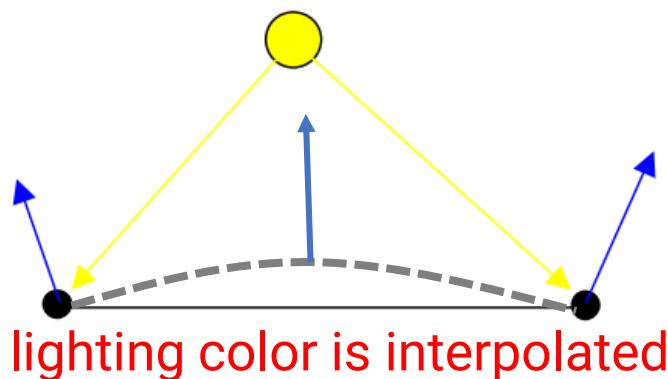
Vertex Shader and Fragment Shader

- **Important concepts**
 - The vertex shader runs **per vertex**
 - **Goal: at least transform a vertex to Clip Space**
 - The fragment shader runs **per (rasterized) fragment**
 - **Goal: calculate the color of the fragment**

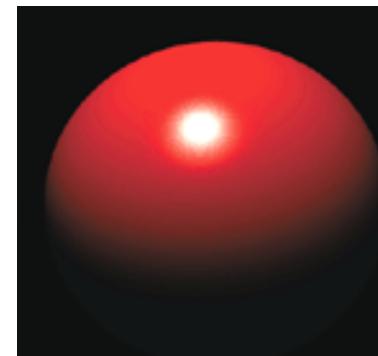
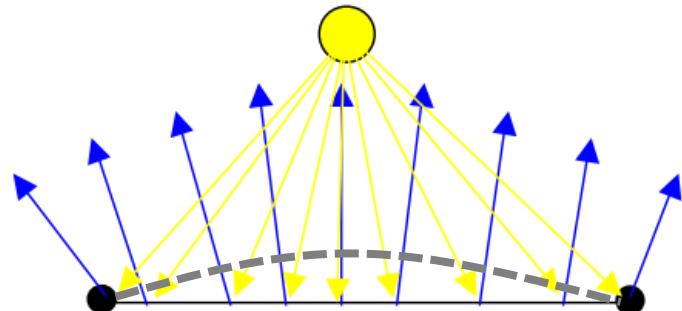


Per-Fragment Shading

- **Gouraud shading** (interpolated vertex lighting)



- **Phong shading**



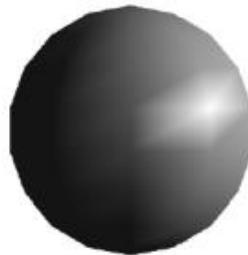
surface normal is interpolated (how? Rasterization!)

Per-Fragment Shading (cont.)

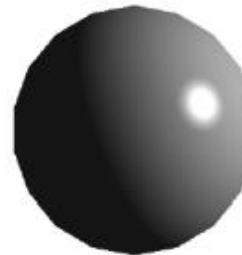
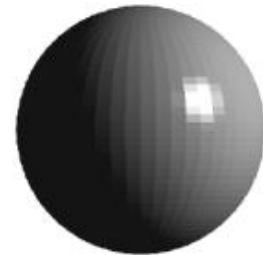
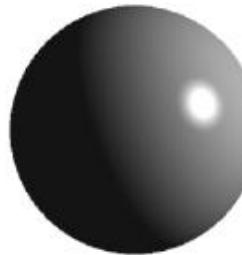
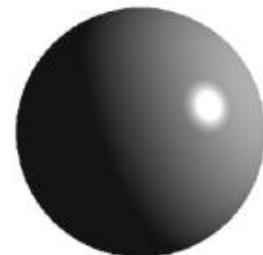
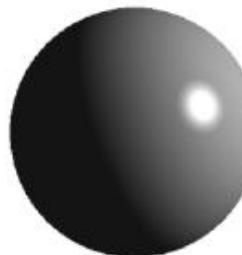
flat shading

(a₁)

Gouraud shading

(b₁)

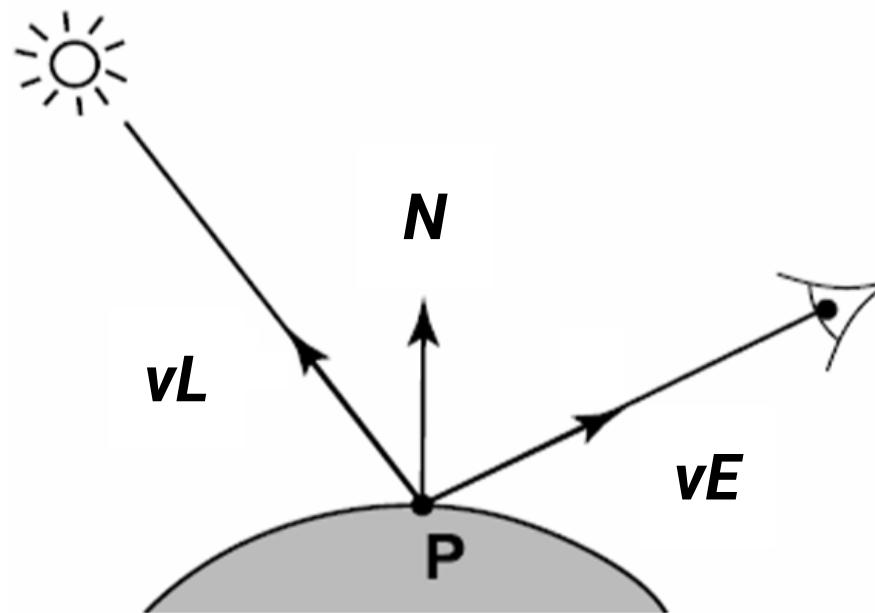
Phong shading

(c₁)(a₂)(b₂)(c₂)(a₃)(b₃)(c₃)

Lighting and Shading

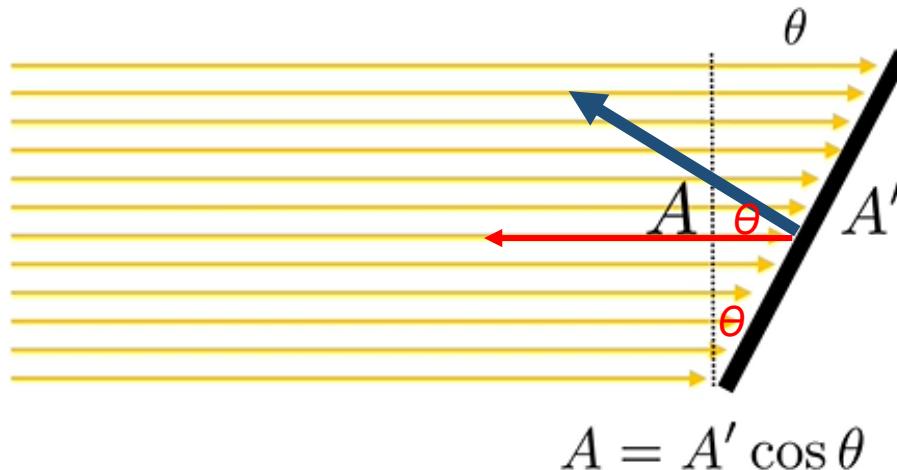
Shading

- The color of a surface is determined by
 - The **light data**
 - The **material property** of the surface
 - The **geometric setting** between the light and the surface



Lambertian Cosine Law

- Illumination on an oblique surface is less than on a normal one
- Generally, illumination falls off as $\cos\theta$



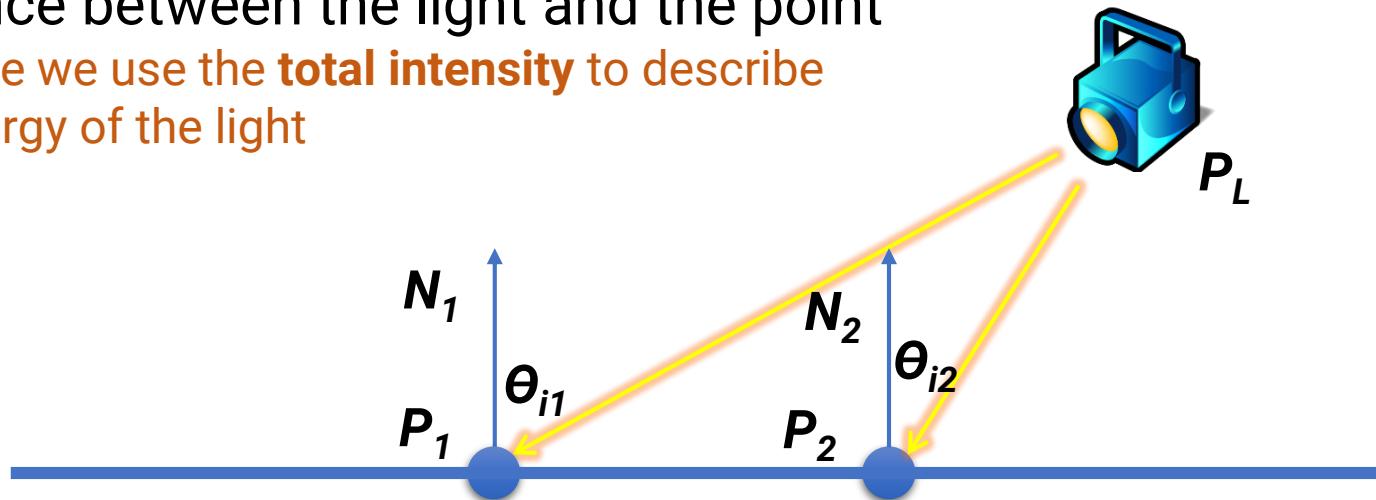
$$E = \frac{\Phi}{A'} = \frac{\Phi \cos \theta}{A}$$

Lights in Computer Graphics

- Point light
 - Spot light
 - Area light
- local lights
-
- Directional light
 - Environment light
- distant lights

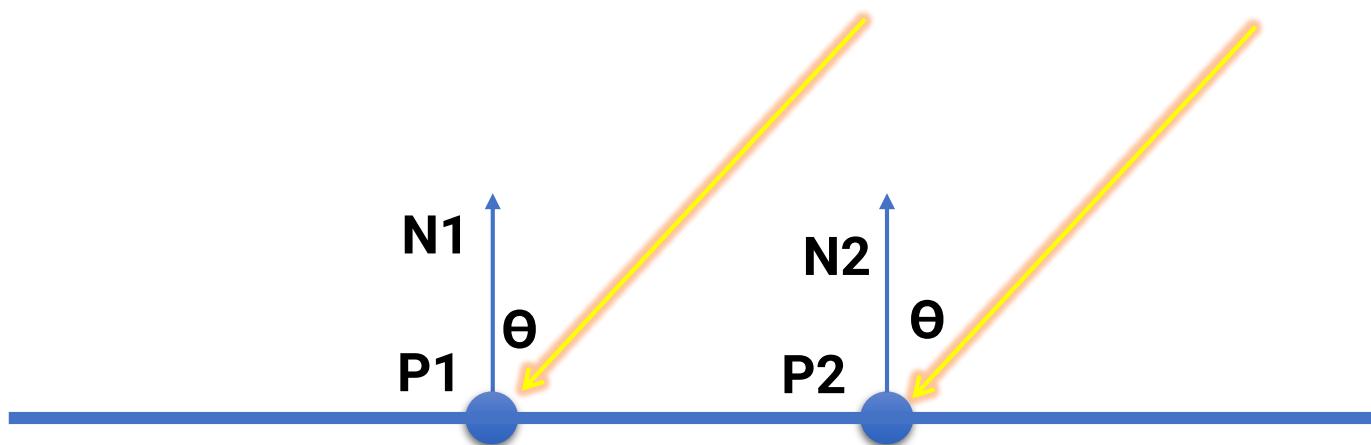
Local Light

- The distance between a light and a surface is **not** long enough compared to the scene scale
- The position of light needs to be considered during shading
 - **Lighting direction** $v_L = |P_L - P|$
 - **Lighting attenuation** is proportional to the square of the distance between the light and the point because we use the **total intensity** to describe the energy of the light



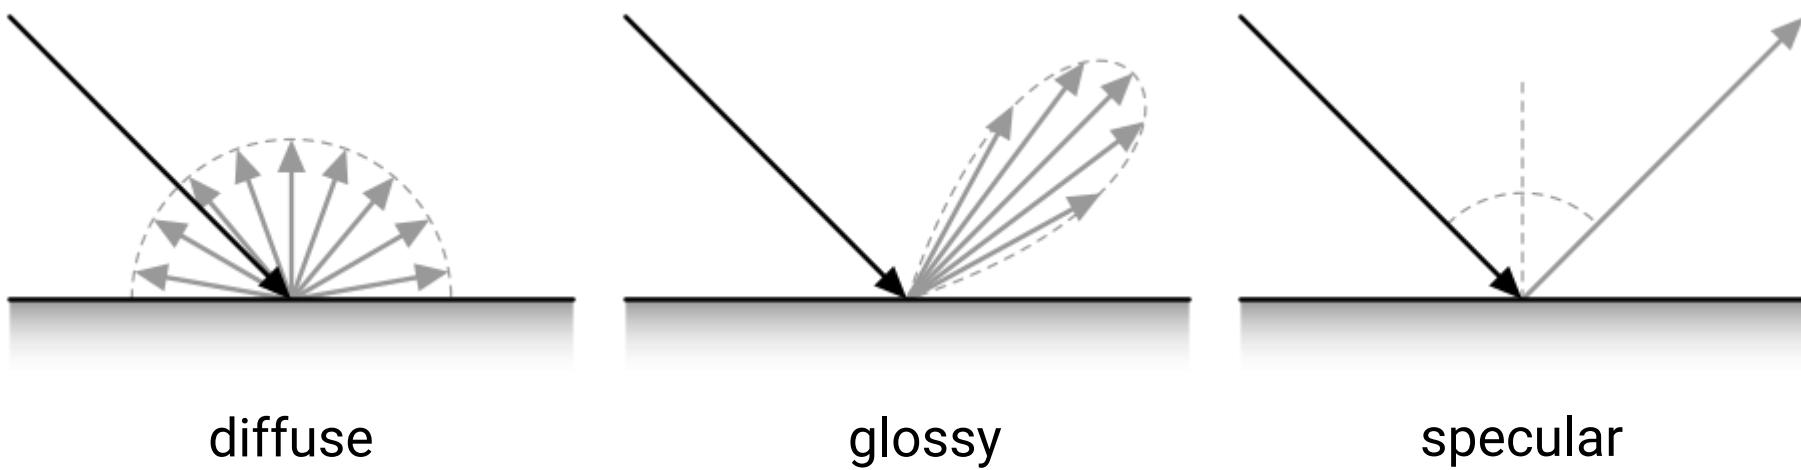
Distant Light

- The distance between a light and a surface is long enough compared to the scene scale and **can be ignored**
 - **Lighting direction is fixed**
 - **No lighting attenuation**
because we use **directional radiance** to describe the energy of the light
- **Directional light (sun)** is the most common distant light



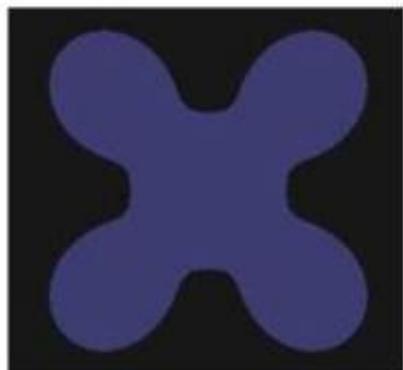
Materials

- Highly related to surface types
- The **smoother** a surface, the more reflected light is concentrated in the direction a **perfect mirror** would reflect the light

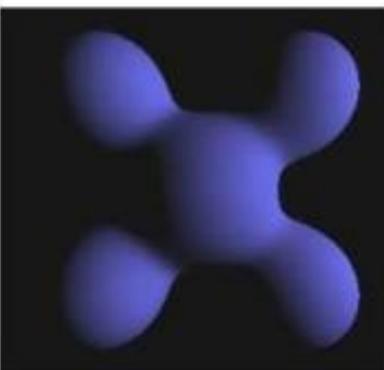


Phong Lighting Model

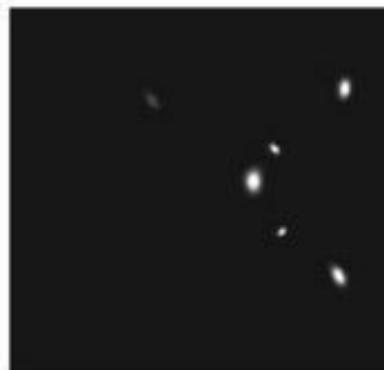
- **Diffuse reflection**
 - Light goes everywhere; colored by object color
- **Specular reflection**
 - Happens only near mirror configuration; usually white
- **Ambient reflection**
 - Constant accounted for global illumination (cheap hack)



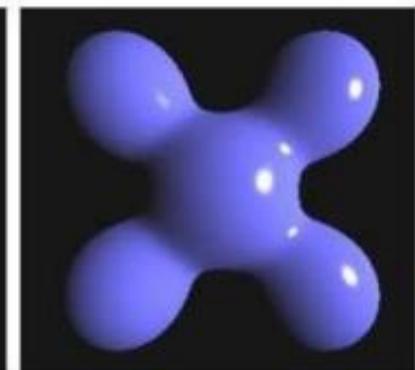
ambient



diffuse

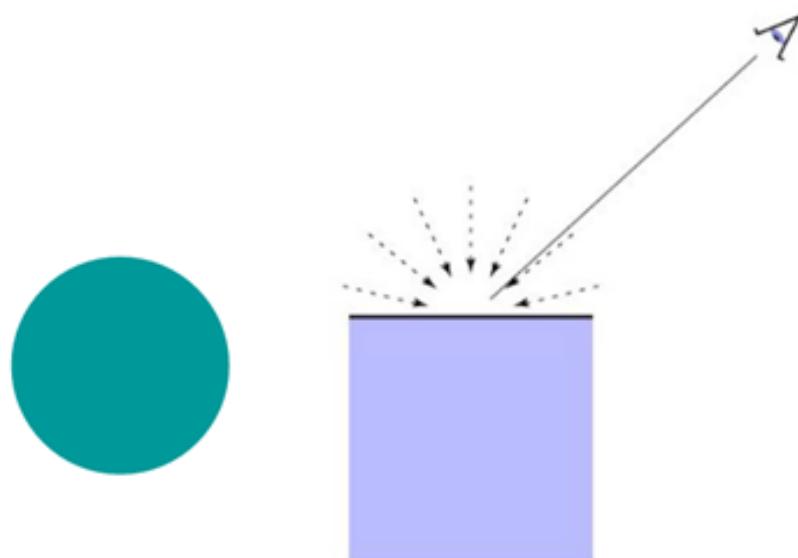


specular



Ambient Shading

- Add constant color to account for disregarded illumination and fill black shadows



$$L_a = k_a \cdot I_a$$

the **intensity** of ambient light

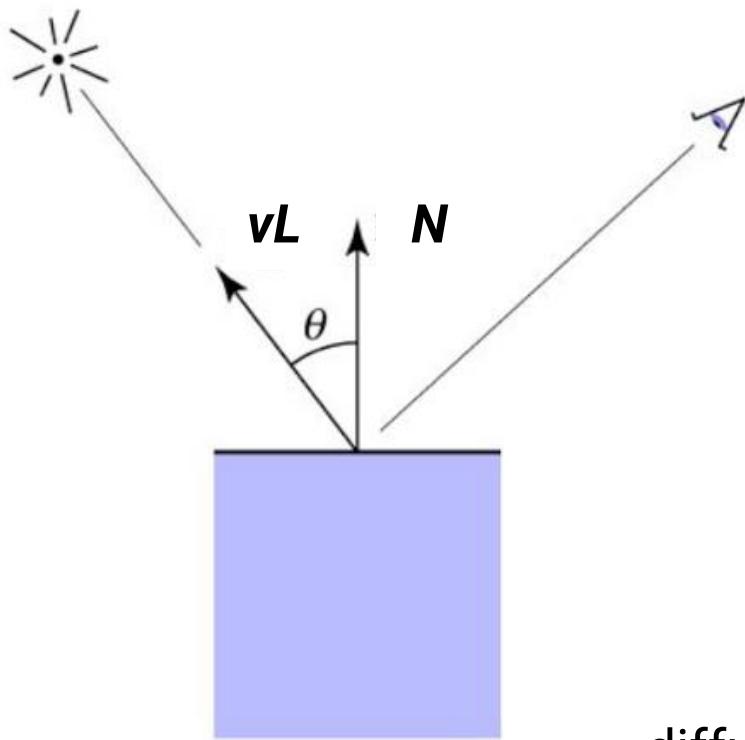
ambient coefficient

reflected ambient light

The equation $L_a = k_a \cdot I_a$ is displayed. To the right of the equation, a blue arrow points downwards from the text "the intensity of ambient light". Another blue arrow points upwards from the text "ambient coefficient". A third blue arrow points upwards from the text "reflected ambient light".

Diffuse Shading

- Applies to diffuse or matte surface



illumination from source

$$L_d = k_d \cdot I \cdot \max(0, N \cdot vL)$$

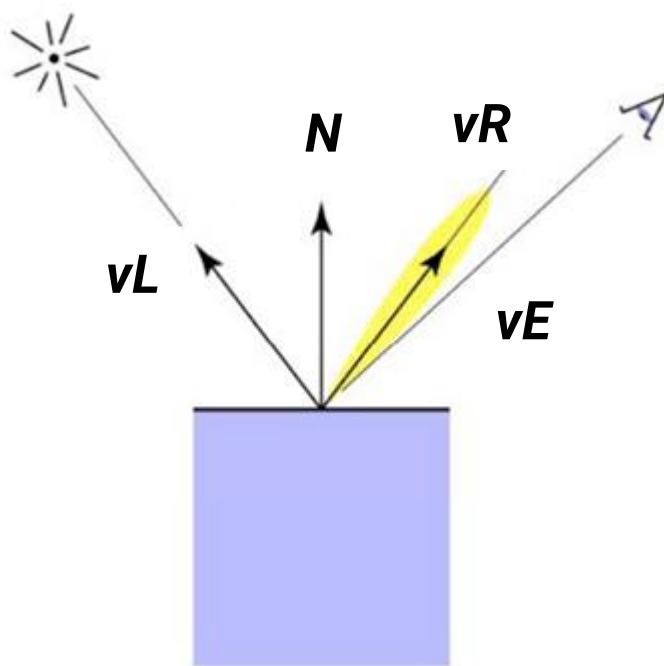
Lambertian law

diffuse coefficient

diffusely reflected light

Specular Shading

- Also known as glossy
- **Phong specular model [1975]**
 - Fall off gradually from the perfect reflection direction



$$\begin{aligned}
 vR &= vL + 2((N \cdot vL)N - vL) \\
 &= 2(N \cdot vL)N - vL
 \end{aligned}$$

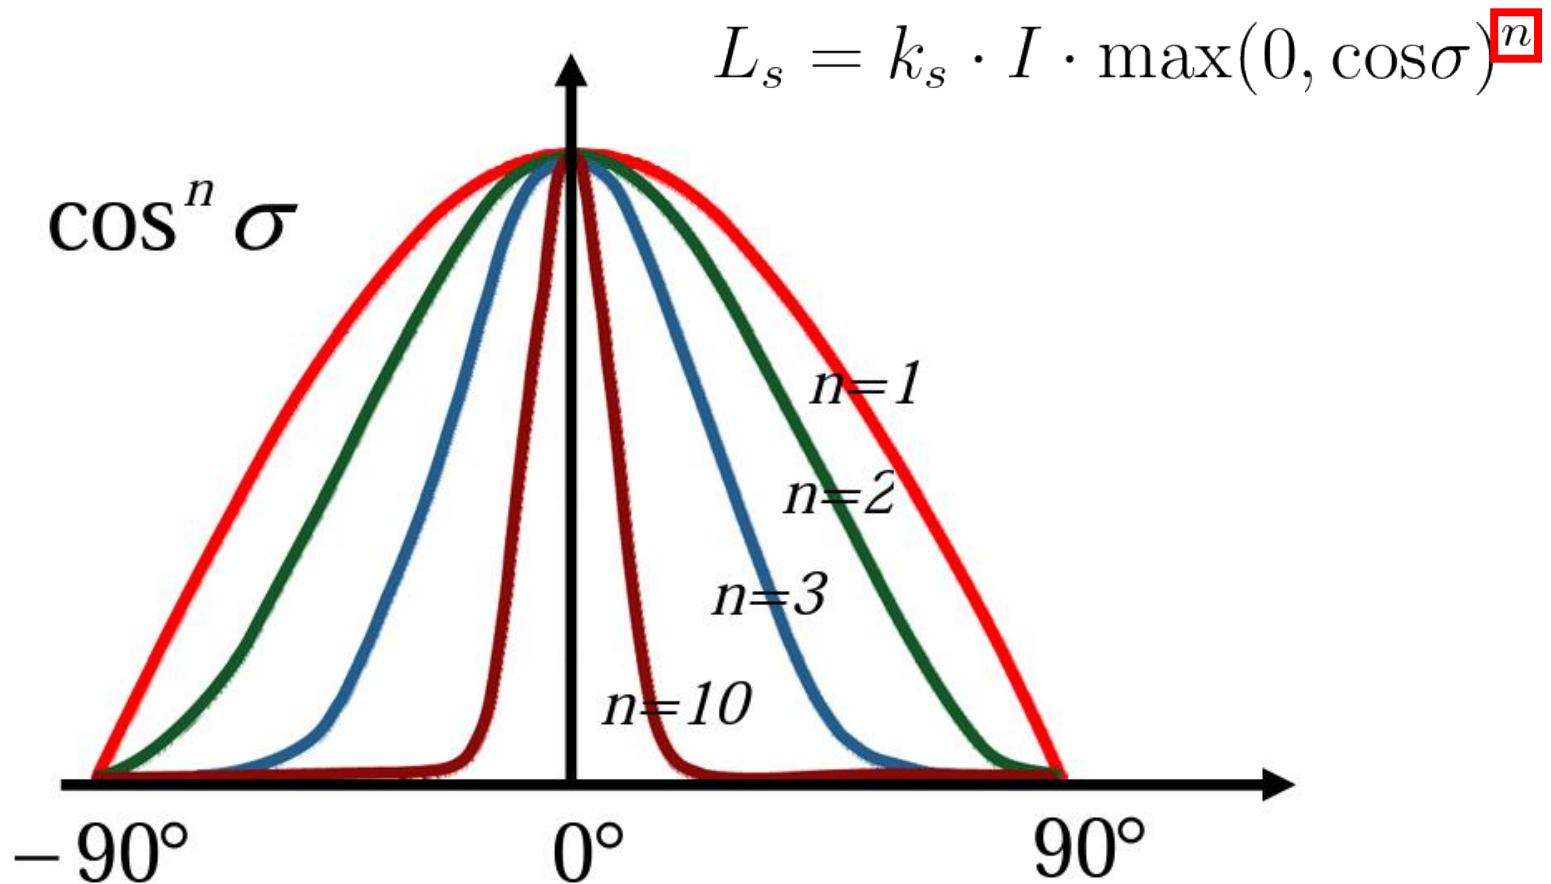
specular exponent

$$\begin{aligned}
 L_s &= k_s \cdot I \cdot \max(0, \cos\sigma)^n \\
 &= k_s \cdot I \cdot \max(0, vE \cdot vR)^n
 \end{aligned}$$

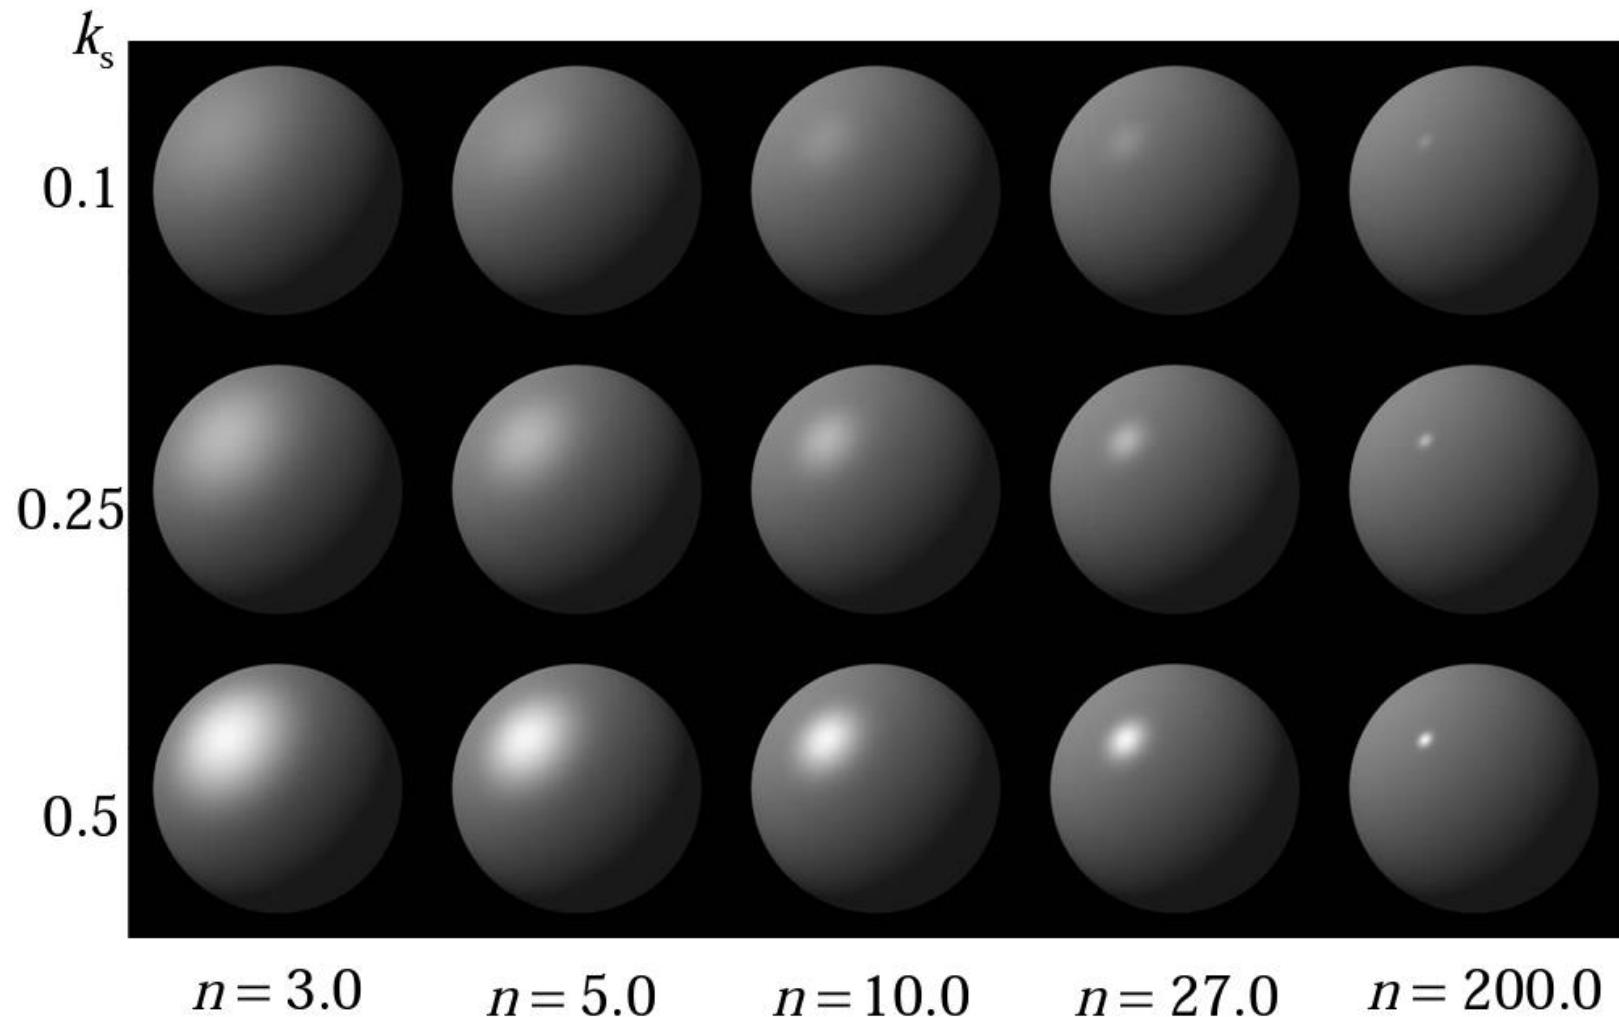
↑
specular coefficient
specularly reflected light

Specular Shading (cont.)

- Increase n narrows the lobe

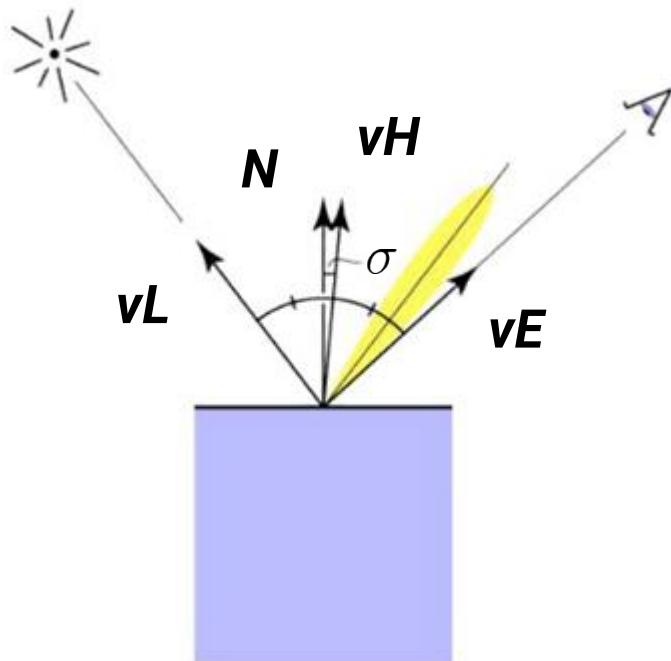


Specular Shading (cont.)



Phong specular Variant: Blinn-Phong

- Rather than computing reflection directly, just compare to normal bisection property
- One can prove $\cos^n(\sigma) = \cos^{4n}(\alpha)$

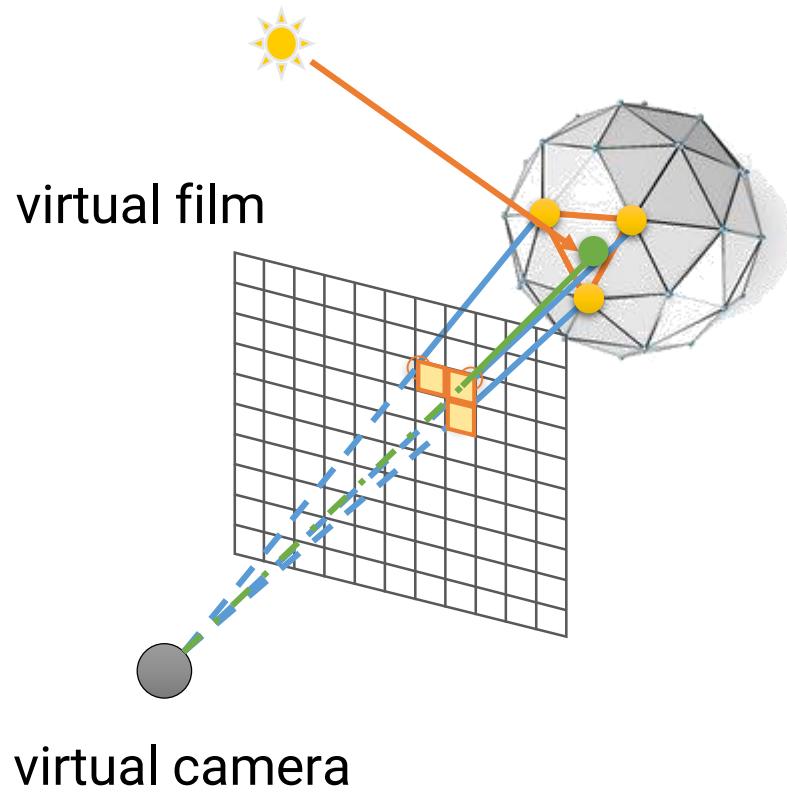


$$\begin{aligned} vH &= \text{bisector}(vL, vE) \\ &= \frac{(vL + vE)}{\|vL + vE\|} \end{aligned}$$

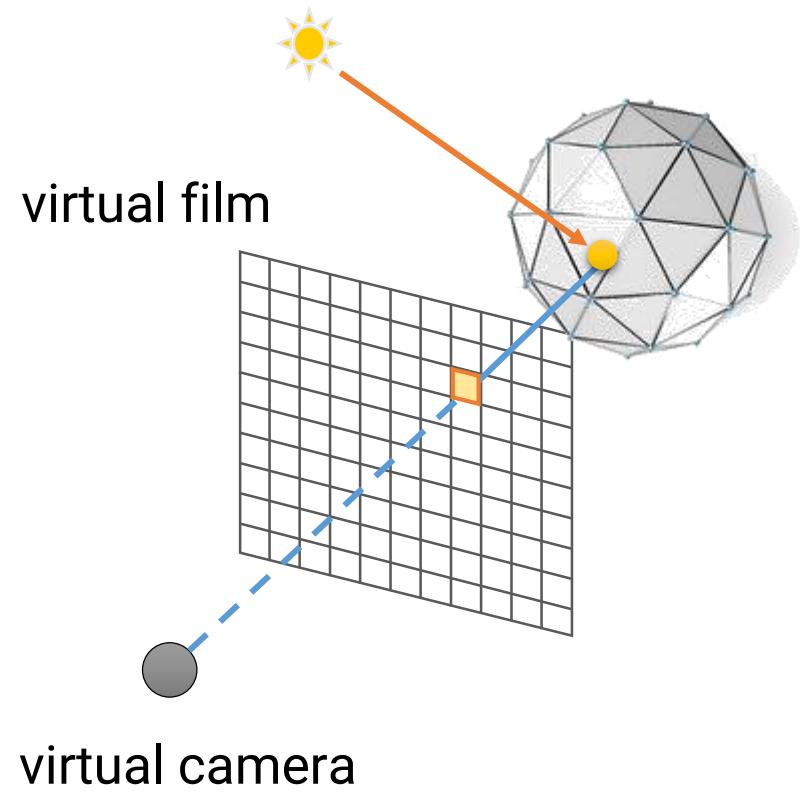
$$\begin{aligned} L_s &= k_s \cdot I \cdot \max(0, \cos\sigma)^n \\ &= k_s \cdot I \cdot \max(0, N \cdot vH)^n \end{aligned}$$

Shading in Rasterization v.s. Ray Tracing

Rasterization

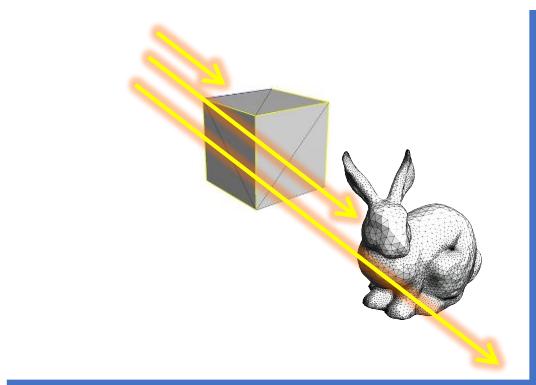


Ray tracing

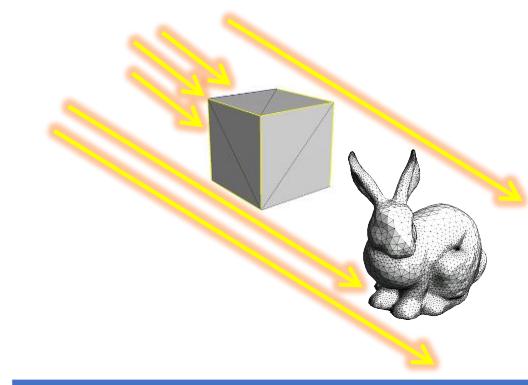


Local, Direct, and Global Illumination

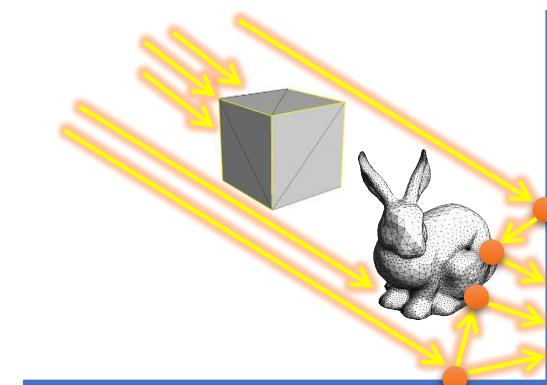
- Direct illumination considers only the **direct** contribution of lights
- Local illumination can be considered as direct lighting **without occlusion** (all lights are fully visible, no shadows)
- Global illumination includes **multi-bounce** illumination reflected from other surfaces (**recursive** computation!)



local illumination



direct illumination

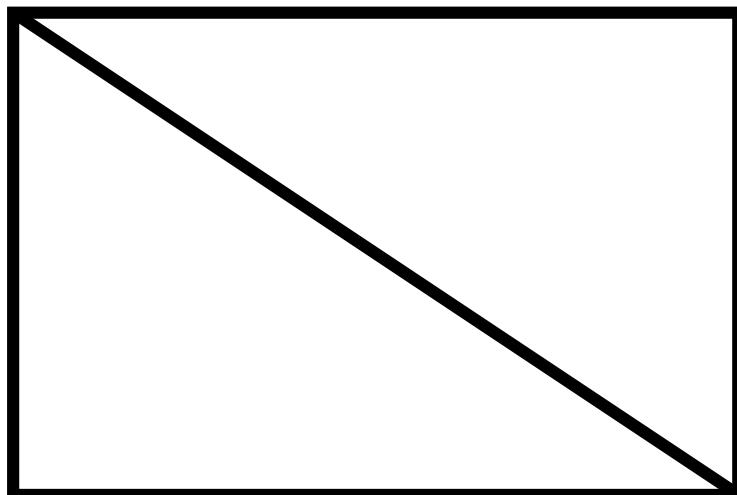


global illumination

Textures

Textures

- Can be used to represent **spatially-varying** data
- Can **decouple** materials from the geometry



Geometry: two triangles

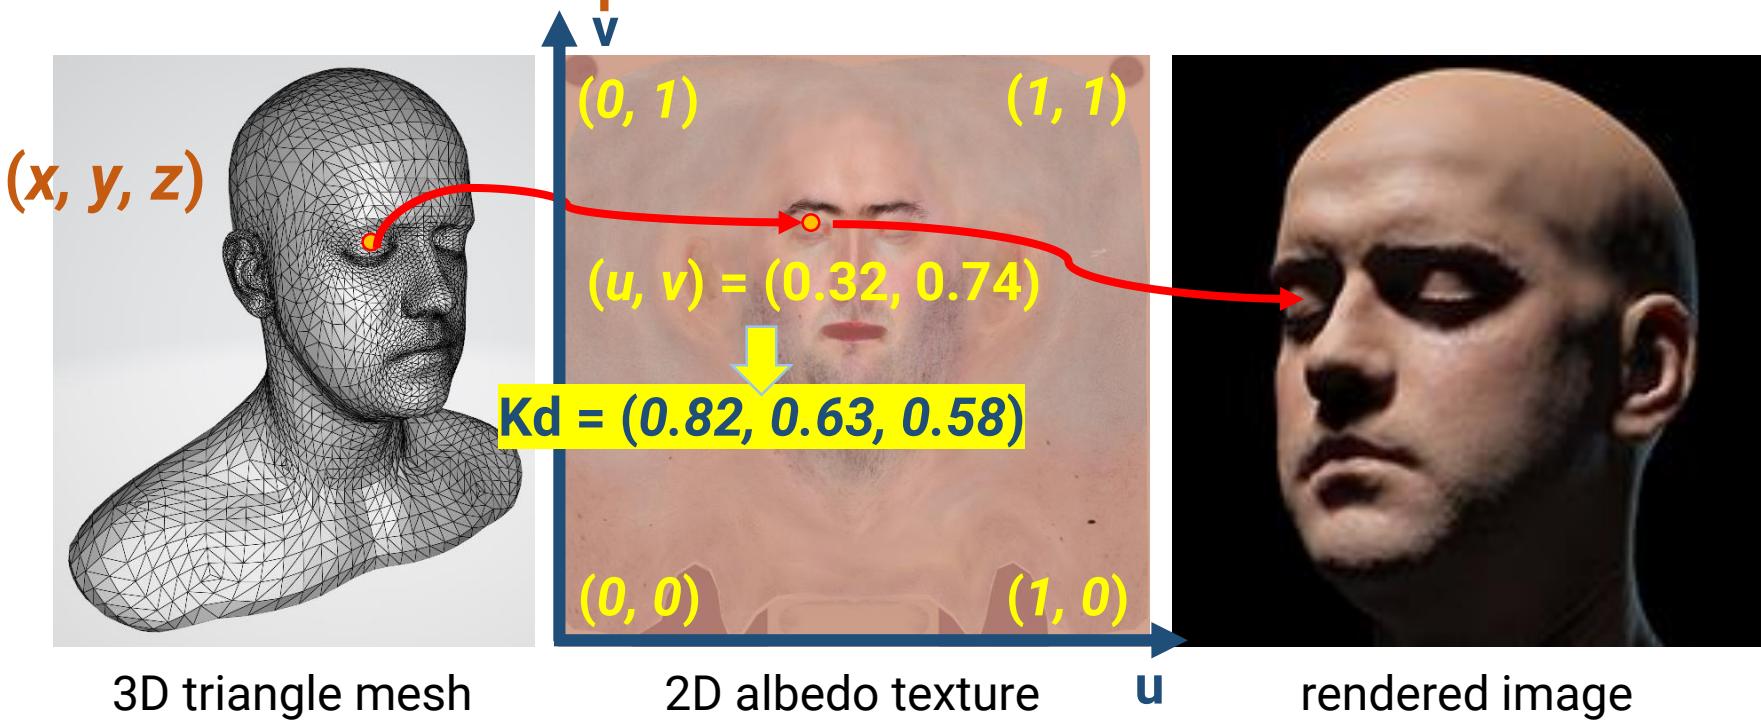
Material: 2D image texture



complex appearance

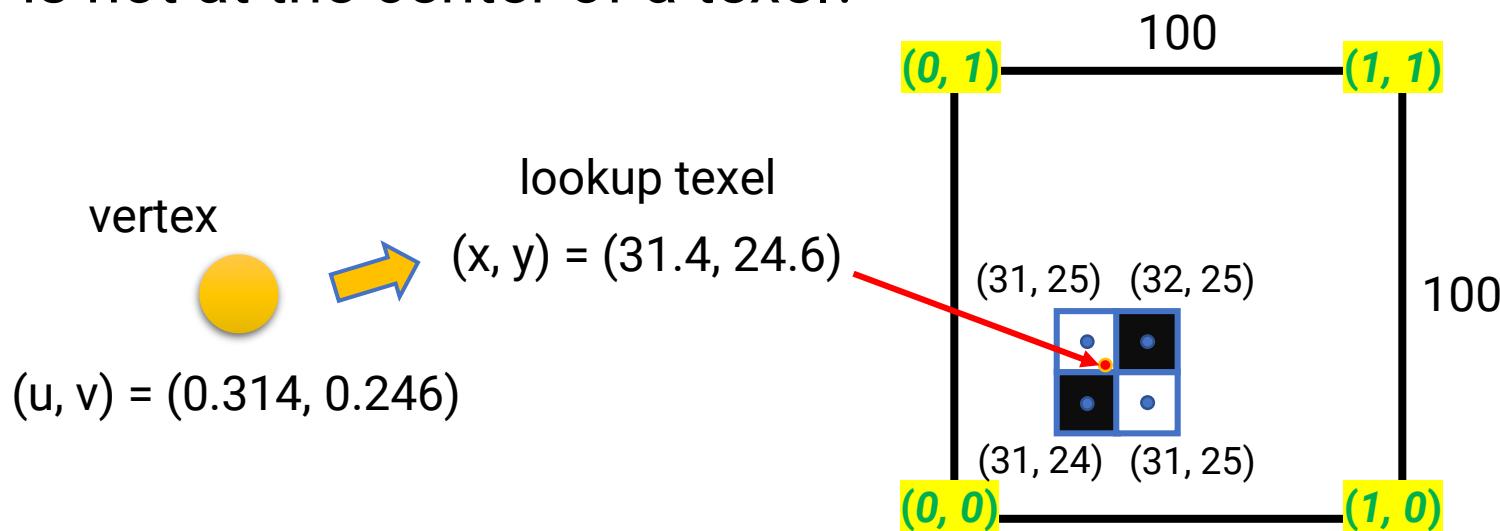
Texture Coordinate

- A coordinate to look up the texture
- The way to map a point on an **arbitrary 3D surface** to a pixel (texel) on an **image** texture
 - Need **surface parameterization**



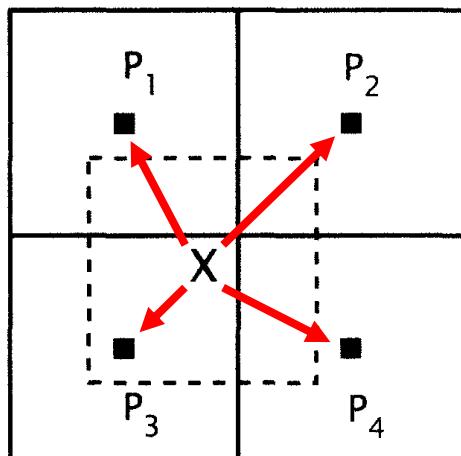
Texture Filtering

- Like an image, the content in a 2D texture is **discretely** represented by texels
- The texture coordinates can be **continuous** (especially after interpolation by the rasterization)
- How to determine the texture value if the lookup point is not at the center of a texel?



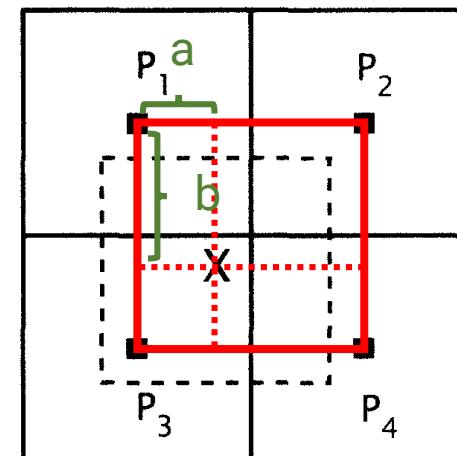
Texture Filtering (cont.)

- Strategies
 - **Nearest neighbor**
 - **Bilinear interpolation**



nearest neighbor

P₃ is closest
Use P₃'s pixel value



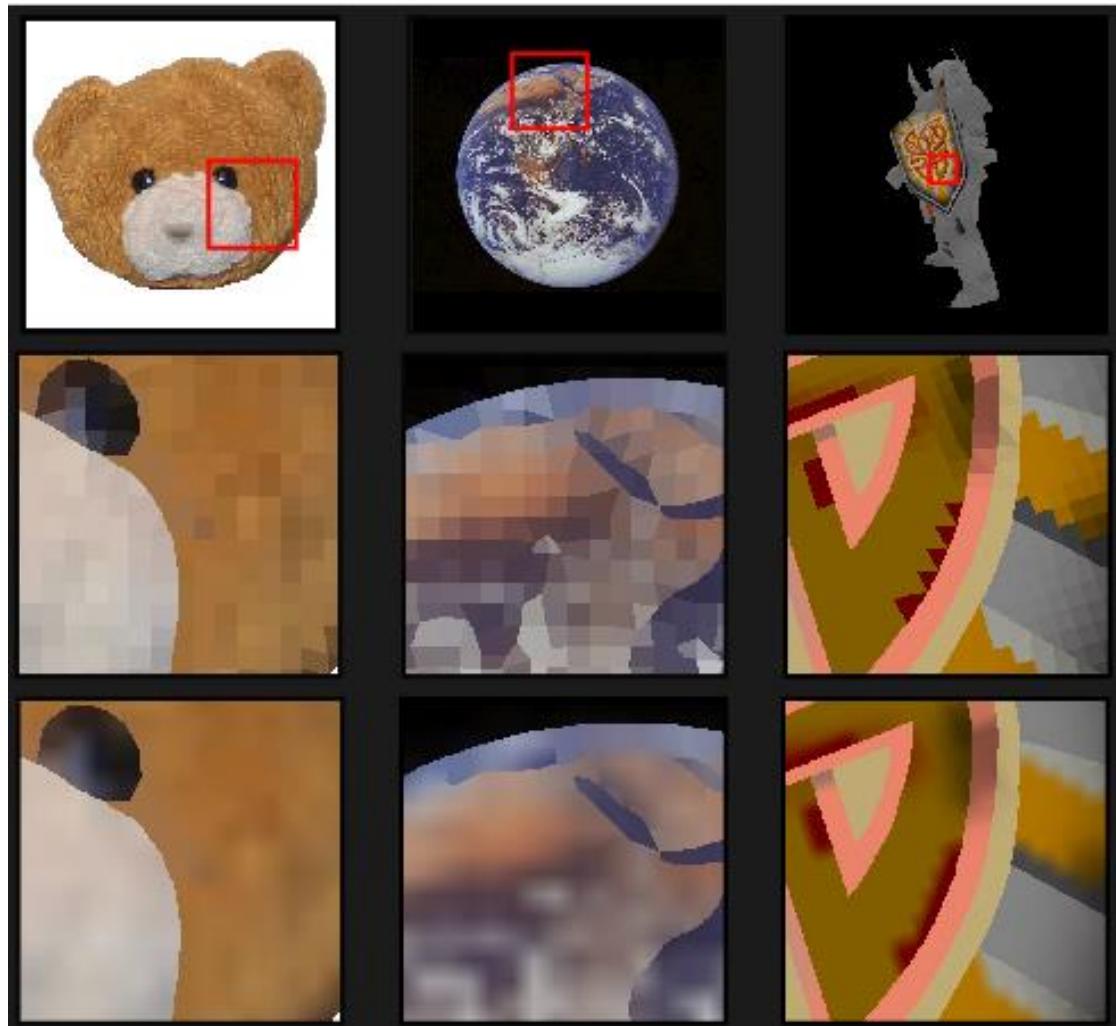
bilinear interpolation

$$(1-a)(1-b)P_1 + (a)(1-b)P_2 + (1-a)(b)P_3 + (a)(b)P_4$$

Texture Filtering (cont.)

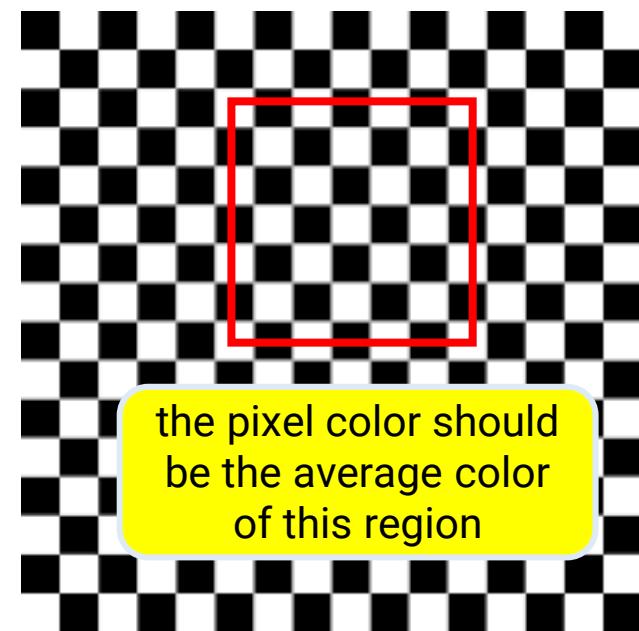
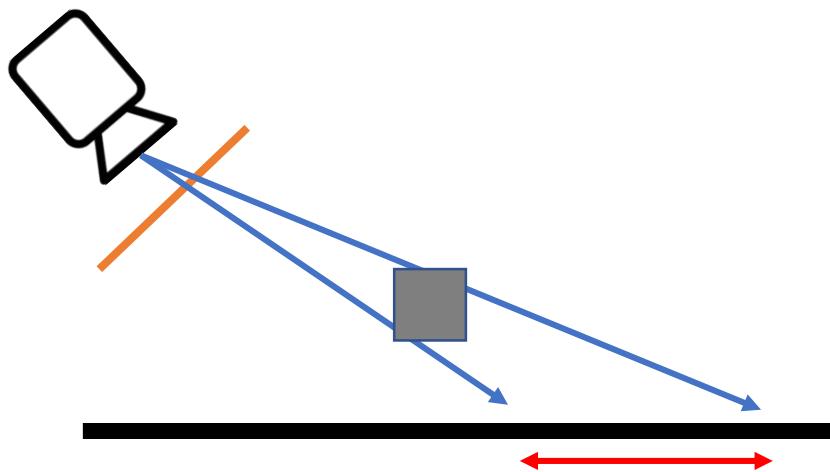
nearest
neighbor

bilinear
interpolation



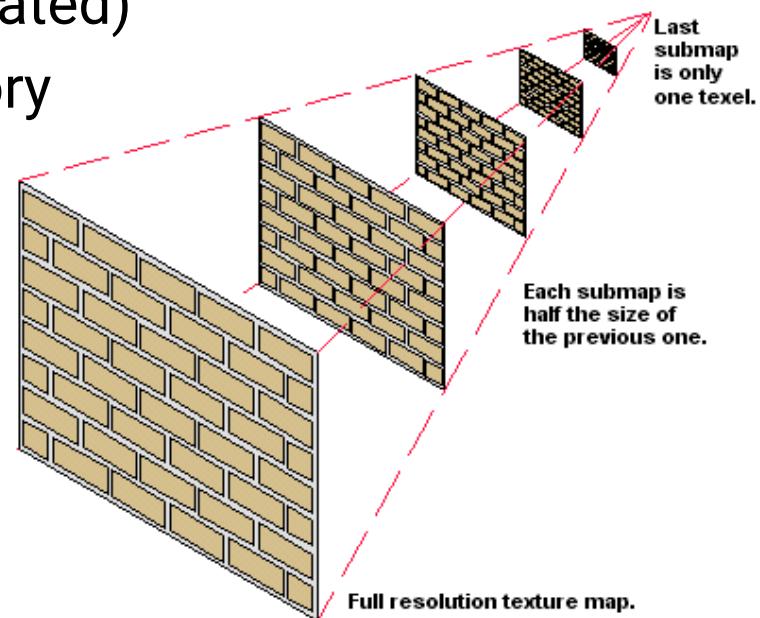
Mipmap

- To avoid aliasing, we should determine the regions a pixel covers (footprint) and average all the texture values inside the regions
- Time-consuming to do this in the run time!



Mipmap (cont.)

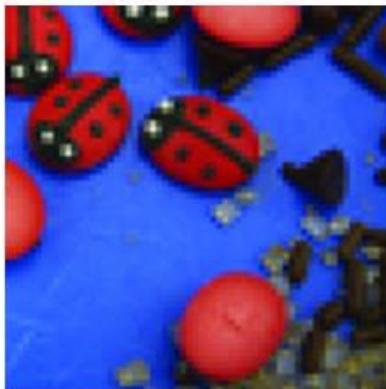
- Mipmap provides a clever way to solve this problem
- **Pre-process**
 - Build a **hierarchical representation** of the texture image
 - Each level has a half resolution of its previous level (generated by linearly interpolated)
 - Take at most **1/3** more memory



Mipmap (cont.)



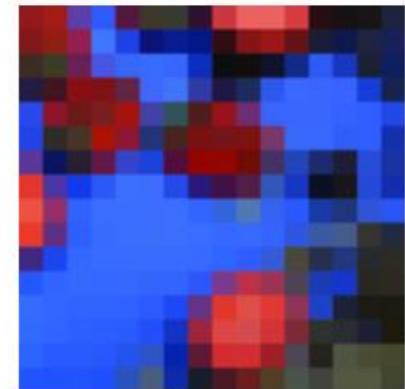
Level 0 = 128x128



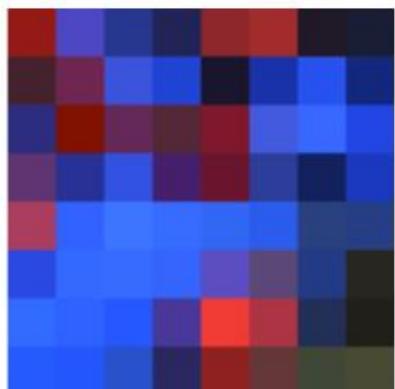
Level 1 = 64x64



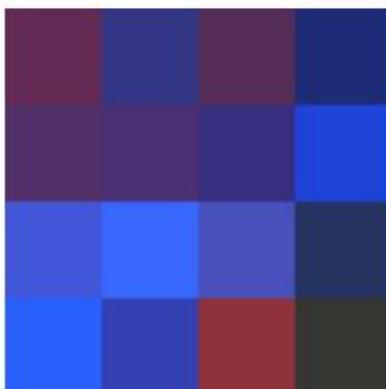
Level 2 = 32x32



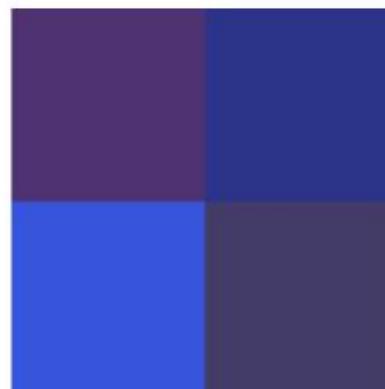
Level 3 = 16x16



Level 4 = 8x8



Level 5 = 4x4



Level 6 = 2x2

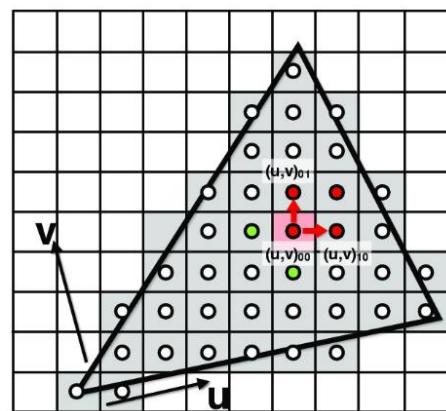


Level 7 = 1x1

Mipmap (cont.)

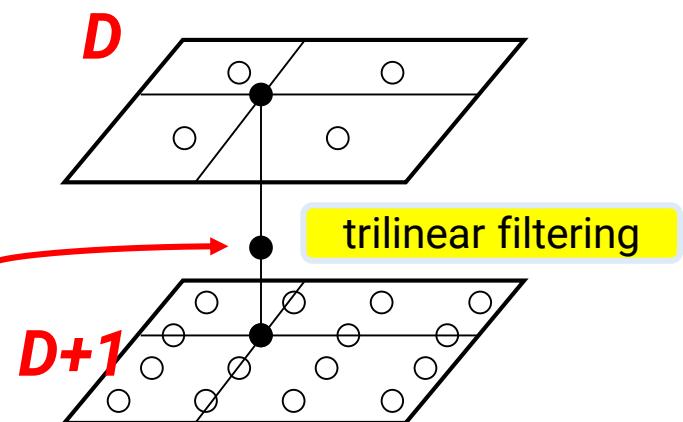
- **Run-time lookup**

- Use **screen-space texture coordinate** to estimate its footprint in the texture space
- Choose two levels D and $D+1$ based on the footprint
- Perform linear interpolation at level D to obtain a value V_D
- Perform linear interpolation at level $D+1$ to obtain V_{D+1}
- Perform linear interpolation between V_D and V_{D+1}

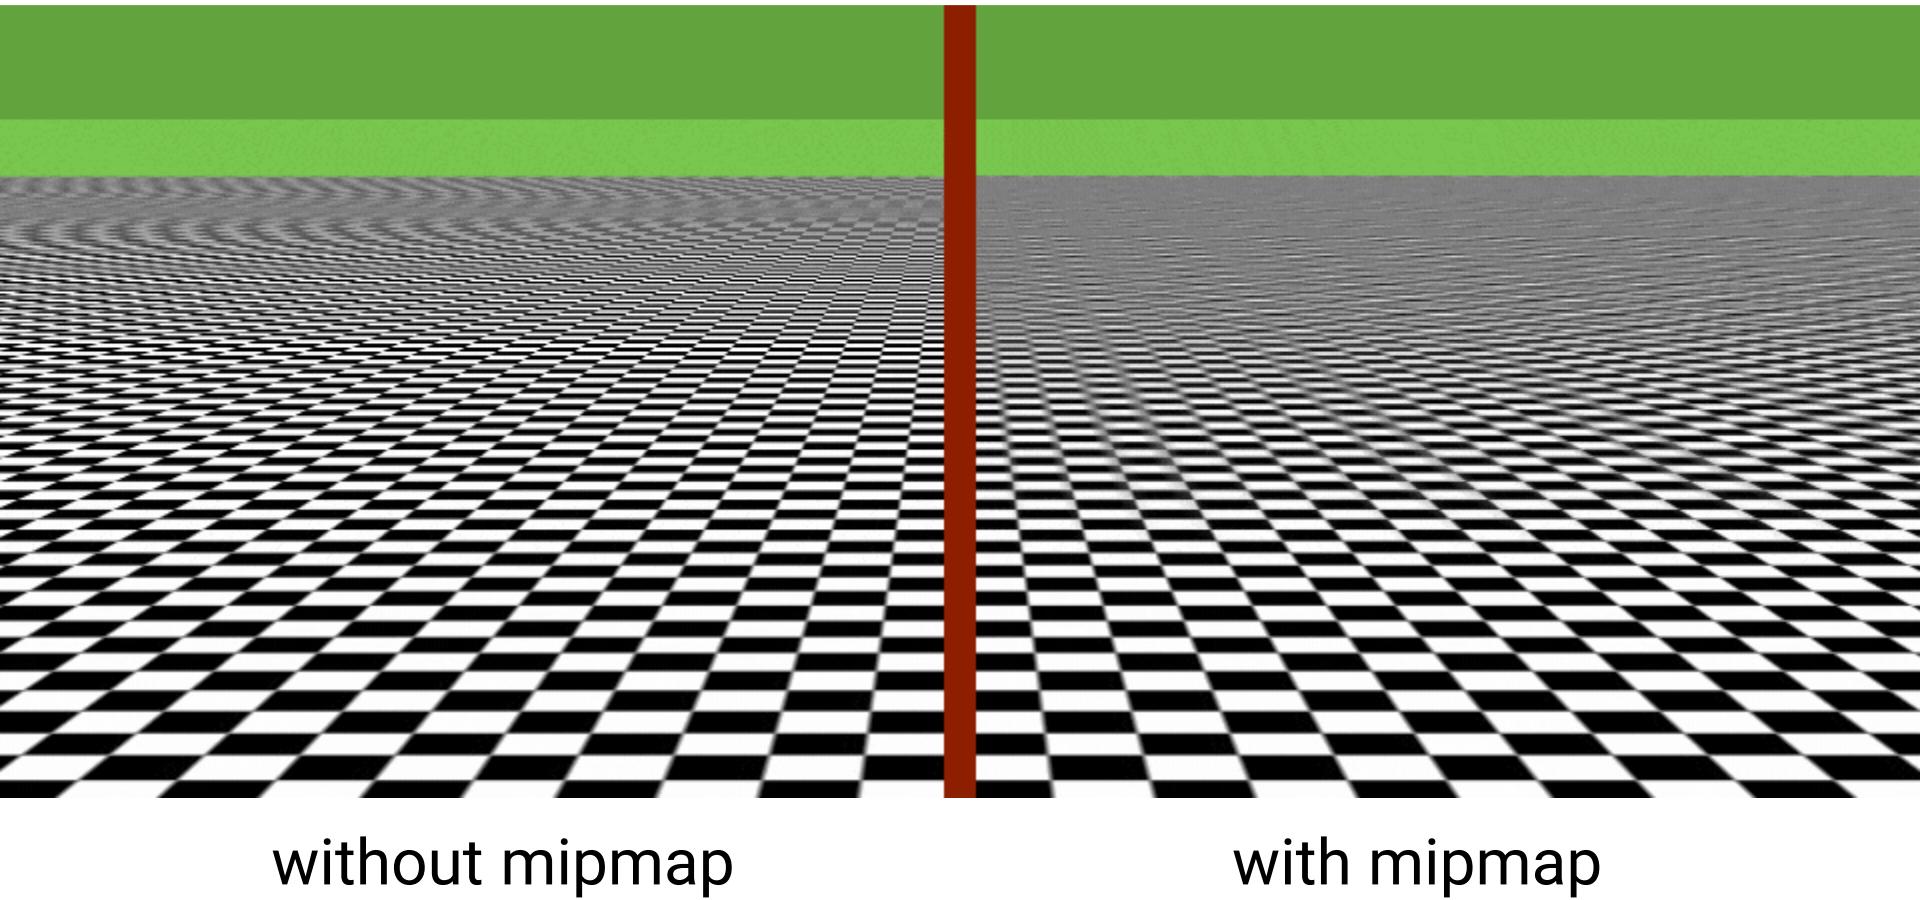


$$\frac{1}{w} = 2^{n-1-l}$$

$$l = n - 1 + \log w$$



Mipmap (cont.)



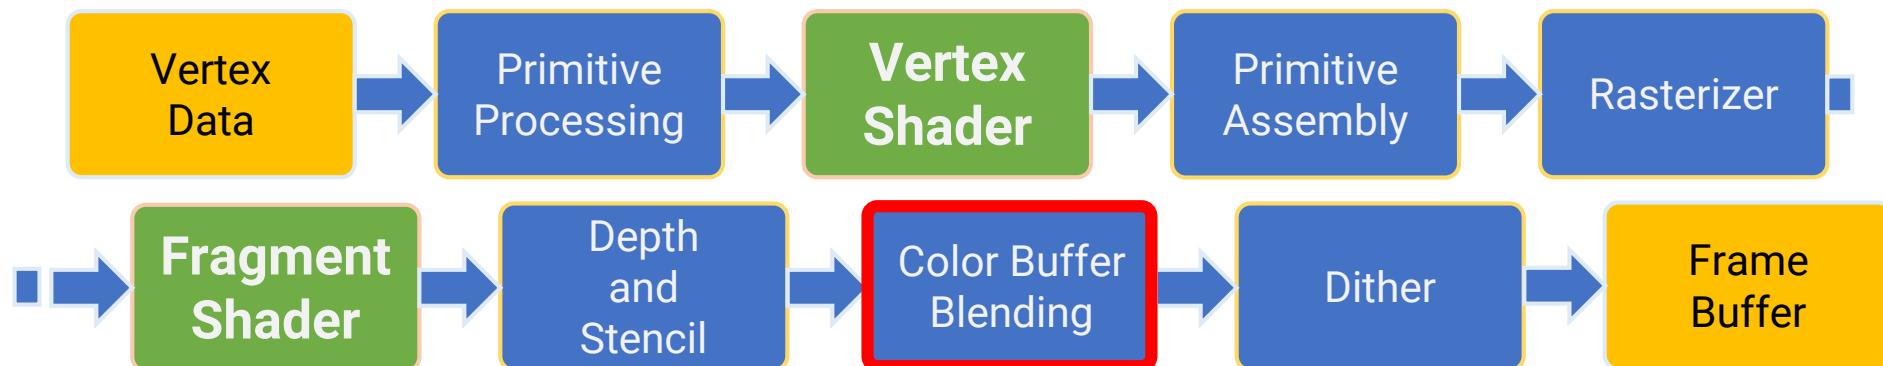
Transparency

Transparency in Rasterization

- Z-buffer provides an efficient way to solve the visibility for opaque objects
 - Keep the closest objects from the camera during the rendering process
- For transparent objects, we will see occluded objects through the transparent one
- However, in rasterization transparency is difficult to resolve **because each polygon only has its own information**
 - It does not know which triangle locates behind, so it cannot determine the pixel color in its fragment shader

Transparency in Rasterization (cont.)

- The GPU rendering pipeline provides a way to simulate transparency
- Major idea
 - Render transparent objects **in an order w.r.t their distance to the camera** (farther objects first)
 - When rendering transparent objects, blend the fragment color with the previous results in the color buffer



Blending Equation

- OpenGL provides flexibility to composite the fragment color when rendering transparent objects

$$\bar{C}_{result} = \bar{C}_{source} * F_{source} + \bar{C}_{destination} * F_{destination}$$

blending equation

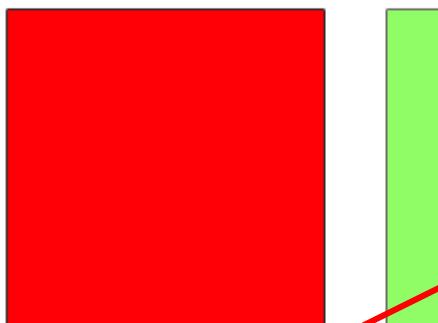
- \bar{C}_{source} : the source color vector (the color output by the fragment shader)
- $\bar{C}_{destination}$: the destination color vector (the color vector currently stored in the color buffer)
- F_{source} : the source factor value (set the impact of the alpha value on the source color)
- $F_{destination}$: the destination factor value (set the impact of the alpha value on the destination color)

Blending Equation (cont.)

- Example

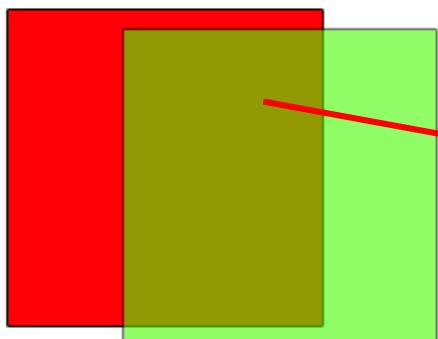
$$\bar{C}_{result} = \bar{C}_{source} * \bar{F}_{source} + \bar{C}_{destination} * \bar{F}_{destination}$$

GL_SRC_ALPHA GL_ONE_MINUS_SRC_ALPHA



Alpha: the transparency of the surface

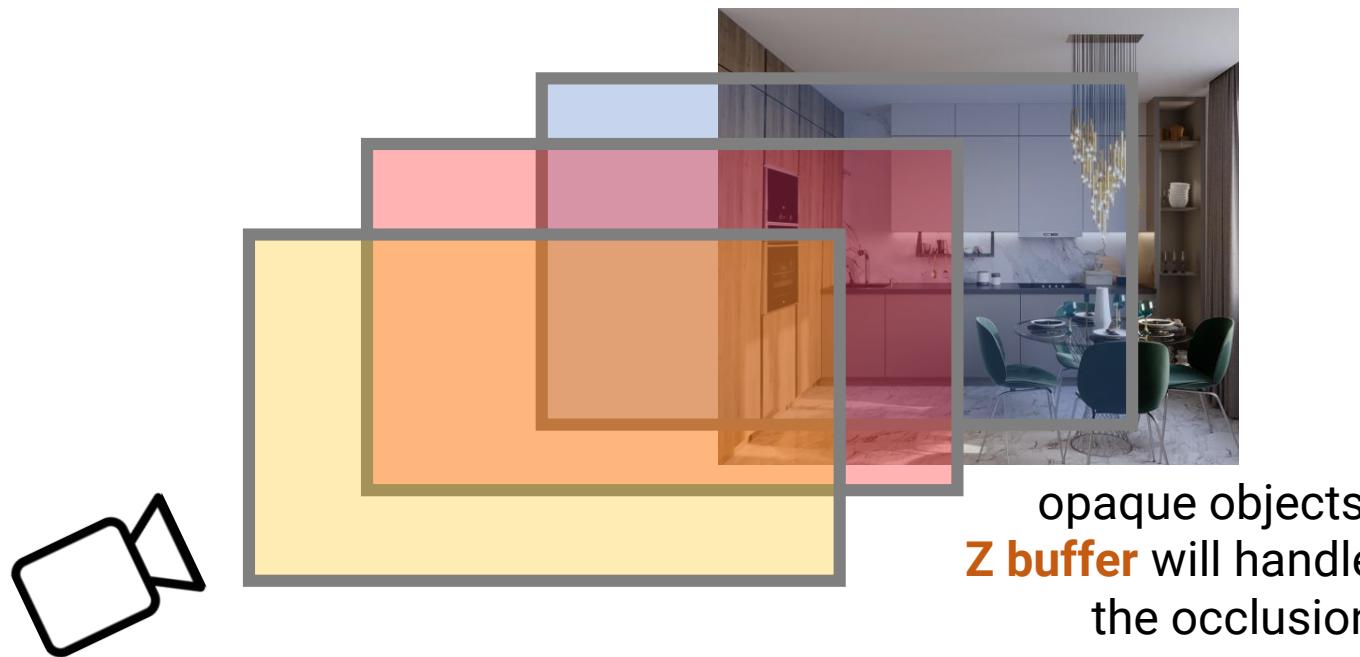
- For opaque surfaces, alpha = 1
- For totally transparent surfaces, alpha = 0
- For semi-transparent surfaces, alpha is between 0 and 1



$$\begin{aligned}
 & (0.0, 1.0, 0.0, 0.6) * 0.6 + \\
 & (1.0, 0.0, 0.0, 1.0) * (1 - 0.6) \\
 = & (0.4, 0.6, 0.0, 0.76)
 \end{aligned}$$

Rendering Algorithm for Transparency

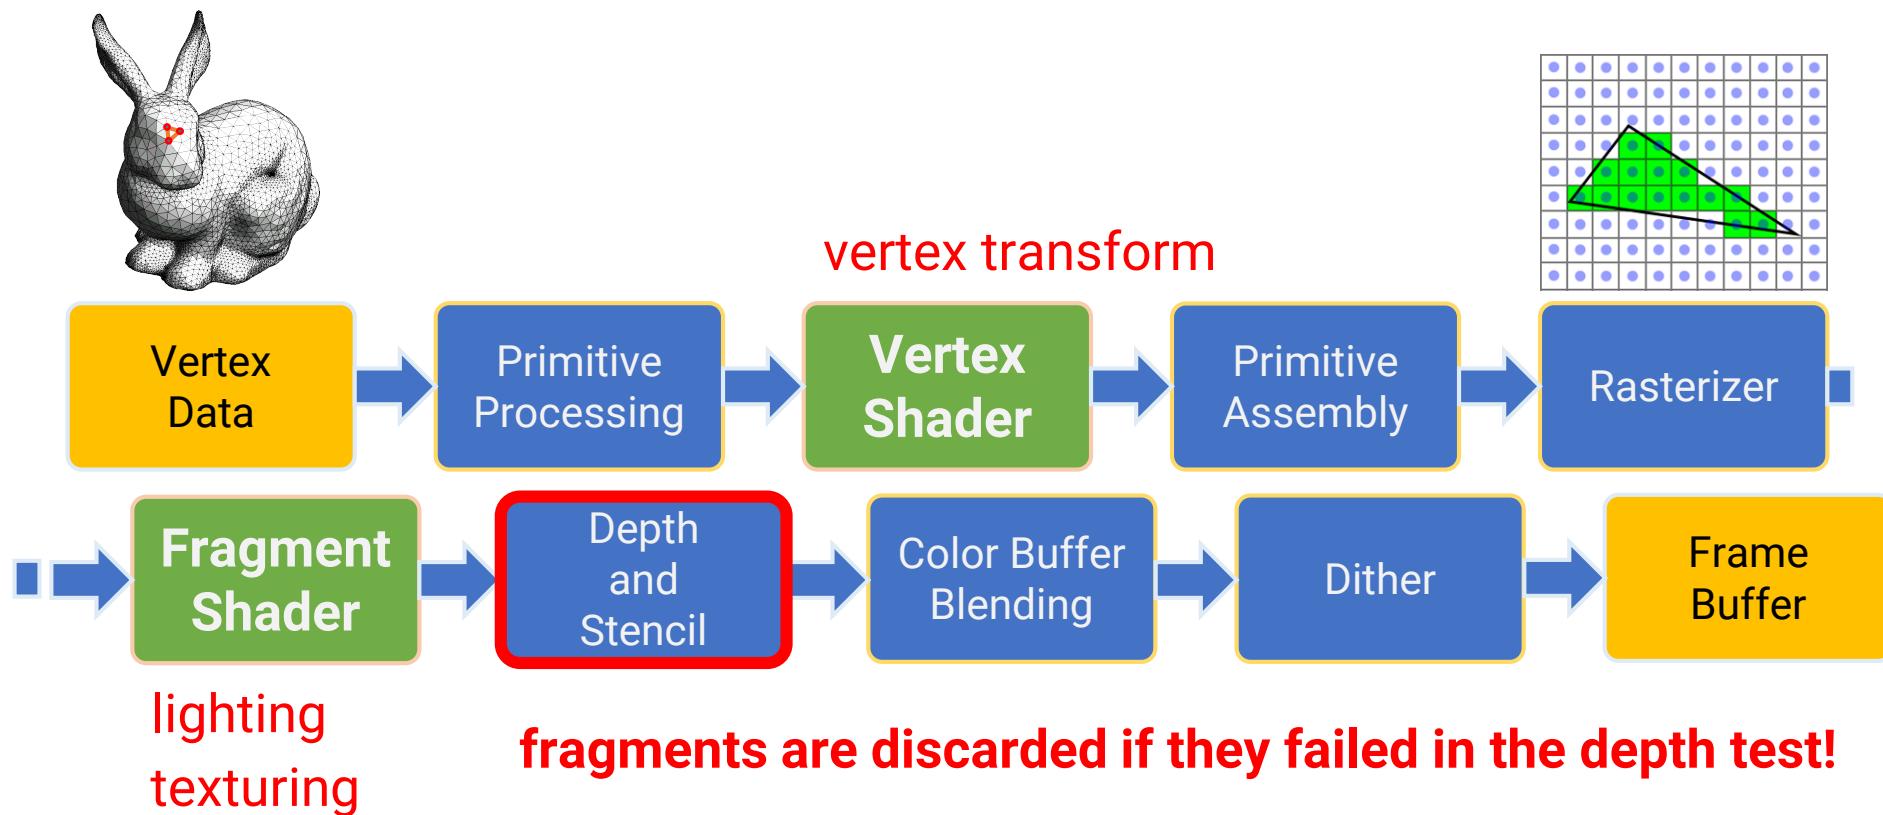
- Render **opaque** objects first **in any order**
- Render **transparent** objects **in an order w.r.t their distance to the camera** (farther objects first)
- Each time, blend with the color buffer



Deferred Shading

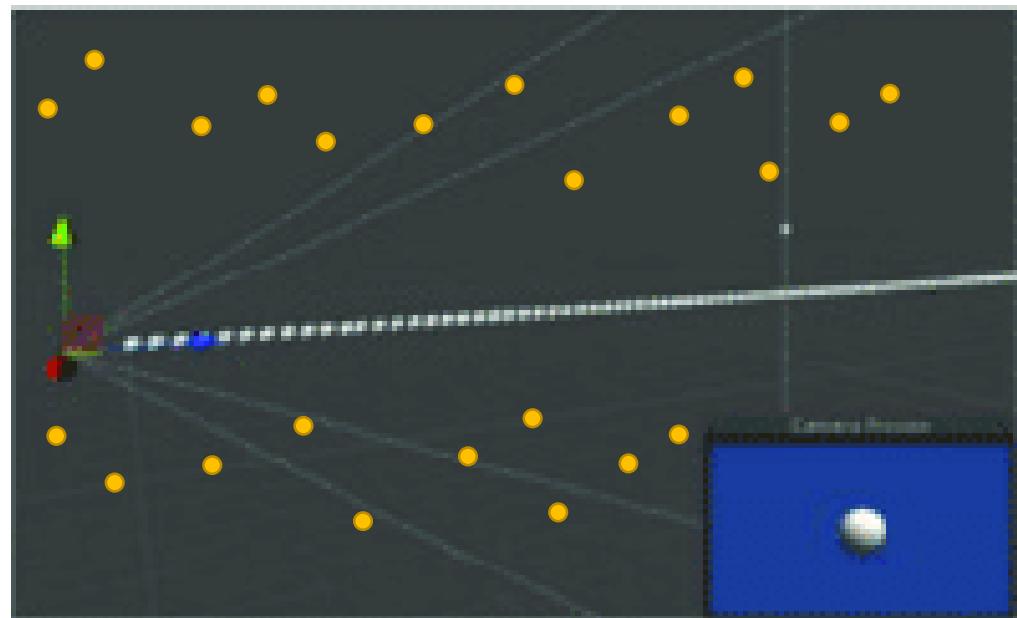
Problem Formulation

- Forward rendering



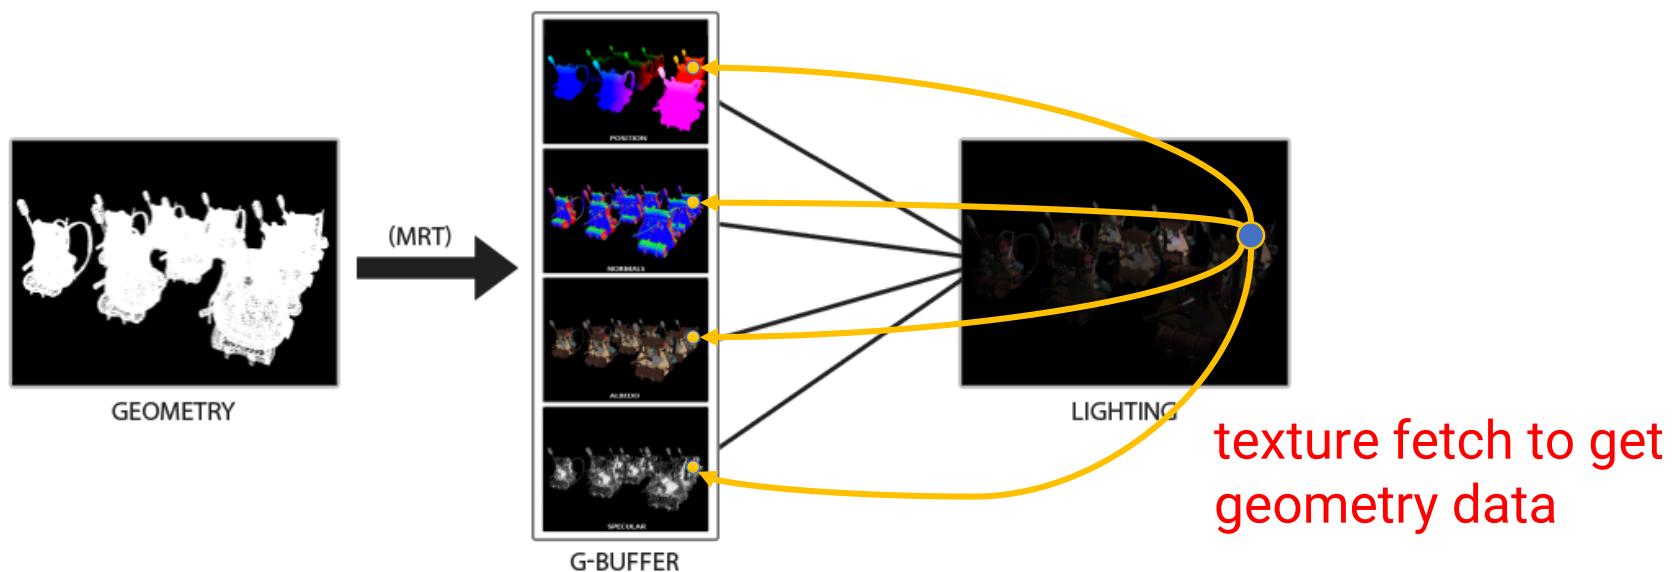
Problem Formulation (cont.)

- Problem of forward rendering
 - In scenes with **many** lights and **complex** layouts, lots of computation resources are wasted on shading the **occluded** surfaces that will finally be discarded!
 - Overdraw per pixel!



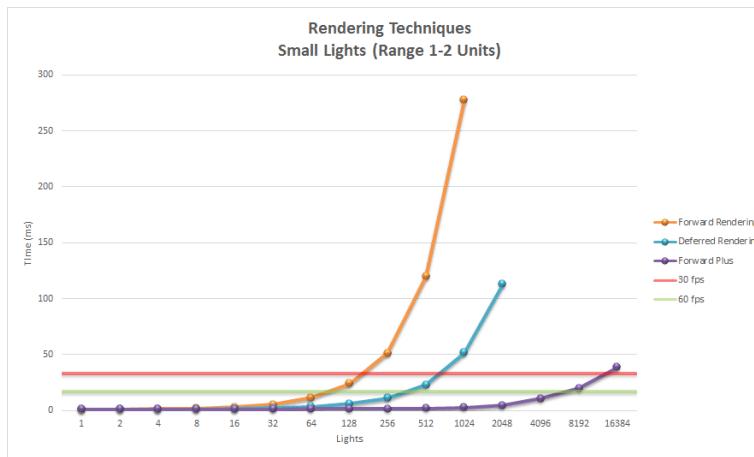
Deferred Shading

- A **Two-pass** rendering algorithm
 - In the first pass, recognize all **visible** surfaces from the camera, store their **geometry** and **material** properties in **geometry buffers (G-buffers)**
 - In the second pass, only compute lighting on the visible surfaces based on the G-buffers



Deferred Shading (cont.)

- Pros
 - Reduce unnecessary lighting computation

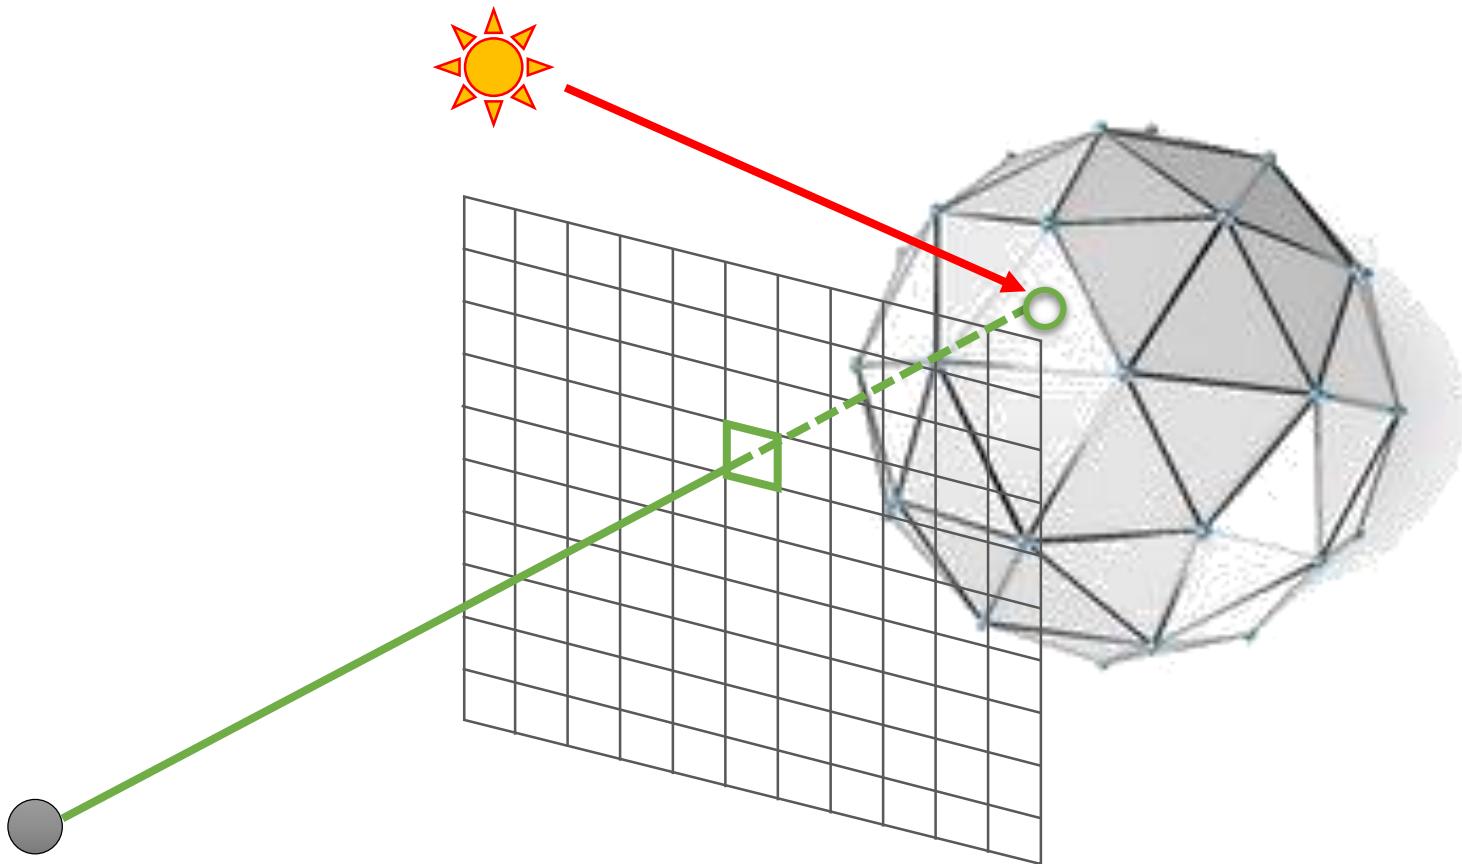


- Cons
 - Large memory bandwidth
 - Difficult for multi-sampled anti-aliasing (MSAA)
 - Transparent objects should be rendered using forward rendering

Shadow Mapping

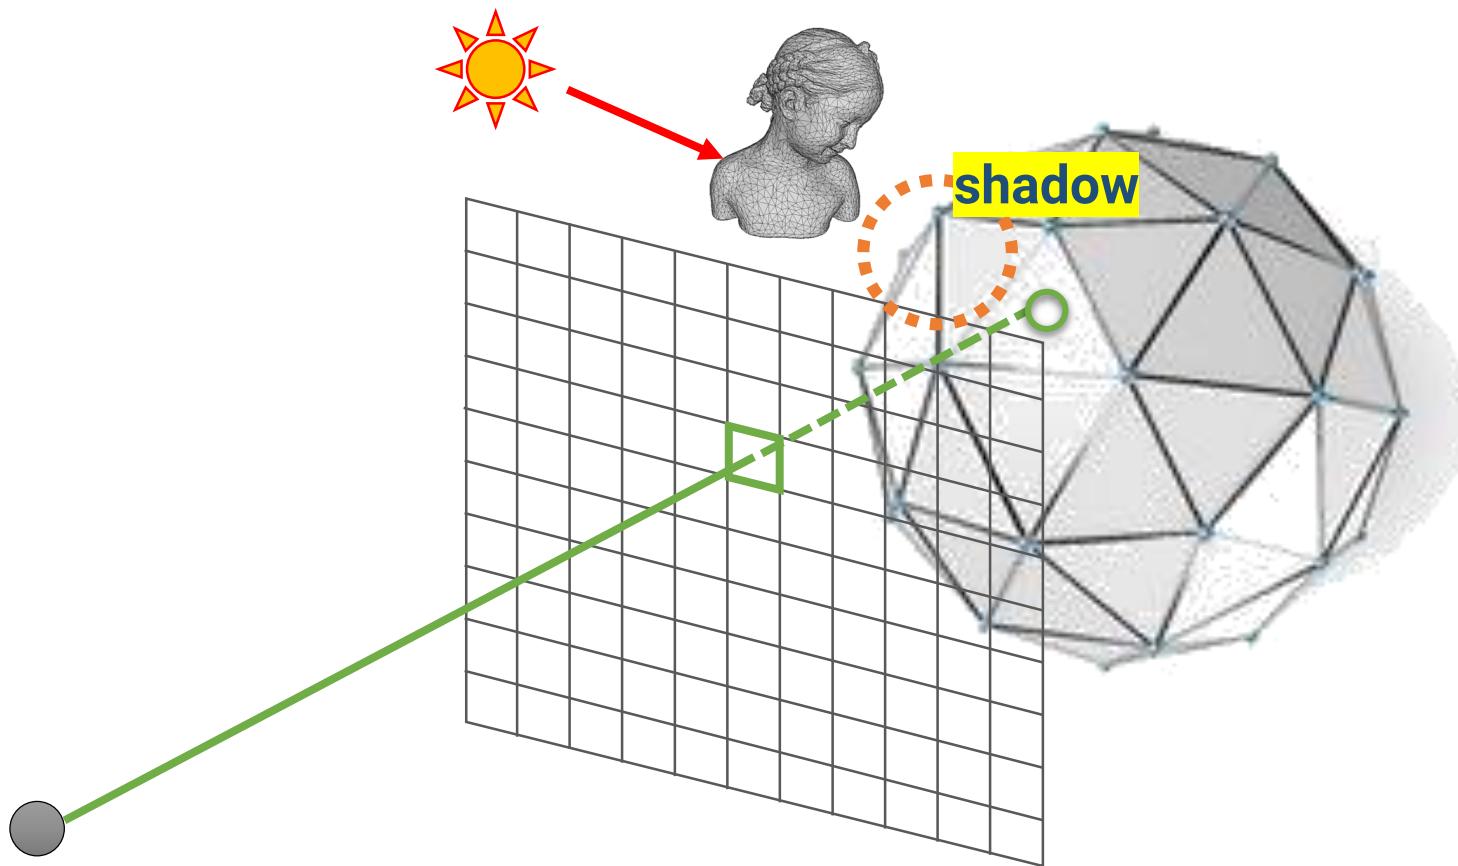
Shadow

- The standard GPU rendering pipeline can only handle unoccluded illumination (local illumination)



Shadow (cont.)

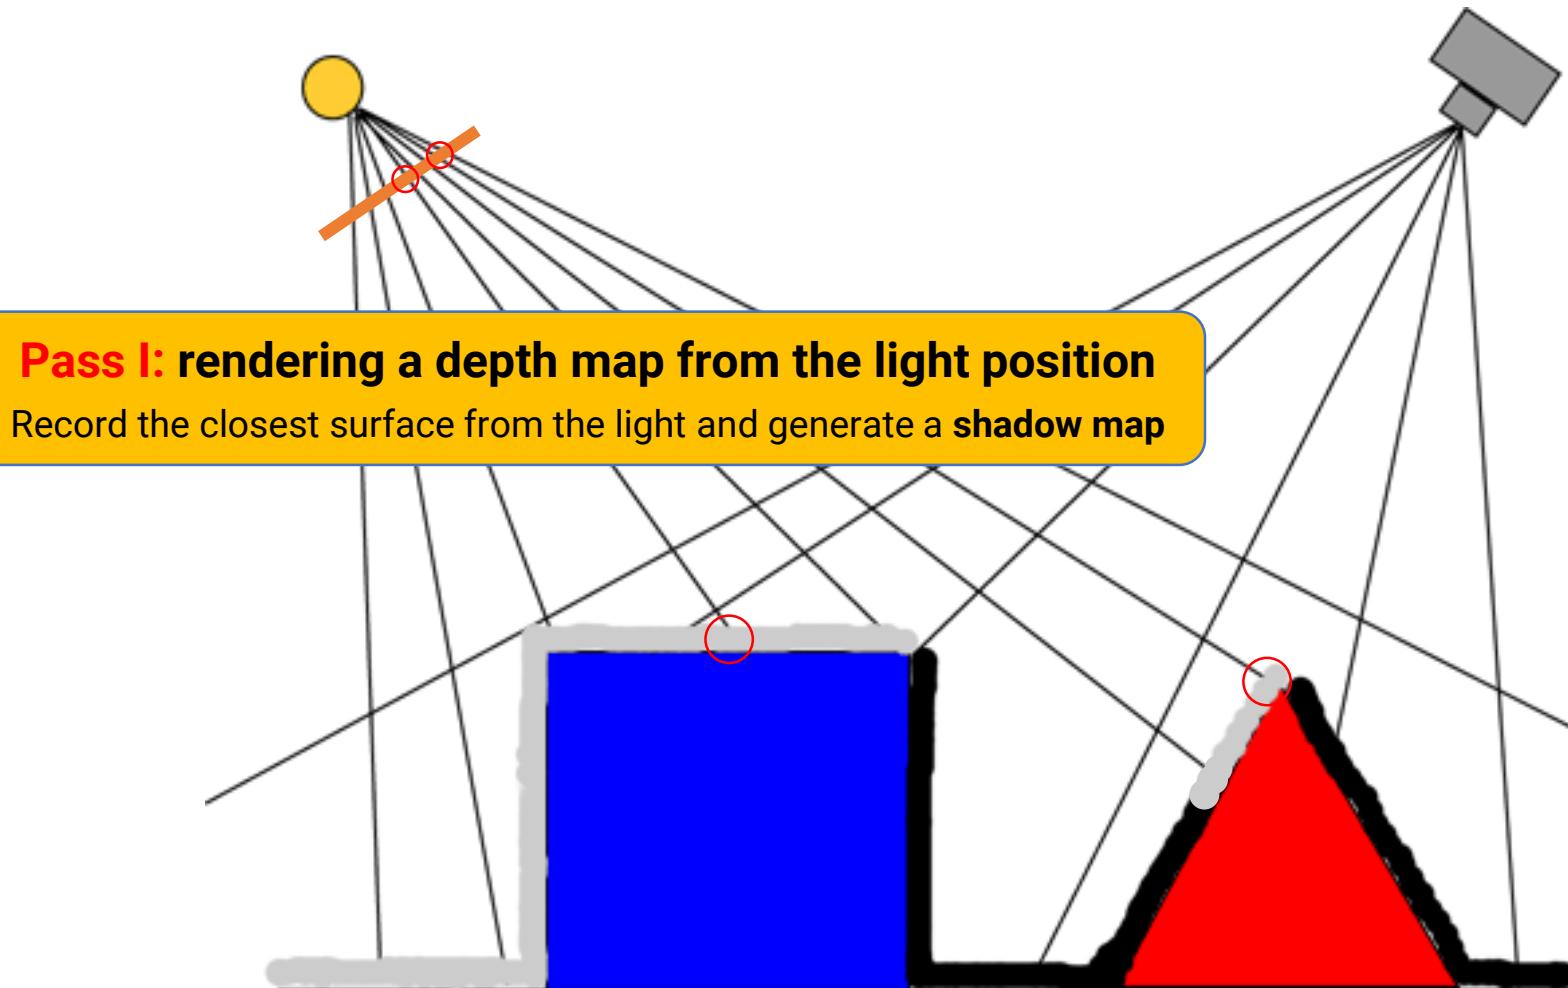
- It is common that a lighting direction is occluded by some other objects



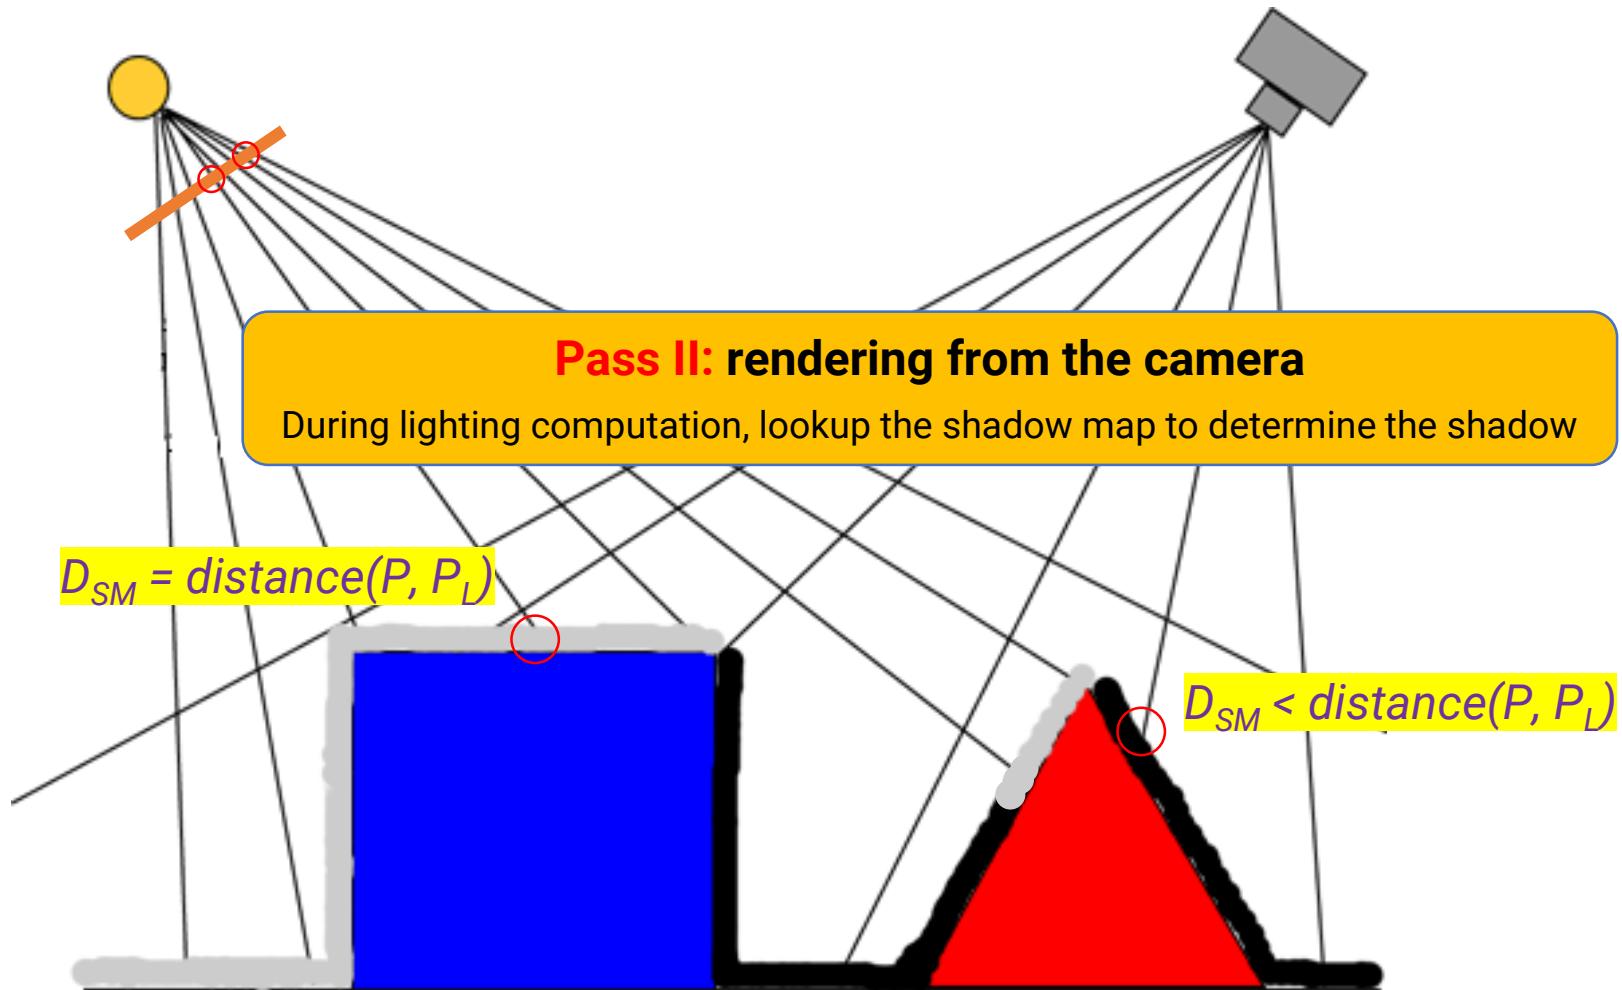
Shadow Map Overview

- Like the case of transparency, rendering shadows is difficult for rasterization because **each polygon only has its own information**
 - It does not know which triangle blocks the light, so it cannot determine the shadow attenuation in its fragment shader
- **Shadow map** is a two-pass rendering technique for simulating shadows using rasterization

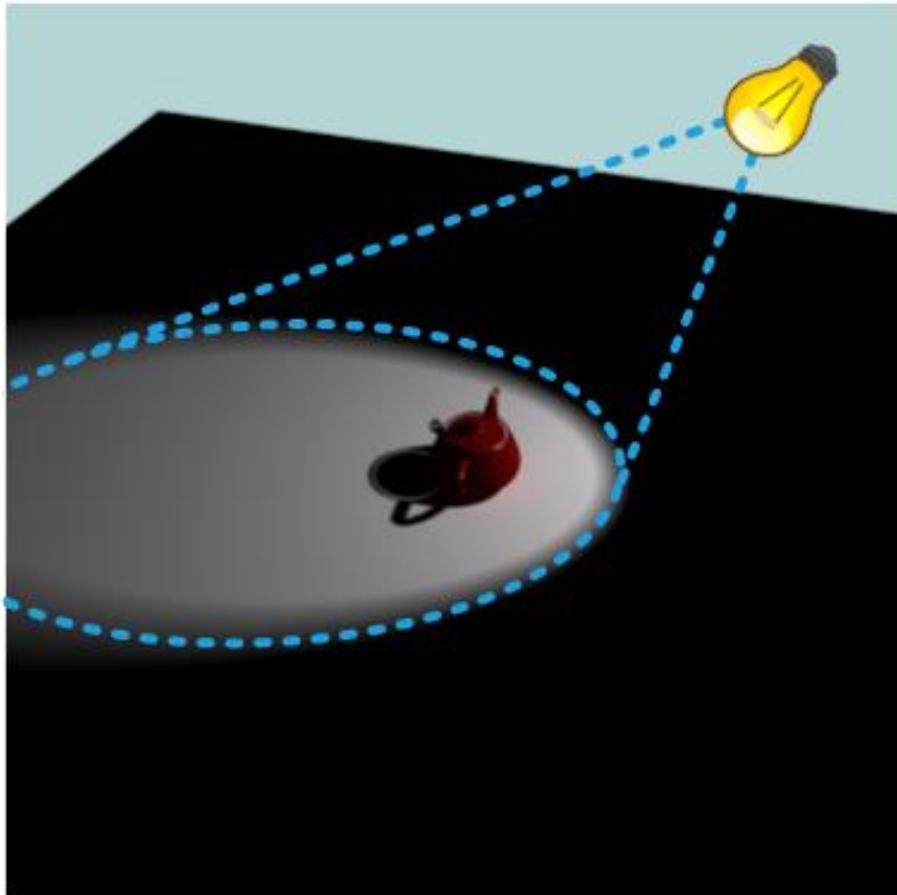
Shadow Map Overview (cont.)



Shadow Map Overview (cont.)



Shadow Map Overview (cont.)



final rendering
(rendering from the camera view)

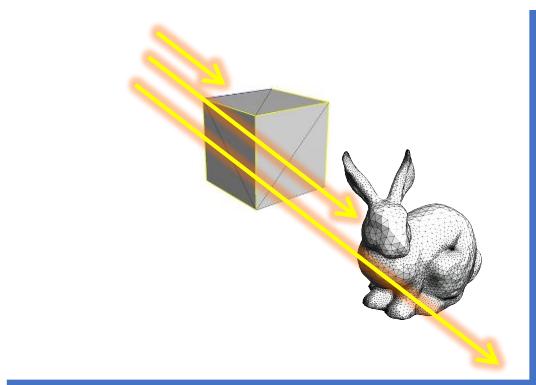


shadow map
(rendering from the light view)

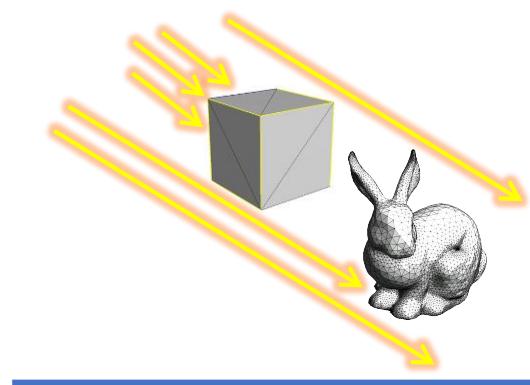
Ambient Occlusion

Recap: Global Illumination

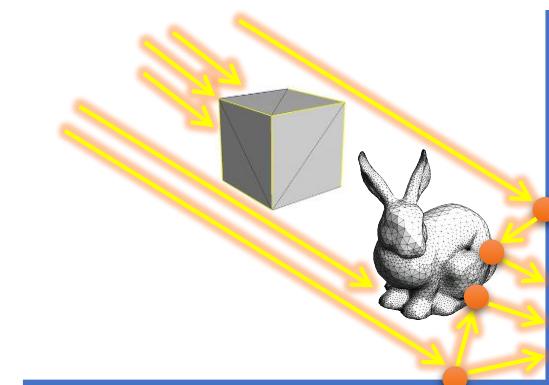
- Global illumination includes multi-bounce lighting
- Very expensive to compute
- In Phong lighting model, a **constant ambient term** is used to account for disregarded illumination
 - However, this produces a “flat”, “non-photo-realistic” appearance



local illumination



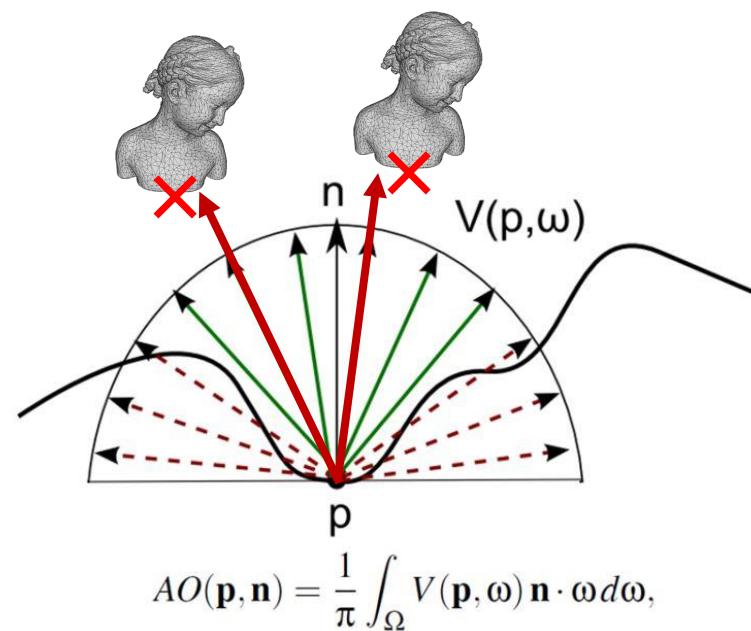
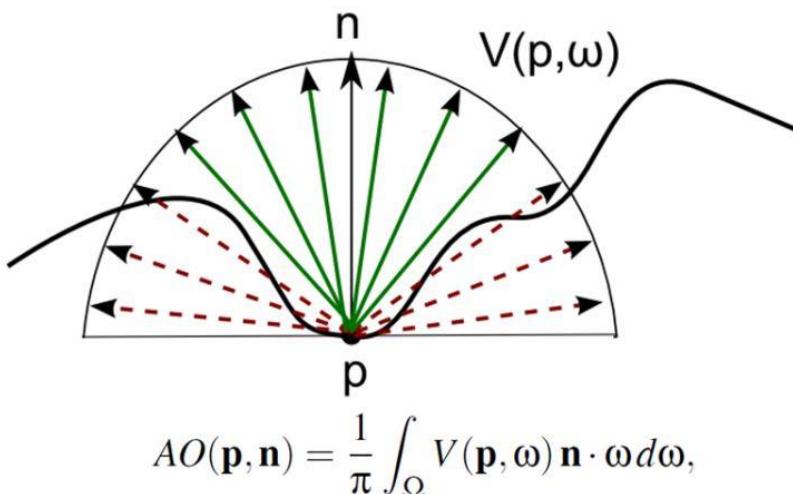
direct illumination



global illumination

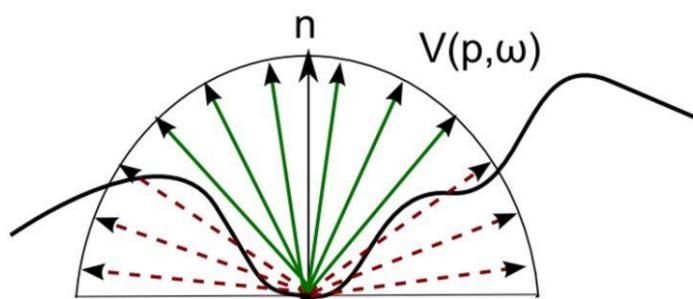
Ambient Occlusion (cont.)

- **Ambient occlusion (AO)** is a popular technique to approximate global illumination
 - Modulate ambient light by the surface's **accessibility**
 - Greatly enhance **depth perception** with a relatively low cost

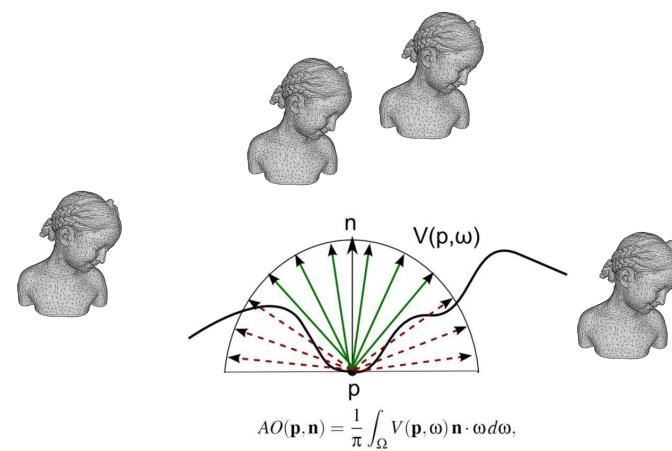


Ambient Occlusion (cont.)

- To compute AO, you need to know whether the ambient light is occluded in a direction
- In ray tracing, you can **trace rays** to determine the visibility
- For rasterization; however, this is difficult because **each polygon only knows its information** (again!)
 - Performance is also an issue!



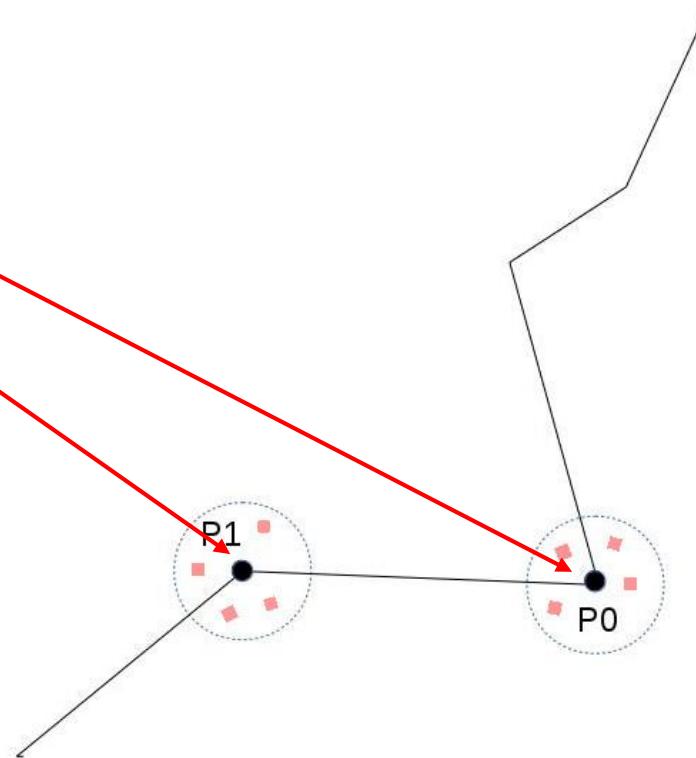
$$AO(\mathbf{p}, \mathbf{n}) = \frac{1}{\pi} \int_{\Omega} V(\mathbf{p}, \omega) \mathbf{n} \cdot \omega d\omega,$$



$$AO(\mathbf{p}, \mathbf{n}) = \frac{1}{\pi} \int_{\Omega} V(\mathbf{p}, \omega) \mathbf{n} \cdot \omega d\omega,$$

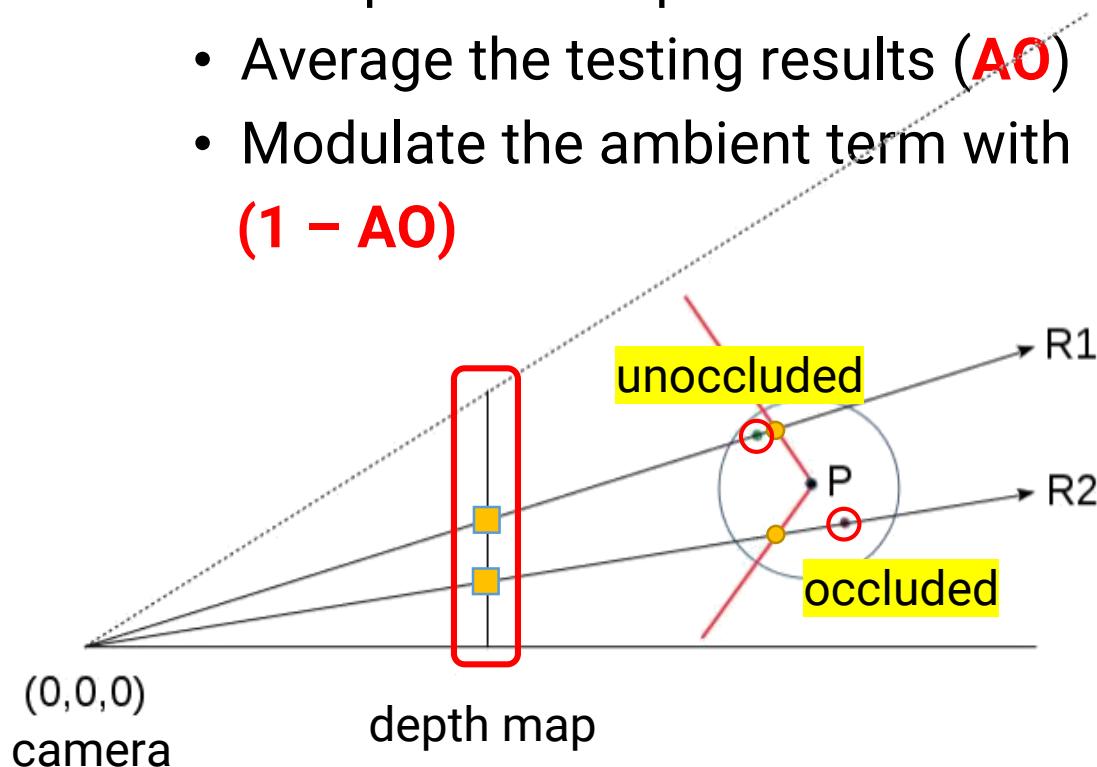
Screen-space Ambient Occlusion (cont.)

- Method
 - Generate samples within a sphere around the shading point (fragment)



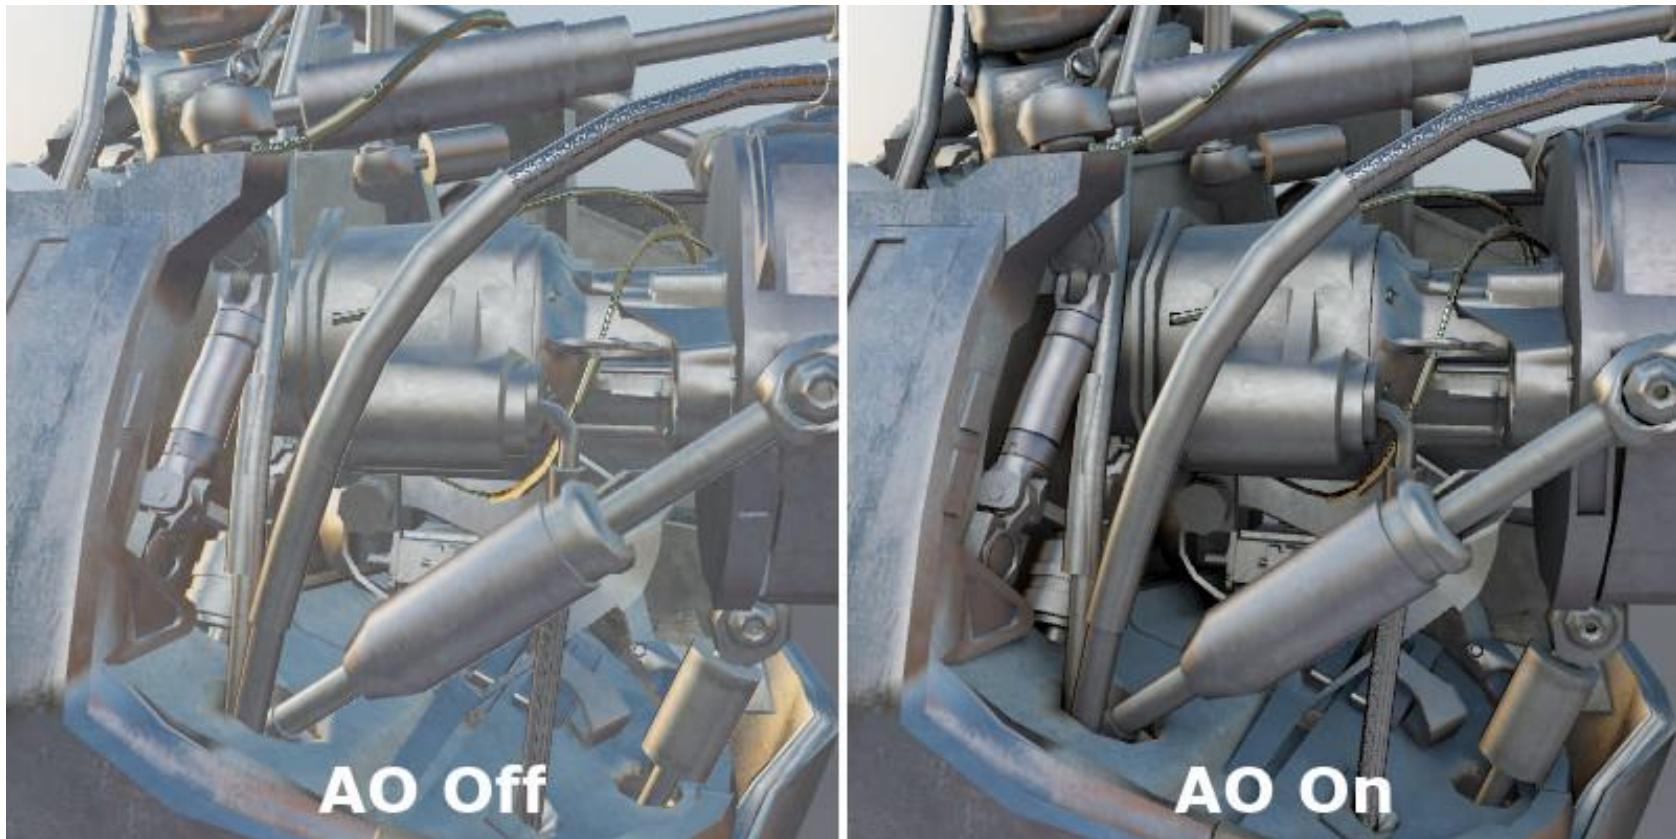
Screen-space Ambient Occlusion (cont.)

- Method
 - Project the samples back to the depth map from the camera
 - Compare the depth values
 - Average the testing results (**AO**)
 - Modulate the ambient term with **(1 – AO)**



Exampled Result

- Ambient occlusion (AO) is a popular technique to approximate global illumination



Global Illumination

Global Illumination

global illumination =

direct illumination + **indirect illumination**

local illumination + shadow map

difficult

constant ambient term

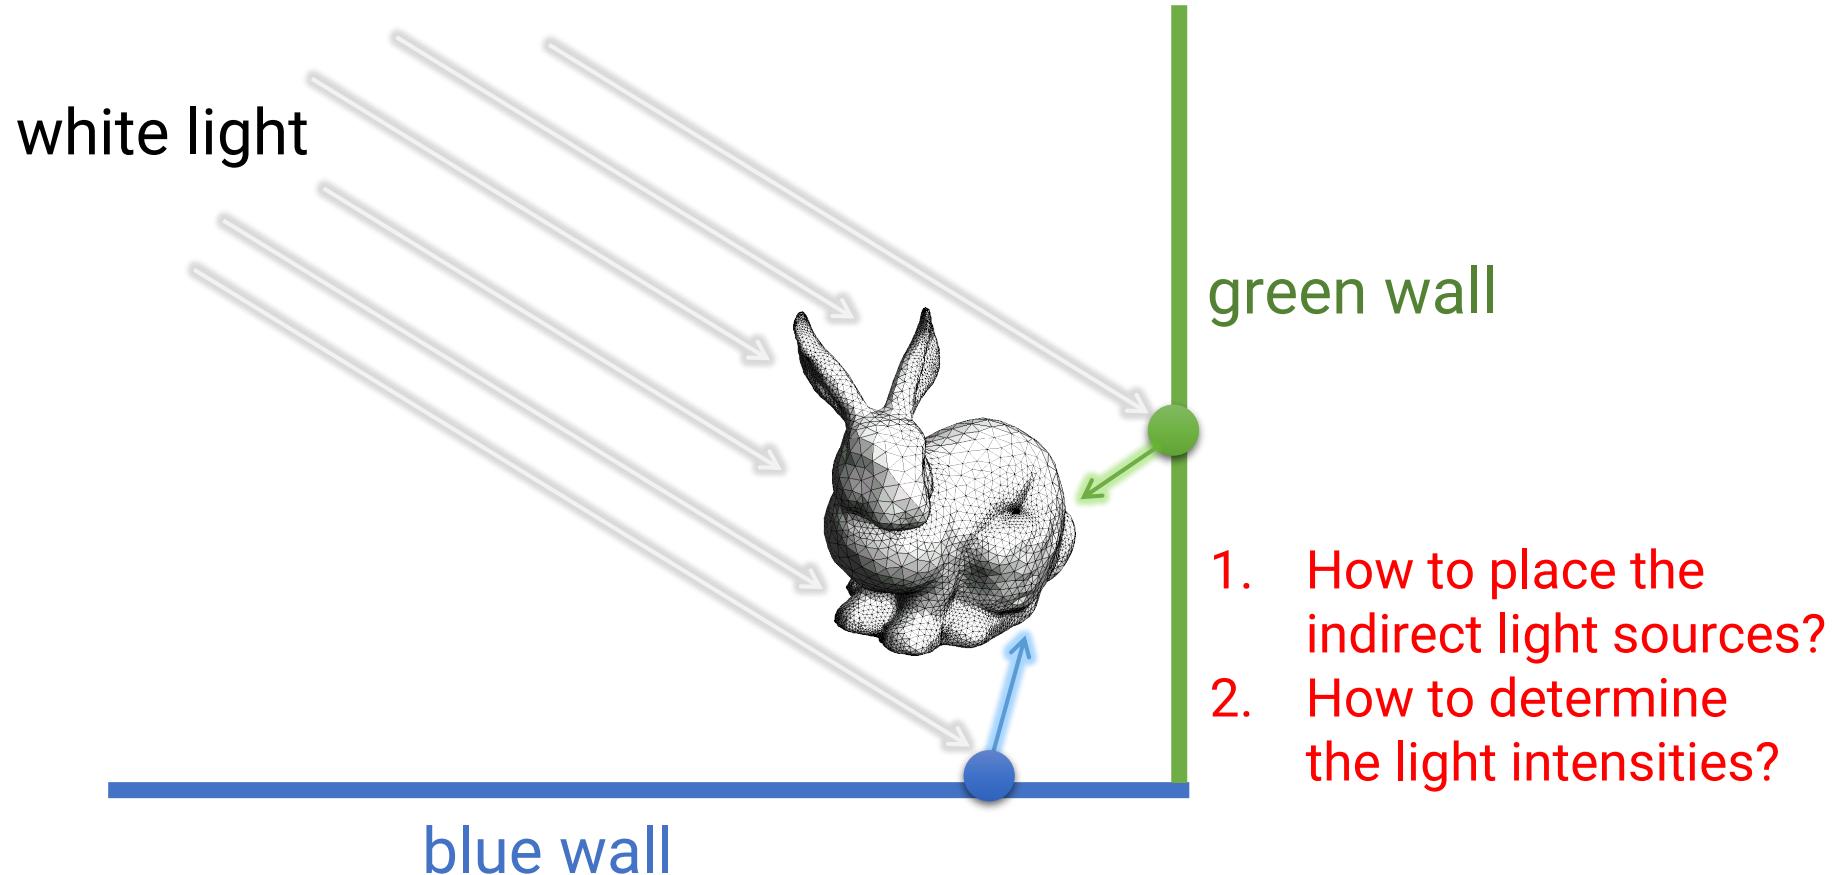
ambient occlusion

both are not good enough

Global Illumination (cont.)



Global Illumination (cont.)

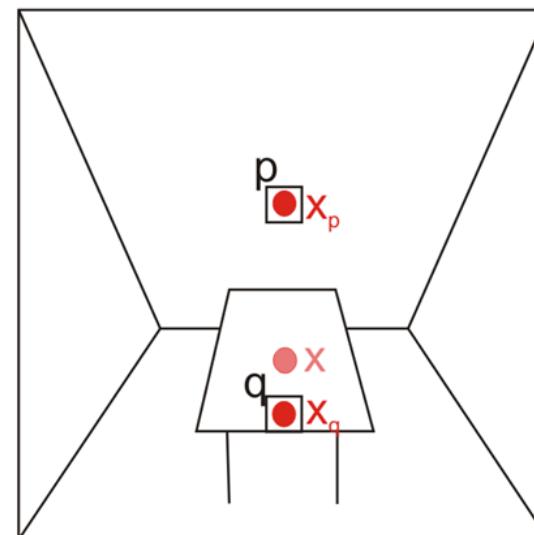
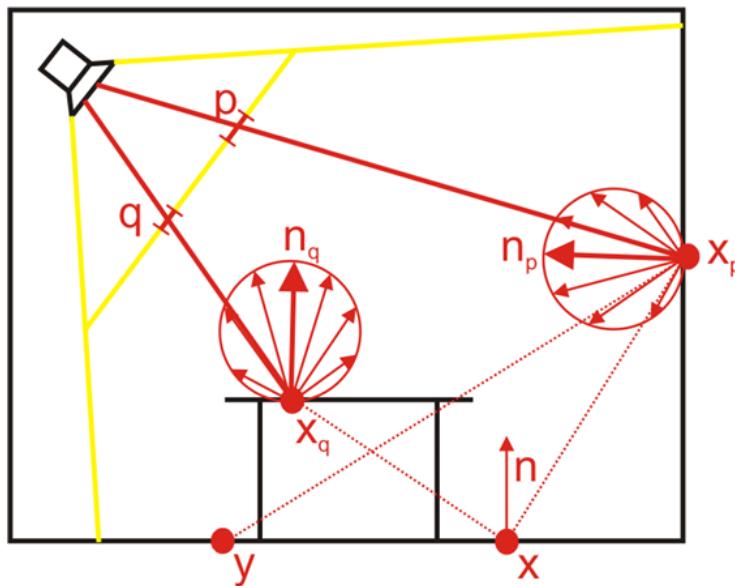


Global Illumination (cont.)

- Indirect illumination is especially difficult for rasterization because ...
 - **Each polygon only has its own information**
 - **It does not know which triangle will cast lighting on it**
- In the last two decades, hundreds of research papers focus on this topic to approximate visually-pleasing global illumination in real-time

Reflective Shadow Map

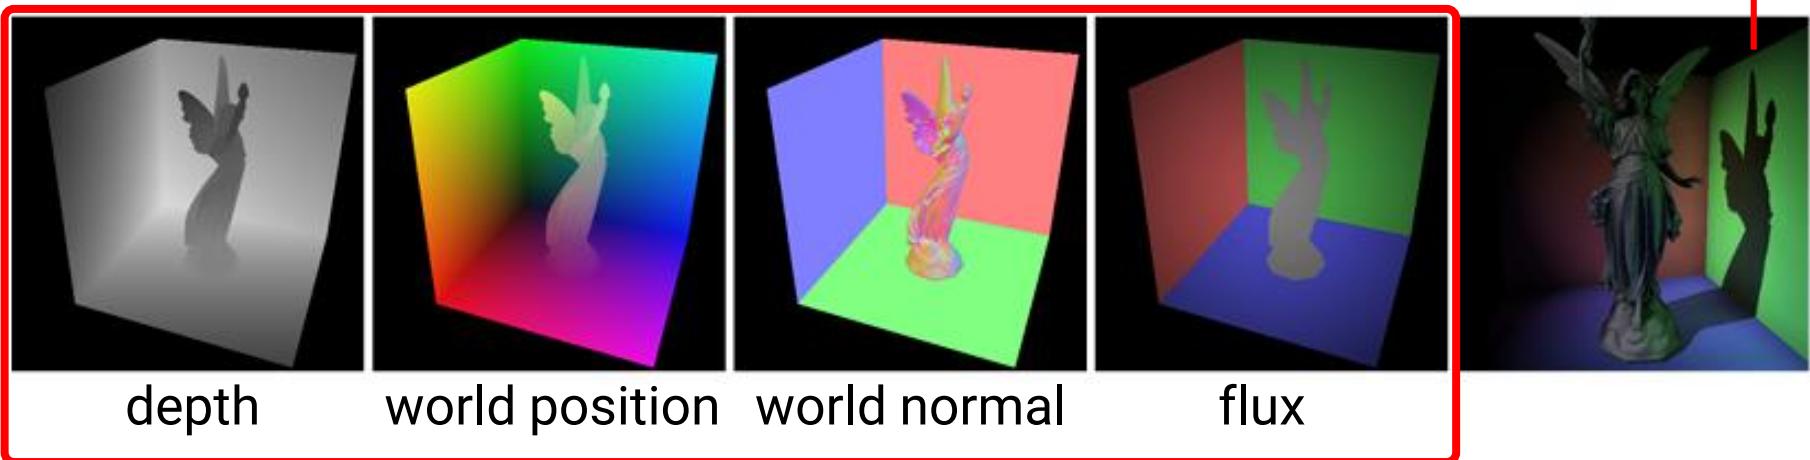
- Proposed by Dachsbacher and Stamminger, I3D 2005
- **Major idea**
 - The closest surfaces from the light can receive the lighting contribution
 - They become the indirect light sources



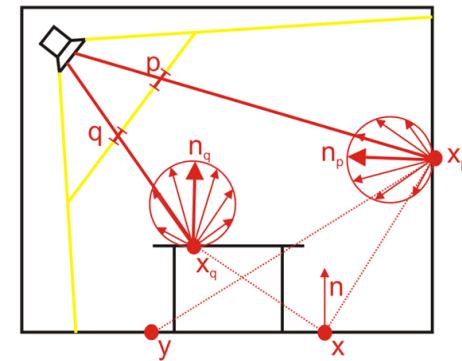
Reflective Shadow Map (cont.)

- Two-pass rendering algorithm

Pass II: render from the camera view



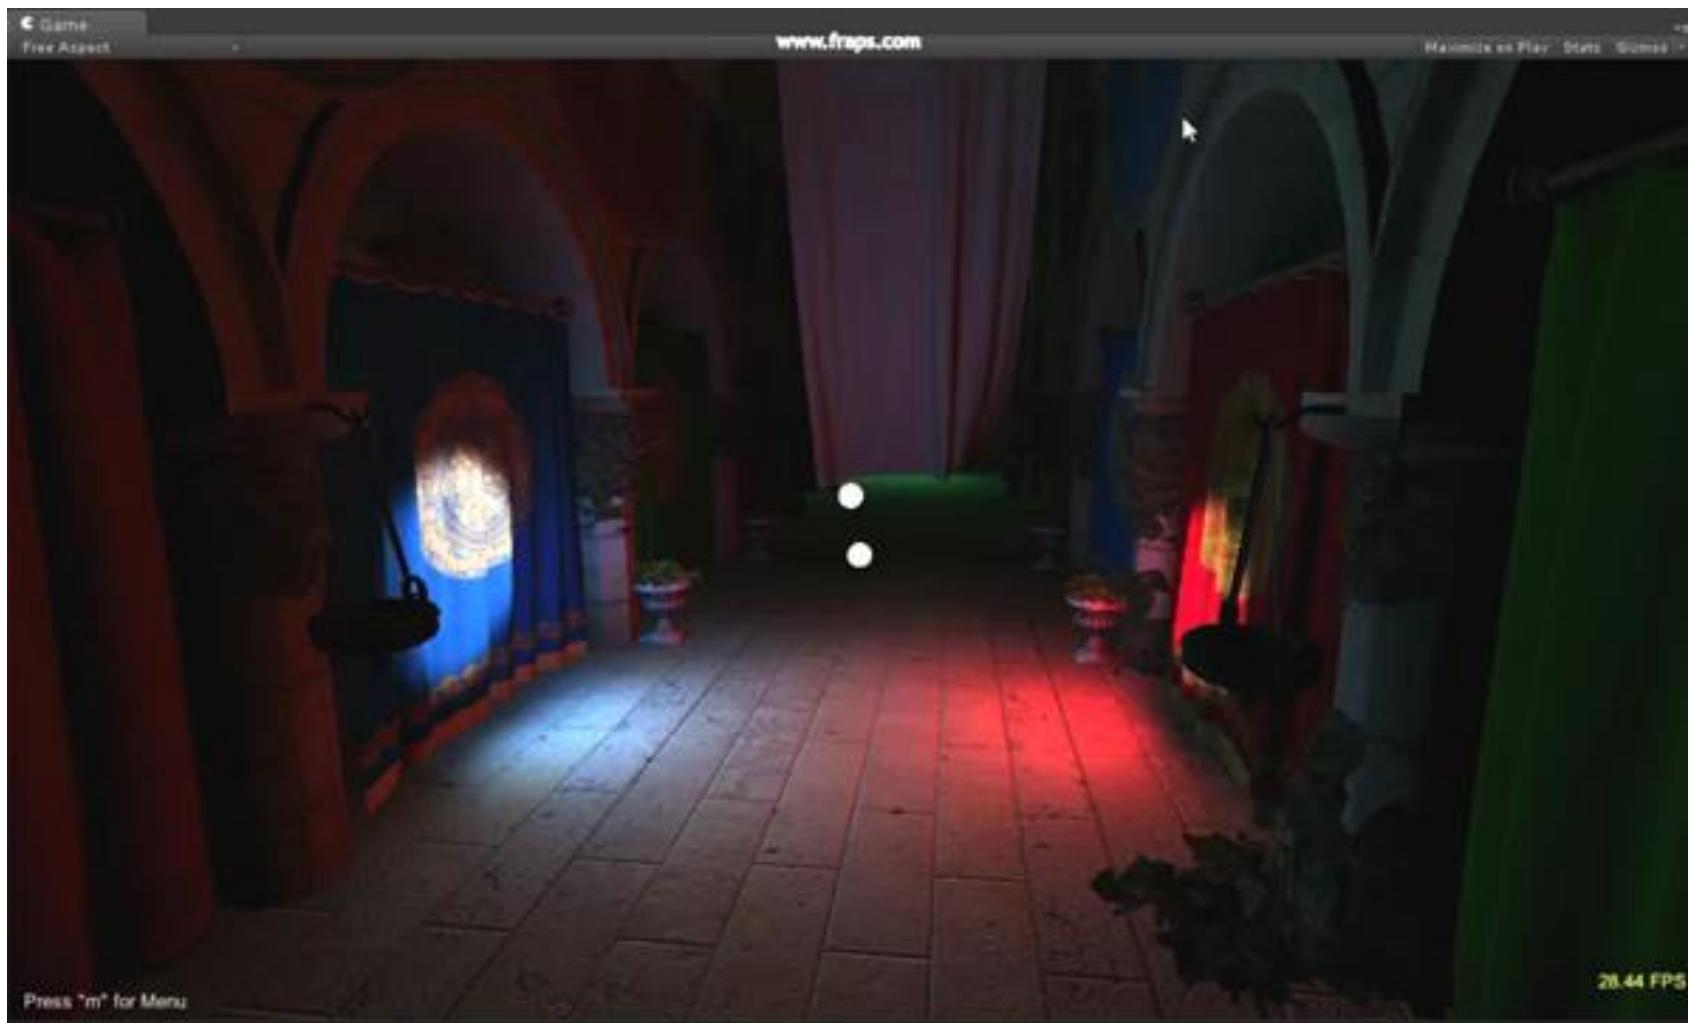
**Pass I: render G-buffer from a light view
(called RSM)**



Reflective Shadow Map (cont.)

- **Pass I:** rendering G-buffer (called **RSM**) from the light view for generating indirect light sources
 - World-space position
 - World-space normal
 - Reflected flux
 - The intensity of the primary light source multiplied by the reflectance of the surface
- **Pass II:** rendering from the camera view
 - Direct lighting is computed by local illumination and shadow mapping
 - Indirect lighting is estimated from the RSM

Reflective Shadow Map (cont.)

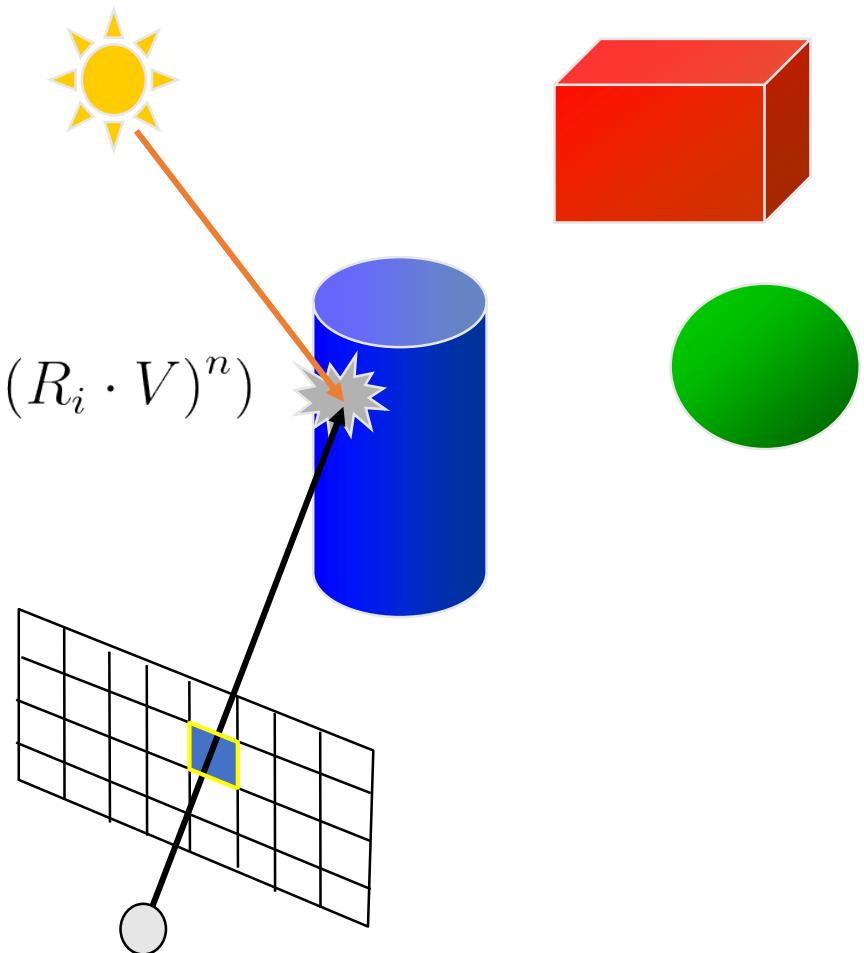


Ray Tracing

Ray Casting

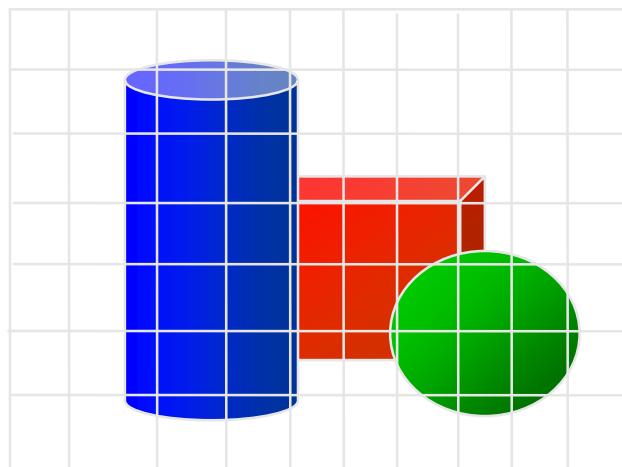
- Proposed by Appel [1968]

$$K_a I_a + \sum_{i=1}^{nls} I_i (K_d (L_i \cdot N) + K_s (R_i \cdot V)^n)$$

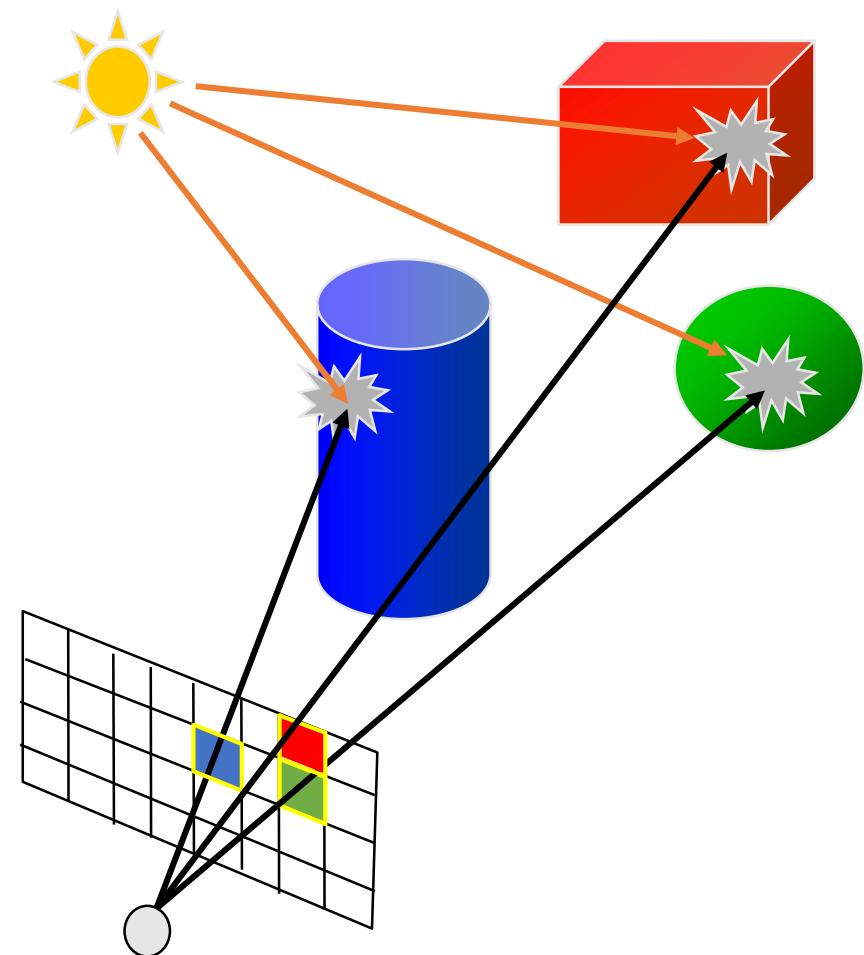


Ray Casting (cont.)

- Proposed by Appel [1968]

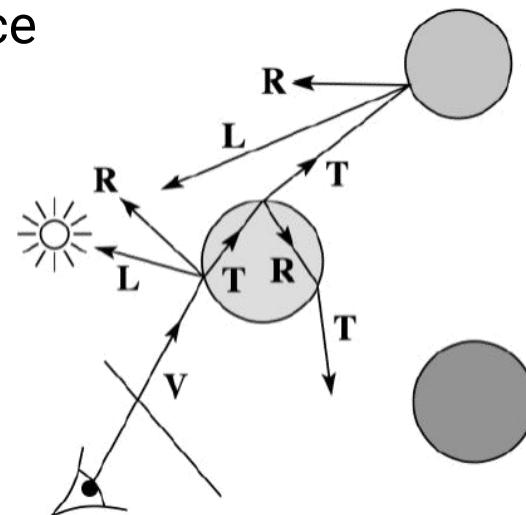


local illumination

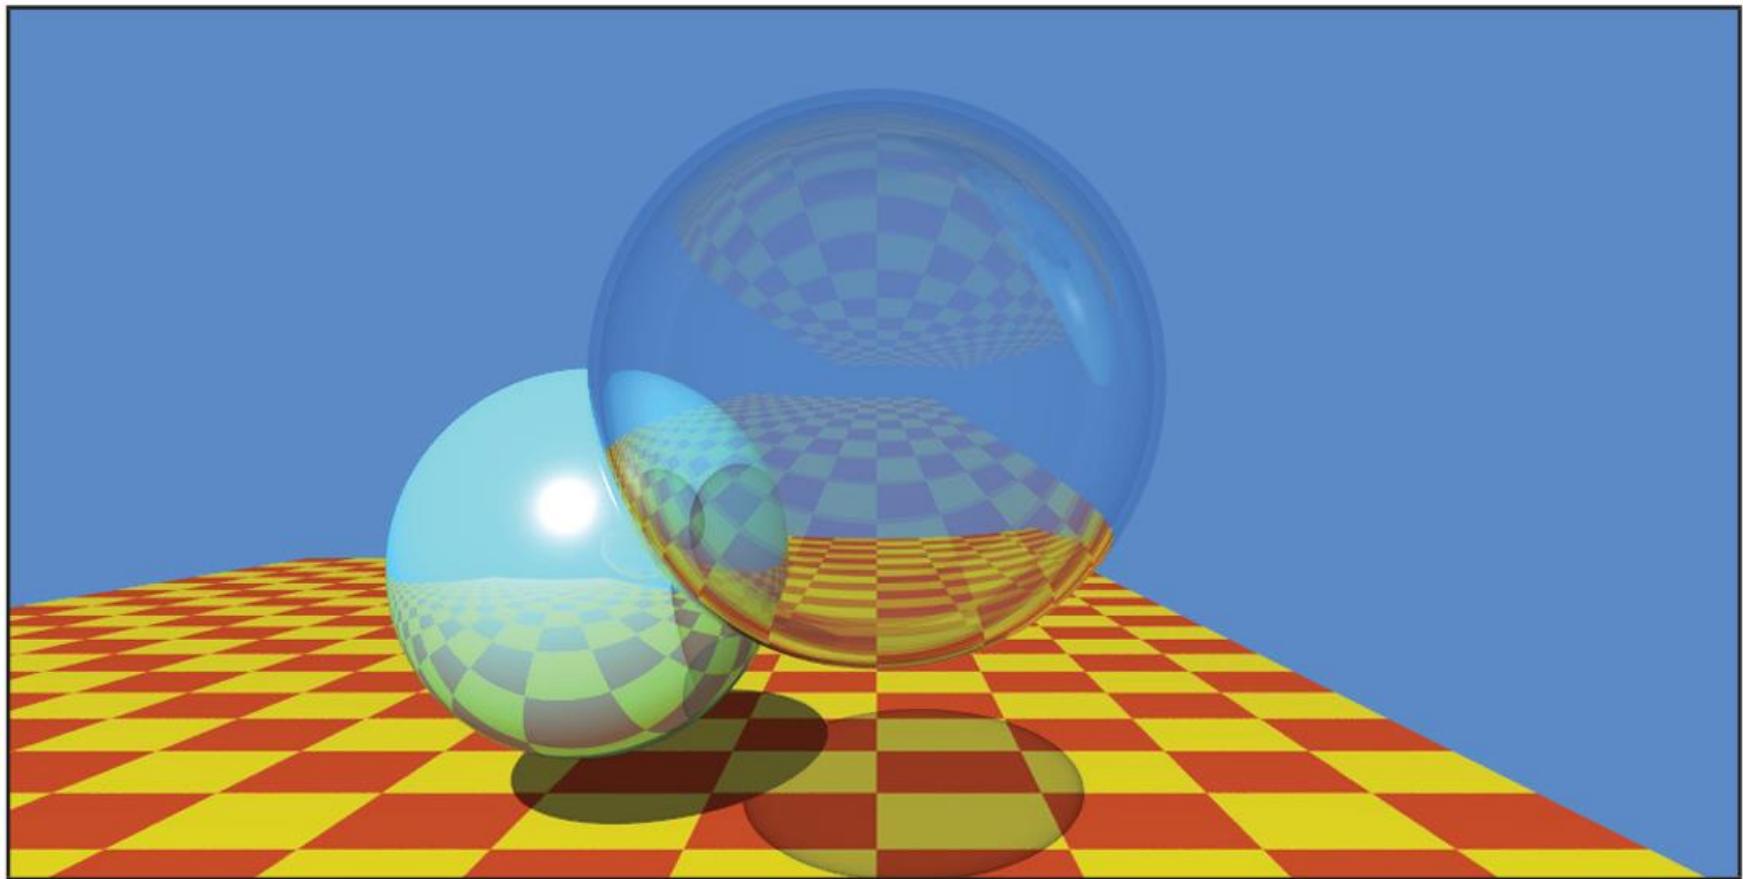


Whitted Ray Tracing

- Proposed by Whitted, 1980
- **Recursive** trace rays for **shadows**, perfect **specular** (e.g., mirror), and perfect **transparent** (e.g., glass) objects
 - For each pixel, trace a primary ray in the direction V to the first visible surface
 - For each intersection, trace secondary rays including
 - Shadow rays (L) to each light source
 - Reflected ray (R)
 - Refracted ray (T)

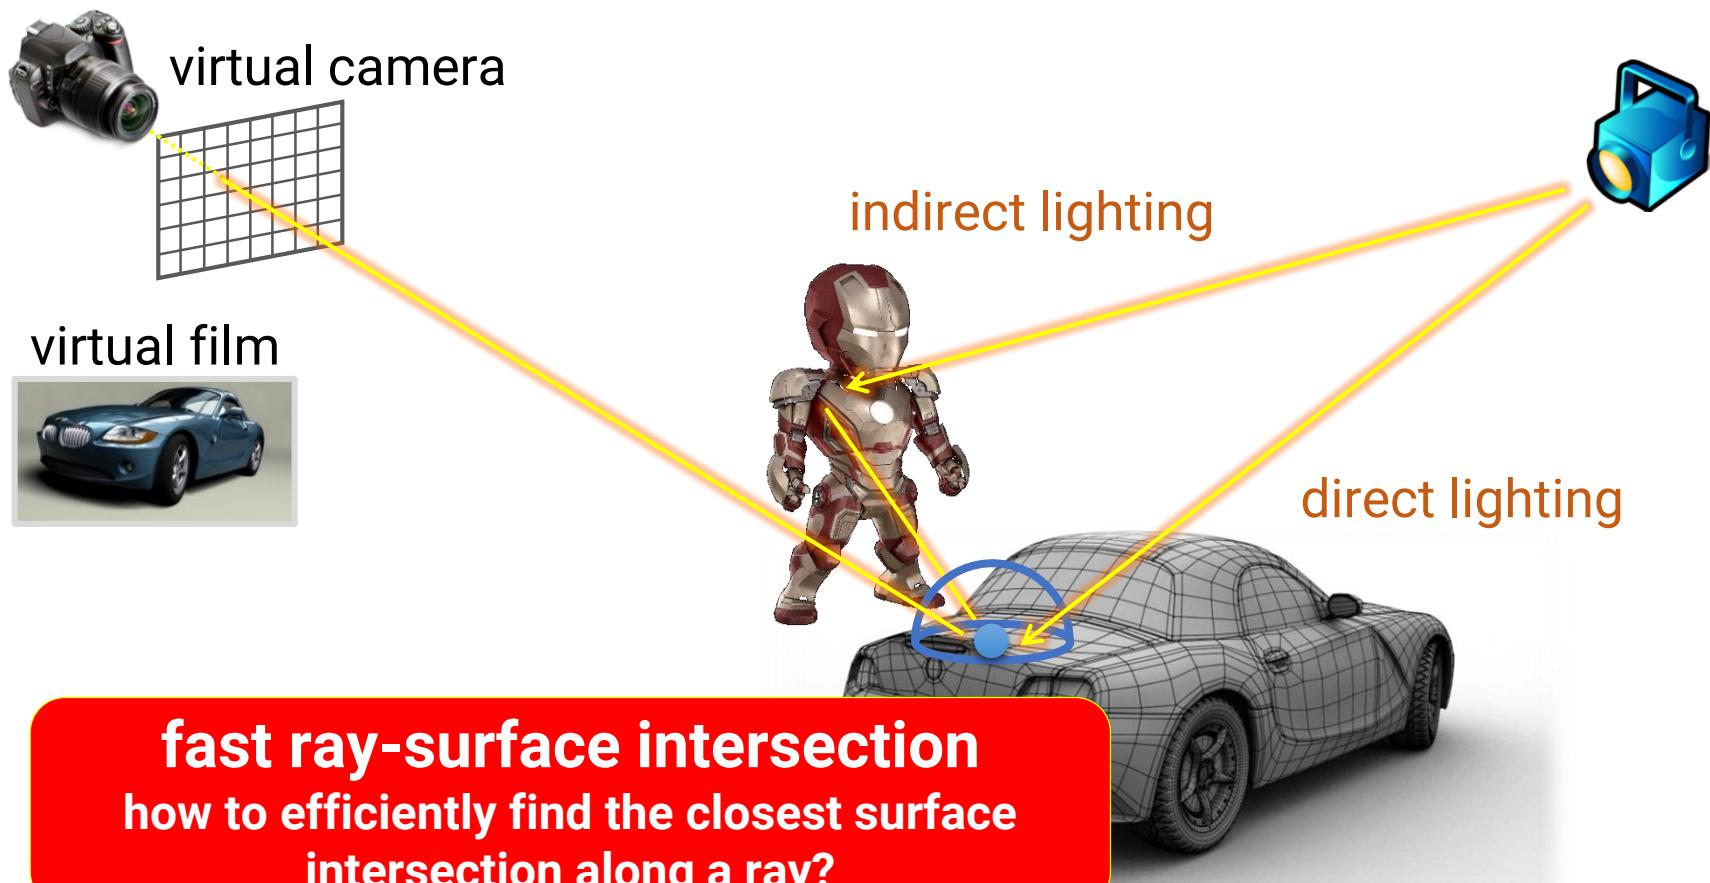


Whitted Ray Tracing (cont.)



Key: Fast Ray-Surface Intersection

- A unified approach for different light transport paths

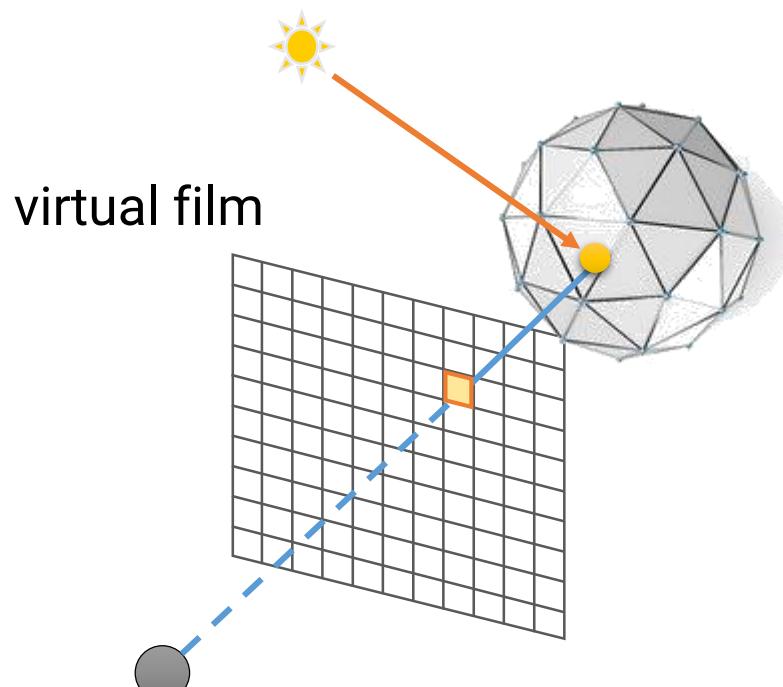


Acceleration Structure

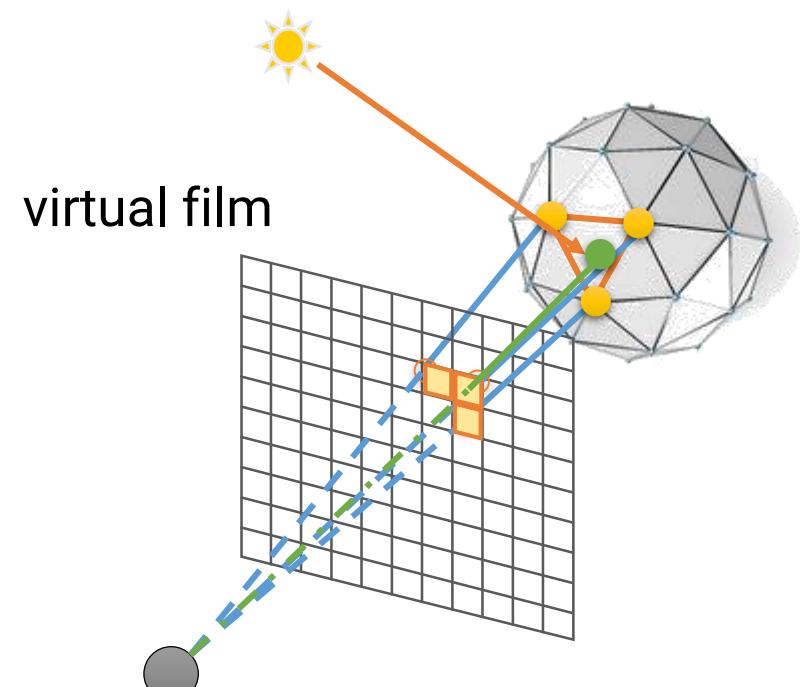
- **Reduce the required number of ray-surface intersection**
- Common acceleration structures
 - Bounding volume hierarchy (BVH)
 - Space subdivision
 - Uniform grid
 - Octree tree
 - Kd tree
- Please go the slides for details and examples

Ray Tracing v.s. Rasterization

Ray tracing

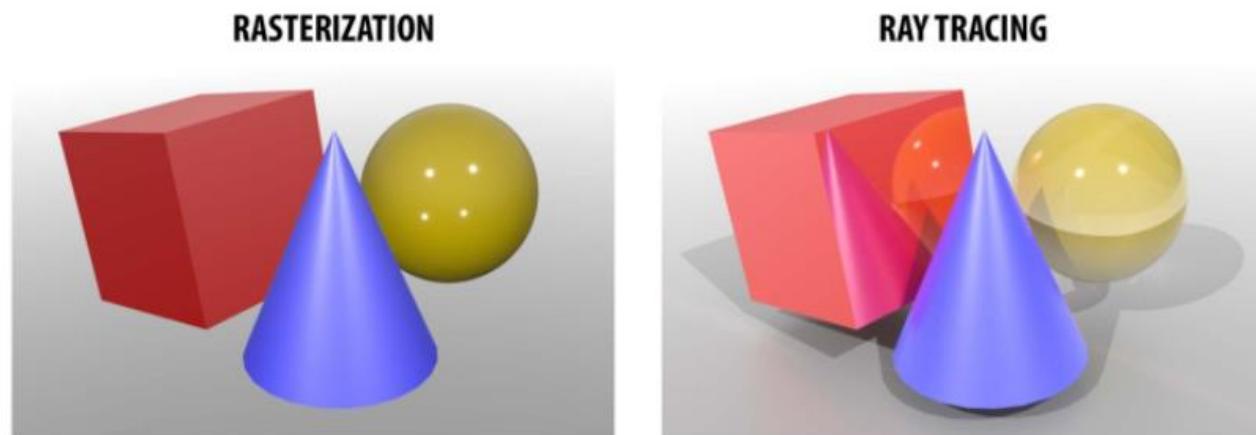


Rasterization



Ray Tracing v.s. Rasterization (cont.)

- Rasterization is **more friendly to hardware** and usually has higher parallelism
- But when we need to **interact with other triangles**, it is much more difficult to simulate effects such as reflection, refraction, shadows, and global illumination
 - Need specialized algorithms



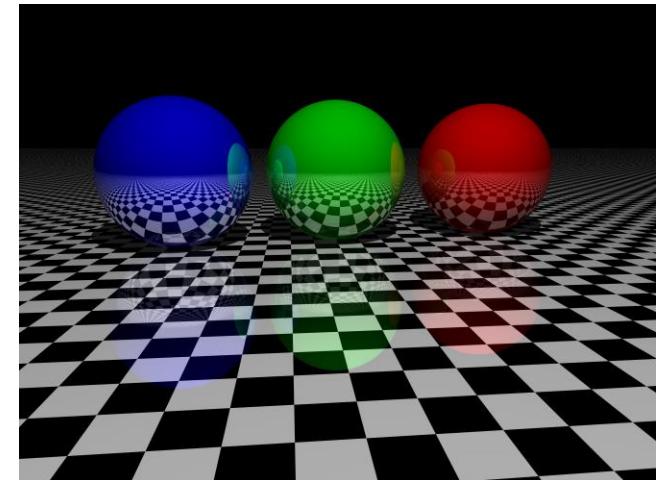
Ray Tracing v.s. Rasterization (cont.)

- Transparency
 - Rasterization
 - Render the object in order (distant objects first) and blend with the previous result in the color buffer
 - Ray-tracing:
 - Trace a secondary (refracted) ray through the object's surface



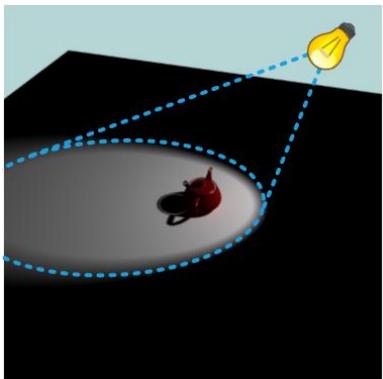
Ray Tracing v.s. Rasterization (cont.)

- **Reflection**
 - **Rasterization**
 - Render the scene into an environment map
 - Look up the environment map in the fragment shader
 - **Ray-tracing:**
 - Trace a secondary (reflected) ray from the object's surface



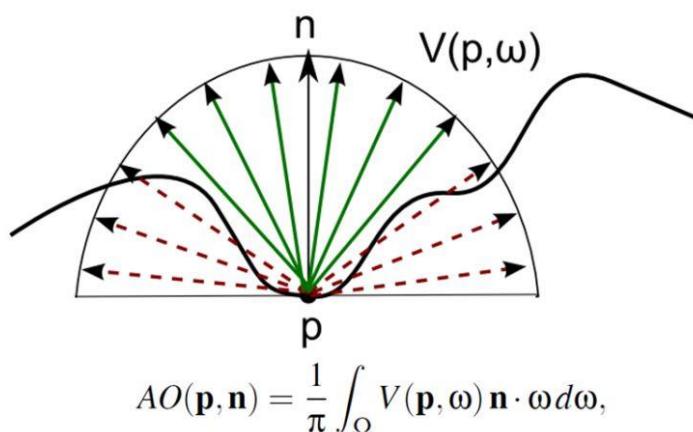
Ray Tracing v.s. Rasterization (cont.)

- **Shadow**
 - **Rasterization (shadow map as an example)**
 - Render a **shadow map** to record the closest surface from each light
 - Look up the map to determine whether a surface point is in shadow or not in the second pass
 - **Ray-tracing**
 - Trace a shadow ray to see if the lighting direction is occluded



Ray Tracing v.s. Rasterization (cont.)

- Ambient occlusion
 - Rasterization (SSAO as an example)
 - Use the **depth map** to find nearby occluders in screen space
 - Ray-tracing
 - Trace shadow rays to see if a direction is occluded

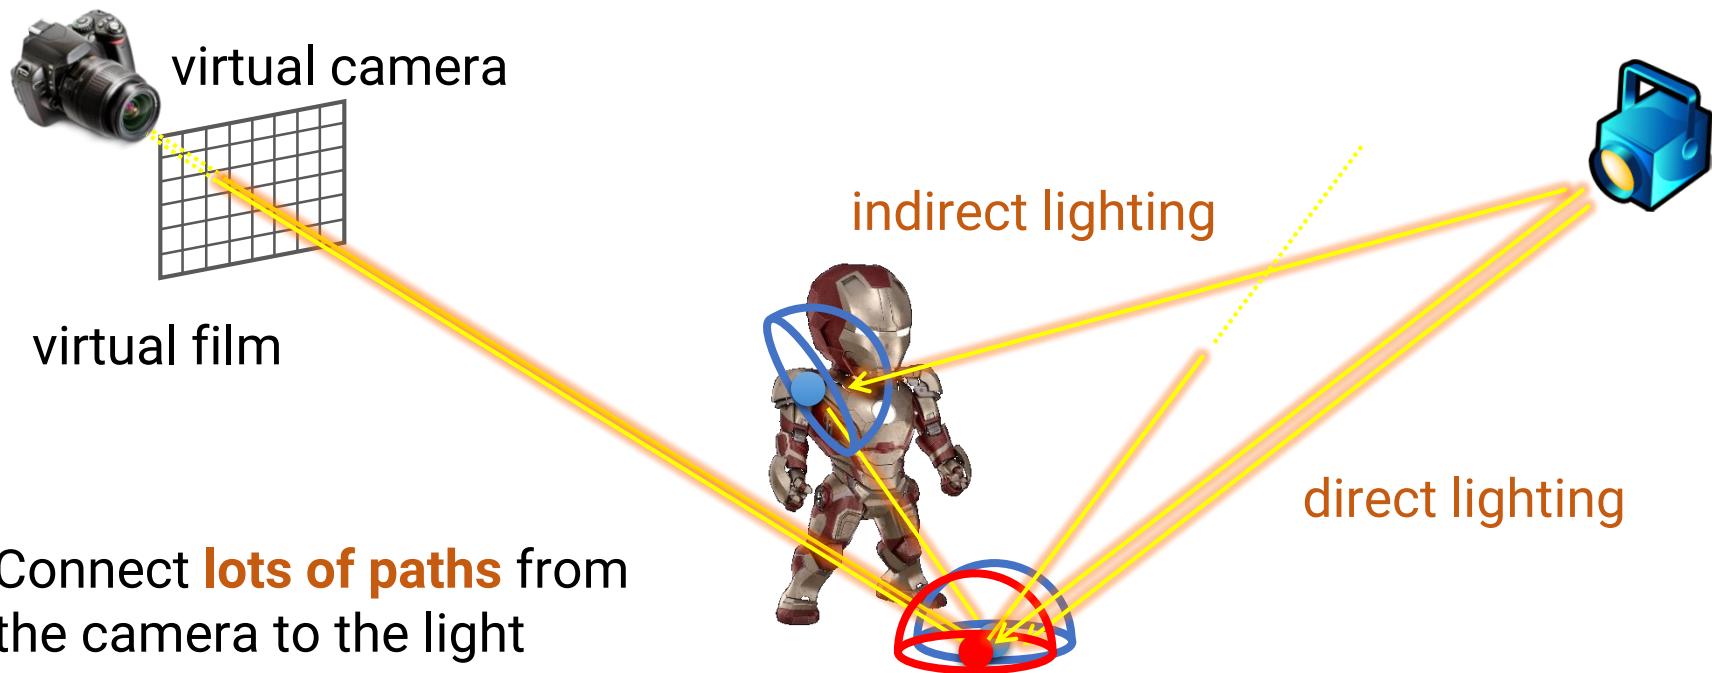


Ray Tracing v.s. Rasterization (cont.)

- Global illumination

- Rasterization (RSM as an example)

- Render the direct lighting result from the light view
 - Use the results in the first pass to render indirect lighting



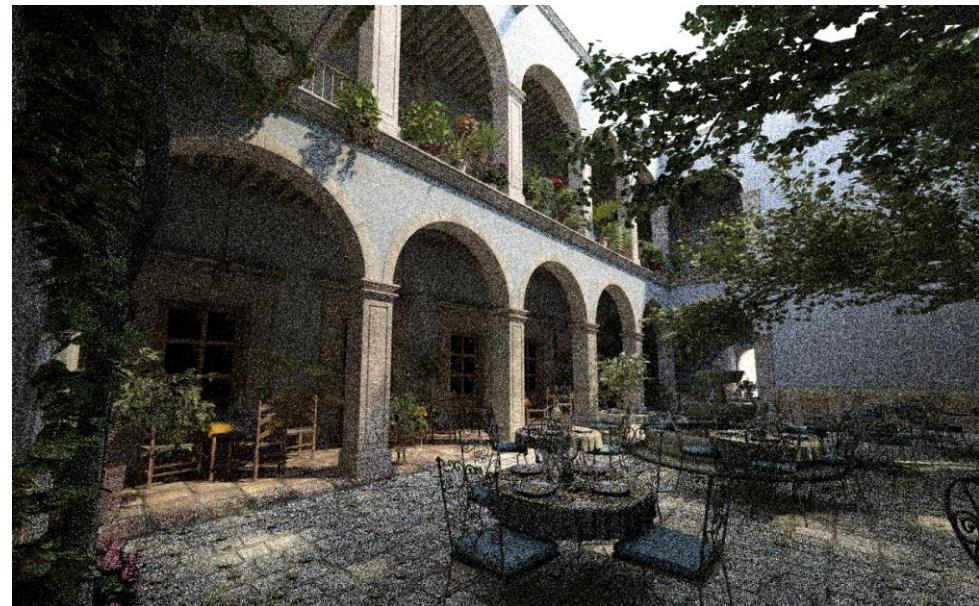
Ray Tracing v.s. Rasterization

- **Problems with ray tracing**

- Ray tracing is more general
- However, its simulator usually has a slow convergence rate and produces lots of noise when samples are not enough

- **Solution**

- More rays (expensive)
- Filtering (bias)



Any Questions?

Good Luck!

