



Bitmapped Images

Multimedia Techniques & Applications

Yu-Ting Wu



Outline

- Overview
- Image compression
- Image manipulation
- Image scaling

Outline

- **Overview**
- Image compression
- Image manipulation
- Image scaling

Recap: Two Approaches for Graphical Modeling



Overview

- Record the value of every pixel in the image
- **Image size** is the main cost for the simplicity
- Images created from external devices are usually in a bitmapped fashion
 - Digital cameras
 - Scanners



digital camera




scanner

Resolution

- A measure of **how finely** a device approximates continuous images using finite pixels
 - Closely related to sampling rates
- Two ways of specifying resolution
 - **Printers and scanners: number of dots per unit**
 - Dots per inch (dpi)
 - E.g., consumer printer (600 dpi), book production (1200 - 2700 dpi), scanners (300 dpi - 3600 dpi)
 - **Video: the size of a frame measured in pixels**
 - E.g., 640 x 480, 768 x 576
 - Can translate into the form of dpi if you know the physical dimension of the display device

Raw Size of Bitmapped Images



128 128 128 **setrgbcolor**

0 0 128 128 **rectfill**

0 255 0 **setrgbcolor**

32 32 64 64 **rectfill**

**Four lines for the
vector graphics format**

128 x 128 x 3 = 49152 bytes for the bitmapped image format

File Formats of Bitmapped Images

- Related to the way of compressing data
 - **Lossless compression**
 - GIF (Graphics Interchange Format)
 - PNG (Portable Network Graphics)
 - BMP (Windows Bitmap)
 - TIFF (Tagged Image File Format)
 - TGA (Truevision TGA, TARGA)
 - **Lossy compression**
 - JPEG (Joint Photographic Experts Group)
 - TIFF (Tagged Image File Format)

Outline

- Overview
- **Image compression**
- Image manipulation
- Image scaling


Image Compression

- **Motivation**

- Faithfully storing all pixel values of an image takes lots of memory space
- Human eyes can tolerate some minor errors in images


- **Observation**

- Images are usually **smooth** and have some **spatial coherence**




Compression Methods

- Spend some computation efforts to earn saving in space
- The effectiveness depends on the content of the compressed image
 - **Image size can become bigger after applying compression**
 - Definitely true, otherwise, any data can be compressed into one byte



128 bytes → 6 bytes
for a row (RLE)




128 bytes → 256 bytes
for a row (RLE)

Compression Methods (cont.)


- **Lossless compression**

- No information will lose during a compression/decompression cycle
- E.g., run-length encoding (RLE), variable-length coding




- **Lossy compression**

- Discard some information during the compression process and the information can **never** be recovered
- E.g., JPEG



Run-Length Encoding (RLE)

- The simplest compression technique
- Each stored value is followed by a **count** to indicate the number of consecutive pixels of that value
- Example
 - RLE for row1: *128 128*
 - *Originally, we need 128 bytes assuming 8-bit intensity is used*
 - RLE for row2: *128 32 0 64 128 32*










Huffman Coding

- The best-known variable length coding
- **Lossless** compression
- Example:

Assume an 8 x 8 image containing 6 different pixel intensities

We can count their occurrence:

naïve encoding	000	001	010	011	100	101
						
intensity	20	60	100	140	180	220
occurrence	5	6	25	16	9	3




Huffman Coding (cont.)


- **Algorithm**







- Build a **Huffman tree**

- Sort the occurrence of intensity




- Merge the two with the smallest occurrence, and sort again



					
20	60	100	140	180	220
5	6	25	16	9	3


Huffman Coding (cont.)

- **Algorithm**
 - Build a **Huffman tree**
 - Keep doing ...



Huffman Coding (cont.)


- **Algorithm**
 - Build a **Huffman tree**
 - Keep doing ...



Huffman Coding (cont.)

- **Algorithm**

- Build a **Huffman tree**
 - Keep doing ...










Huffman Coding (cont.)

• Algorithm

- Build a **Huffman tree**
 - Keep doing ...

Once we have done this, the **leaf nodes** are the initial data items


					
20	60	100	140	180	220
5	6	25	16	9	3









Huffman Coding (cont.)

• Algorithm

- Build a **Huffman tree**
- Label the code (from root)




					
20	60	100	140	180	220
5	6	25	16	9	3

Huffman Coding (cont.)

• Algorithm

- Build a **Huffman tree**
- Label the code (from root)
- We will obtain

20: 11111	60: 1110
100: 0	140: 10
180: 110	220: 11110




The average number of bits per data:

$$1 \cdot (25/64) + 2 \cdot (16/64) + 3 \cdot (9/64) + 4 \cdot (6/64) + 5 \cdot (3/64) + 5 \cdot (5/64) = 2.31 \text{ (compared to raw: 8 bit, and naïve encoding: 3 bit)}$$

JPEG Compression

- JPEG is the most important **lossy** compression technique, which stands for ***Joint Photographic Experts Group***
 - Related file formats: *.jpg / *.jpeg / *.jpe / *.jfif / *.jfi / *.jif
- It works because image data can tolerate a certain amount of data loss

JPEG Compression (cont.)




- **RGB → YCbCr**


- People are more sensitive to intensity (Y) and less sensitive to color (Cb, Cr)
- Cb and Cr are lower frequency and have more spatial coherence
- Compress Cb and Cr; while keeping Y as it is




JPEG Compression (cont.)



- **Divide into 8x8 blocks (for Cb & Cr)**
 - The entire image is too difficult to compress
 - Small image block has higher coherence



JPEG Compression (cont.)




- **Discrete Cosine Transform (DCT)**


- A method for transforming an image into its frequency domain
- The DCT of an image block is the coefficients of different cosines of the image block



image




DFT result



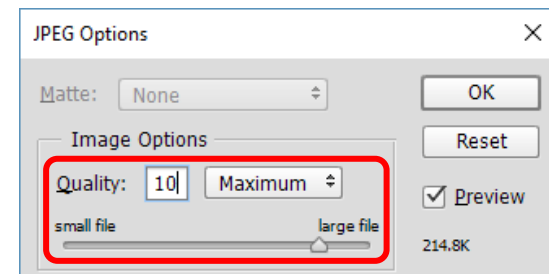
DCT result

JPEG Compression (cont.)




- **Quantization**

- Humans are less sensitive to the high-frequency signal
- Use fewer bits for high-frequency signals in the DCT result and vice versa
- **This step is the reason for lossy compression**
- After this step, many components will end up with zero coefficients





JPEG Compression (cont.)




- **Zigzag ordering**

- For later Huffman encoding
- Result in a longer zero sequence





JPEG Compression (cont.)



- **RLE / Huffman encoding**

- Different strategy for DC and AC term
- The DC components of different blocks are encoded using Huffman algorithm
- The AC components within a block are encoded using RLE



JPEG Compression (cont.)

- The decompression of JPEG data is done by reversing the compression process
- We can control the degree of compression by altering the amount of quantization
- JPEG compression usually achieves very high compression rate for natural images (5% of the original size)

JPEG Compression (cont.)



original image



lossy JPEG format
with "artifacts"

PNG



4.16MB

JPG



528KB

PNG stands for Portable Network Graphics (lossless compression)

Outline

- Overview
- Image compression
- **Image manipulation**
- Image scaling

Image Manipulation

- **Motivations**

- Correct deficiencies in an image (e.g., noise, red-eye)
- Create images that are difficult or impossible to make naturally (e.g., glow)

- Type of image manipulations

- Pixel **point** processing
- Pixel **group** processing

Pixel Point Processing

- Compute a pixel's new value solely on the basis of its old value

mapping function


$$p' = f(p)$$

- Some examples
 - Adjustment of brightness
 - Adjustment of contrast
 - Change the black/white levels




Color Curve

- The operations can be considered generally as altering the mapping function f
- **Color curve**



Color Curve (cont.)

- The operations can be considered generally as altering the mapping function f
- **Color curve**




Histogram

- An approximate representation of the **distribution** of numerical data




Intensity Image




Histogram Equalization

- Remap pixel values to produce a new distribution with higher contrast




Pixel Group Processing

- Compute each pixel's new value as a function not just of its old value, but also of **the values of neighboring pixels**
- Usually related to **filtering**
 - For a pixel of an image, specify a two-dimensional array of weights of its neighbors
 - Several types of filters
 - Smoothing
 - Sharpening
 - Detecting edge



Box Filter

- Each neighbor has the same weight



Box Filter (cont.)

convolution mask
convolution kernel

↓

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

filter

0	0	9	9
0	0	9	9
0	0	9	9
0	0	9	9

image (signal)

$$(1/9*0+1/9*0+1/9*0+1/9*0+1/9*0+1/9*0+1/9*0+1/9*9+1/9*9+1/9*9)$$

=

0	3	6	9
0	3	6	9
0	3	6	9
0	3	6	9


filtered image (signal)



Input $f(x,y)$

$$* \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} =$$


Convolution operator,
not multiplication!





Output $g(x,y)$

Box Filter (cont.)


- Visualization of kernel weight



box filter






gaussian filter



Gaussian Filter

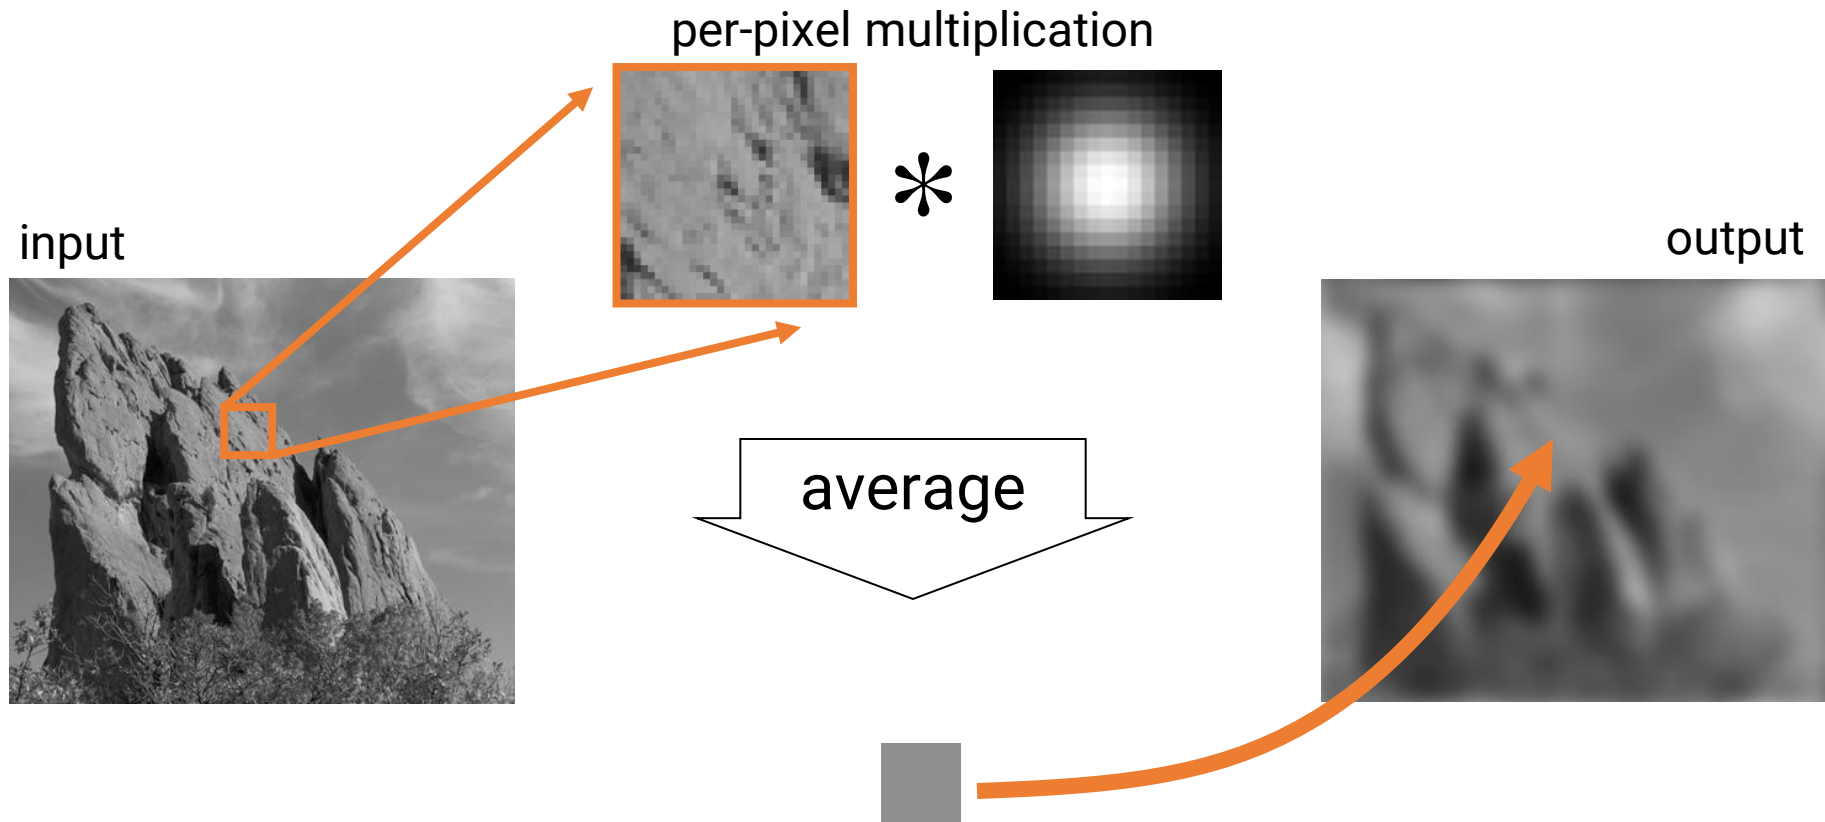
- The weight of neighbor falls exponentially with its distance to the filtered pixel
 - Standard deviation σ controls the speed of decreasing

$$g(x, y) = \underbrace{\frac{1}{2\pi\sigma^2}}_{\text{normalization (sum of weight = 1)}} e^{-\frac{\underbrace{x^2+y^2}_{\text{distance}}}{2\underbrace{\sigma^2}_{\text{standard deviation}}}}$$



$$\frac{1}{273} \begin{bmatrix} 1 & 4 & 7 & 4 & 1 \\ 4 & 16 & 26 & 16 & 4 \\ 7 & 26 & 41 & 26 & 7 \\ 4 & 16 & 26 & 16 & 4 \\ 1 & 4 & 7 & 4 & 1 \end{bmatrix}$$

Gaussian Filter (cont.)




Box Filter v.s. Gaussian Filter


original image



box filter



gaussian filter




Bilateral Filter

- Properties of Gaussian filter

$$g(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{\underbrace{2\sigma^2}_{\text{range of the filter}}}}$$




input




small σ



limited smoothing



large σ



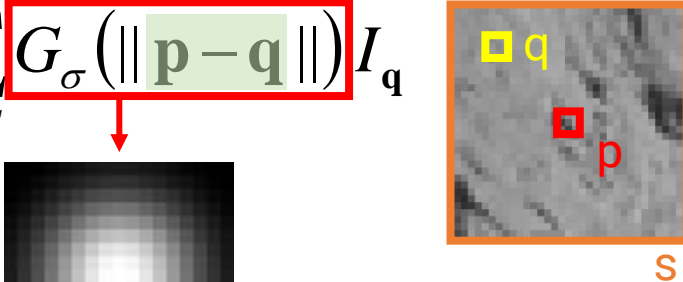
strong smoothing

Bilateral Filter (cont.)

- Problems of Gaussian filter
 - Does smooth images
 - But smooths too much:
edges are blurred
 - Only spatial distance matters
 - No edge term

$$GB[I]_p = \sum_{q \in S} G_{\sigma}(\| \mathbf{p} - \mathbf{q} \|) I_q$$

spatial distance



$$g(x, y) = \frac{1}{2\pi\sigma^2} e^{-\frac{x^2+y^2}{2\sigma^2}}$$

input




output



Bilateral Filter (cont.)

- Problems of Gaussian filter
 - Same Gaussian kernel everywhere



Bilateral Filter (cont.)

- Combine another Gaussian weight computed by **intensity difference** (edge preserving)

$$GB[I]_p = \sum_{q \in S} G_{\sigma} \left(\overset{\text{spatial distance}}{\| \mathbf{p} - \mathbf{q} \|} \right) I_q$$




$$BF[I]_p = \frac{1}{W_p} \sum_{q \in S} G_{\sigma_s} \left(\overset{\text{spatial distance}}{\| \mathbf{p} - \mathbf{q} \|} \right) G_{\sigma_r} \left(\overset{\text{range (intensity) distance}}{\| I_p - I_q \|} \right) I_q$$

NEW!

Bilateral Filter (cont.)


- Visualization

$$\underbrace{BF[I]_p}_{\text{Result}} = \frac{1}{W_p} \sum_{q \in S} \underbrace{G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|)}_{\text{Spatial}} \underbrace{G_{\sigma_r}(\|I_p - I_q\|)}_{\text{Range}} \underbrace{I_q}_{\text{Input}}$$



Bilateral Filter (cont.)

- Parameters

$$BF[I]_{\mathbf{p}} = \frac{1}{W_{\mathbf{p}}} \sum_{\mathbf{q} \in \mathcal{S}} G_{\sigma_s}(\|\mathbf{p} - \mathbf{q}\|) G_{\sigma_r}(\|I_{\mathbf{p}} - I_{\mathbf{q}}\|) I_{\mathbf{q}}$$


- Spatial sigma σ_s** : the spatial extent of the kernel, size of the considered neighborhood
- Range sigma σ_r** : how the weight decreases with respect to the range difference

Bilateral Filter (cont.)

- Parameters

$$\sigma_s = 2$$



$$\sigma_r = 0.1$$



$$\sigma_r = 0.25$$



$$\sigma_r = \infty$$

(Gaussian blur)





input

$$\sigma_s = 6$$




$$\sigma_s = 18$$



Bilateral Filter Application (cont.)

- Denoising




noisy input



Median Filter 5x5

Bilateral Filter Application (cont.)

- Denoising



noisy input



Bilateral Filter 7x7

Bilateral Filter Application (cont.)

- Denoising



noisy input



Gaussian Filter



Bilateral Filter

Sobel Filter


- Negative weights are commonly used for edge detection

-1	0	+1
-2	0	+2
-1	0	+1

Gx


+1	+2	+1
0	0	0
-1	-2	-1

Gy



Photographic Style Transfer

- Two-scale Tone Management for Photographic Look, Bae et al. SIGGRAPH 2006



Photographic Style Transfer (cont.)

- Motivation



Ansel Adams
Clearing Winter Storm




an amateur photographer

Photographic Style Transfer (cont.)

- Motivation



Ansel Adams
Clearing Winter Storm



style transferred result
(an imitation of Ansel Adams)

Photographic Style Transfer (cont.)

- **Observation**

- Different photographers have different tonal styles
 - Global contrast



Ansel Adams




Kenro Izu

Photographic Style Transfer (cont.)

- **Observation**

- Different photographers have different tonal styles
 - Global contrast
 - Local contrast



Ansel Adams



Kenro Izu

Photographic Style Transfer (cont.)

- **Goal**
 - Transfer look between photographers



input image



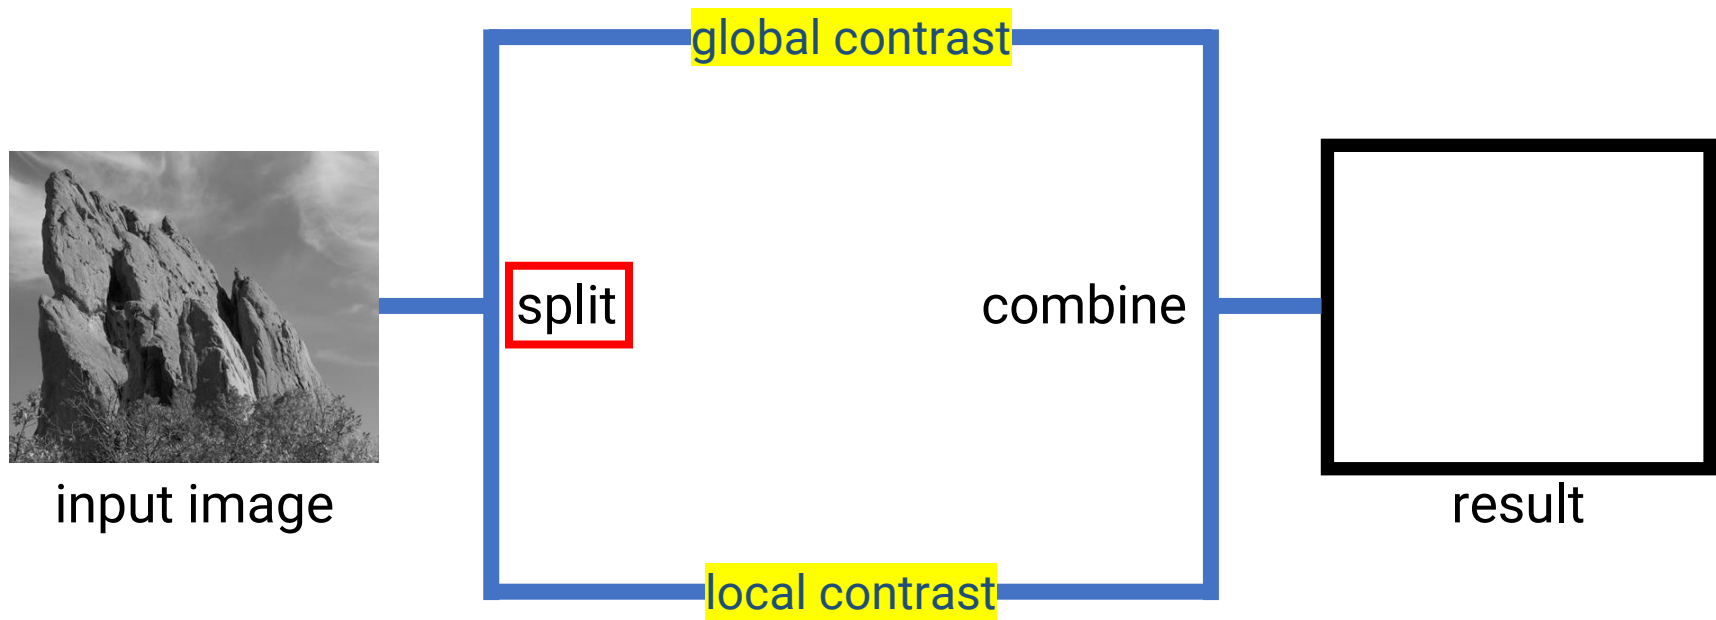
model



result

Photographic Style Transfer (cont.)

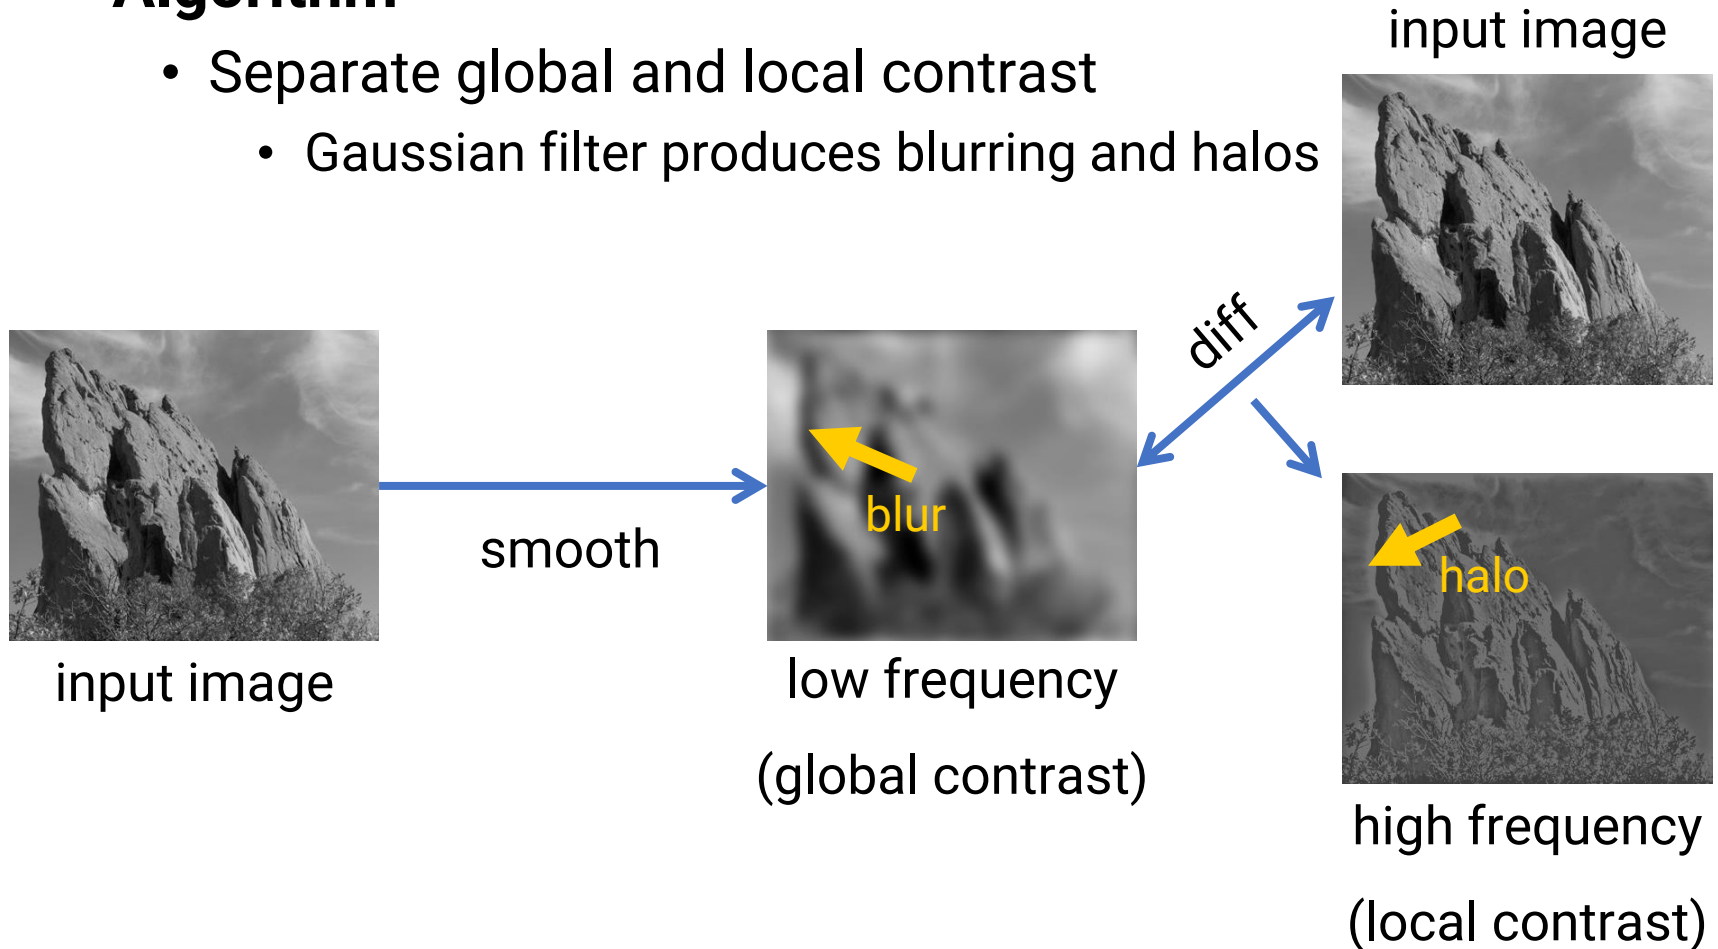
- Algorithm



Photographic Style Transfer (cont.)

- **Algorithm**

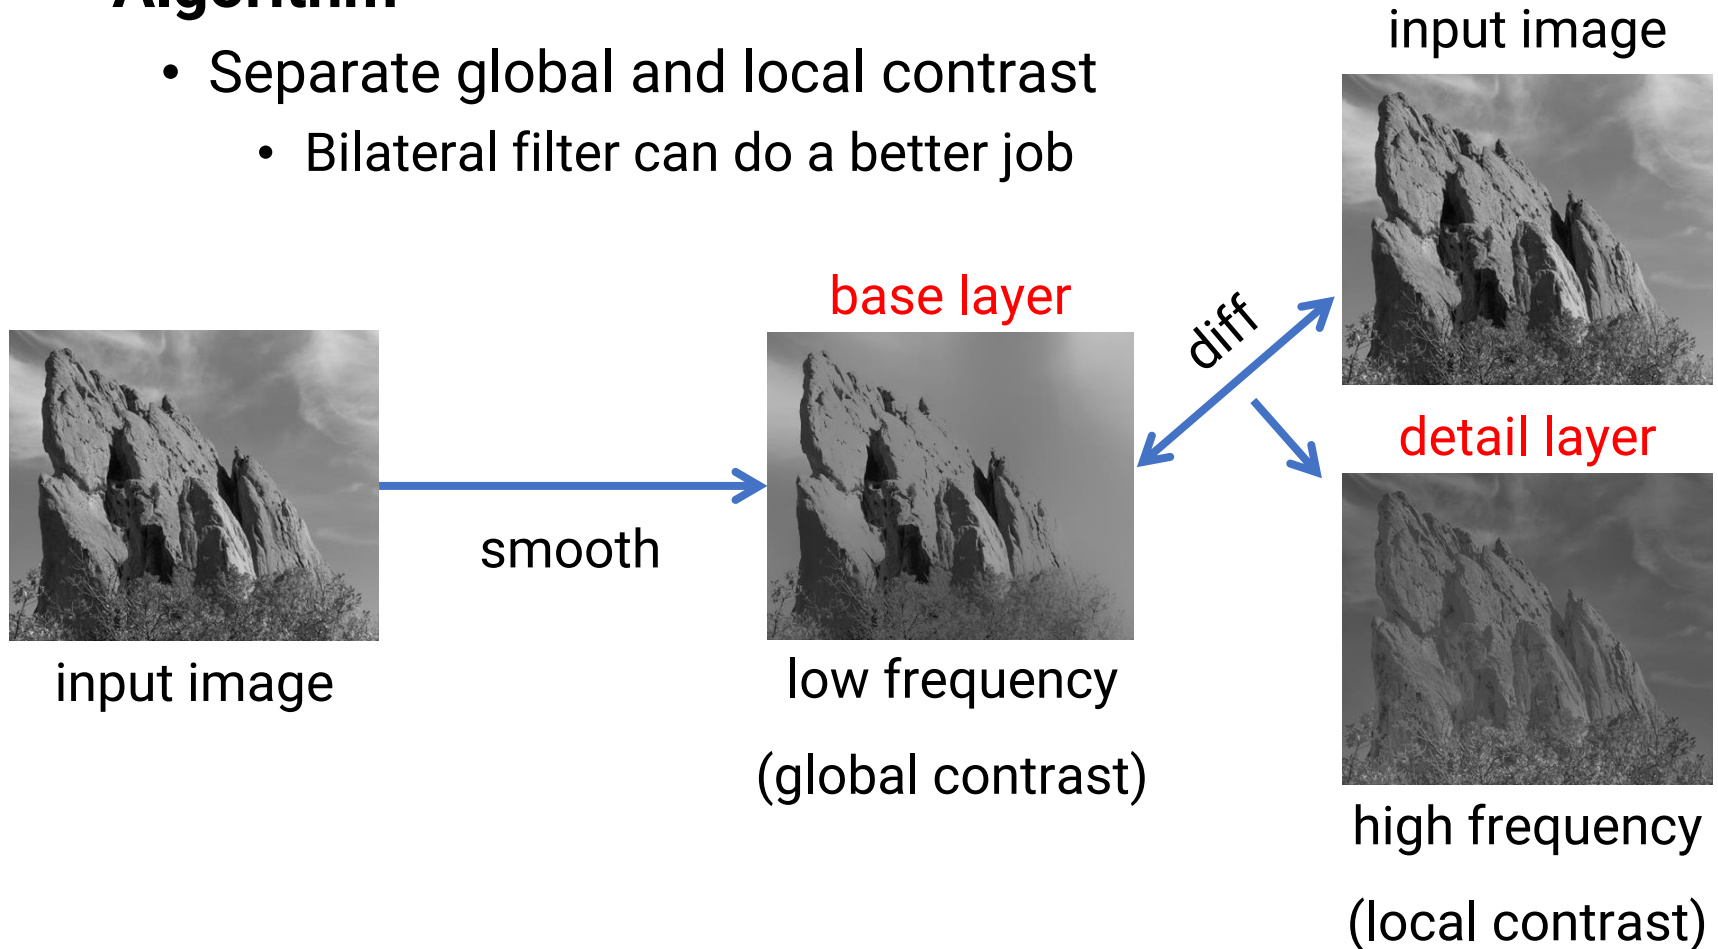
- Separate global and local contrast
 - Gaussian filter produces blurring and halos



Photographic Style Transfer (cont.)

- **Algorithm**


- Separate global and local contrast
 - Bilateral filter can do a better job



Photographic Style Transfer (cont.)

- **Algorithm**


- Separate global and local contrast



Photographic Style Transfer (cont.)

- **Algorithm**


- Adjust global contrast by histogram matching



Photographic Style Transfer (cont.)

- **Algorithm**


- Adjust local contrast by uniformly scaling



Photographic Style Transfer (cont.)

- **Algorithm**

- Sometimes the local contrast is not uniform



smooth region

⇒ low textureness

textured region


⇒ high textureness

Photographic Style Transfer (cont.)

- **Algorithm**
 - Textureness computation



input




textureness

Photographic Style Transfer (cont.)

- **Algorithm**

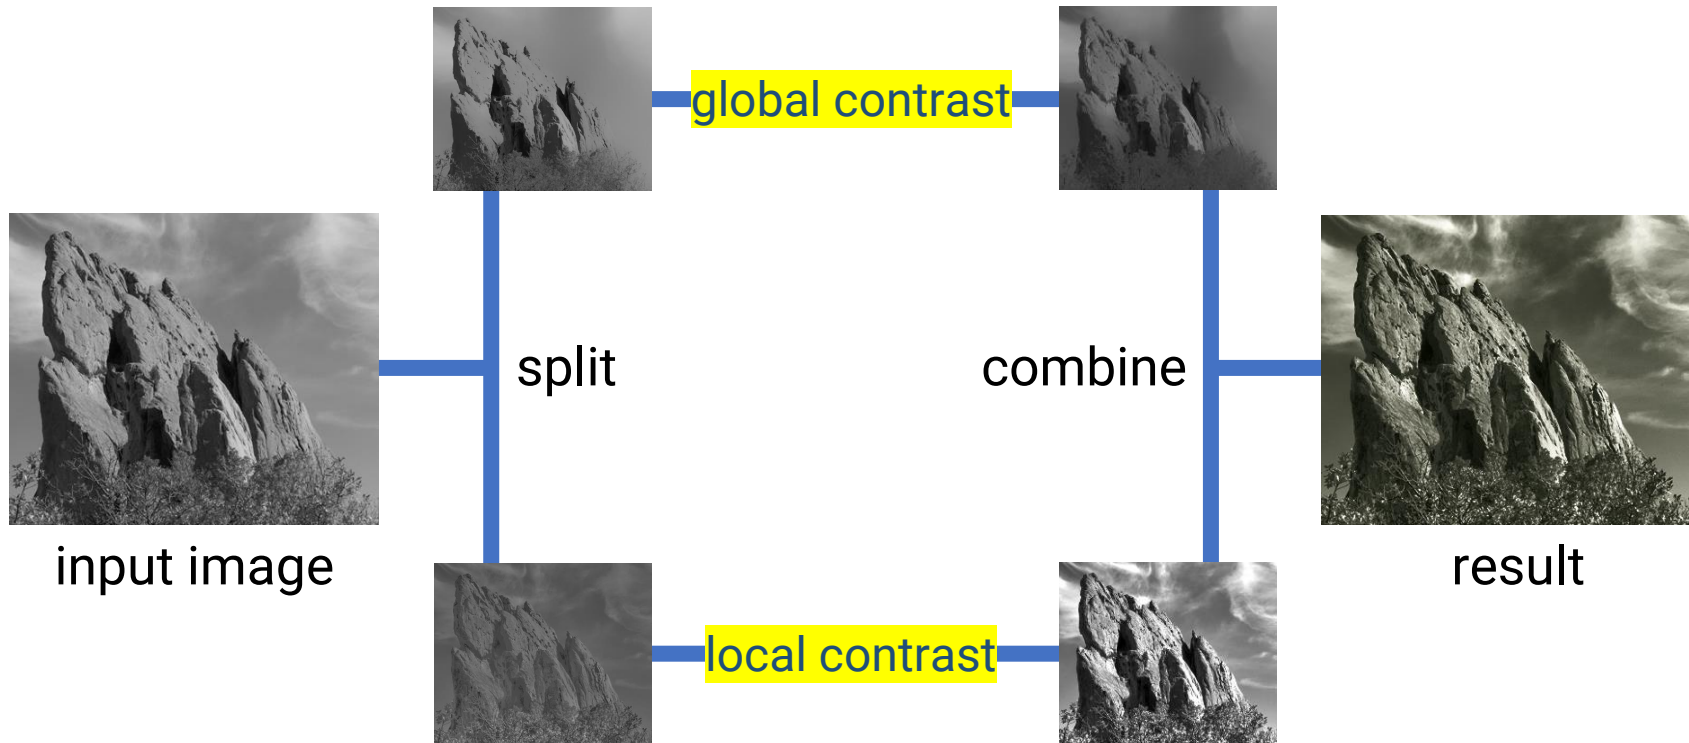
- Non-uniformly increase local contrast based on textureiness



Photographic Style Transfer (cont.)

- **Algorithm**

- Combine global and local contrast




Photographic Style Transfer (cont.)

- Results



Photographic Style Transfer (cont.)

- Results



Outline

- Overview
- Image compression
- Image manipulation
- **Image scaling**

Forward Mapping

5 x 5

(0, 0)

<div></div>	<div></div>			
<div></div>	<div></div>			
	(1, 1)			

Holes!

10 x 10


(0, 0)

<div></div>	?	<div></div>							
?	?								
<div></div>		<div></div>							
		(2, 2)							


$$(x', y') = (2x, 2y)$$

Inverse Mapping

5 x 5



10 x 10




$$(x, y) = (x'/2, y'/2)$$

No holes!
But how to determine
the color?


Image Scaling (cont.)

- Three strategies to obtain an estimation of X




nearest neighbor

P_3 is closest
Use P_3 's pixel value



bilinear interpolation

$$(1-a)(1-b)P_1 + (a)(1-b)P_2 + \\ (1-a)(b)P_3 + (a)(b)P_4$$




bicubic interpolation


using curves to
compute the weight
(nonlinear)

Image Scaling (cont.)


- Three strategies to obtain an estimation of X




1D nearest-neighbour




Linear




Cubic



2D nearest-neighbour



Bilinear



Bicubic

Image Scaling (cont.)

- Example: scale an image from 160 x 120 to 800 x 600





Image Scaling (cont.)

- Example



nearest neighbor



bilinear interpolation



bicubic interpolation

Image Scaling (cont.)

- Example



nearest neighbor



bilinear interpolation



bicubic interpolation

original  × 10

Image Scaling (cont.)

- Example



original



nearest neighbor



bilinear interpolation



bicubic interpolation

